

Graphs: Undirected Graphs

- Path: Sequence of vertices connected by edges
 - Cycle: Path whose first and last vertices are the same
 - Two vertices are connected if there is a path between them
 - If an edge exist between two vertices, they are adjacent to each other
 - The degree of a vertex in an undirected graph is the number of edges that are directly connected to it
 - Self Loop: An edge loop that connects a vertex to itself
- ~~• Disjointed Graphs~~
- Edges that directly connect the same pair of vertices are parallel

- Graphs can be represented by maintaining a list of edges via Array or Linked List
 - Basic operations are not performed efficiently
- Graphs can be represented through an adjacency-matrix
 - Needs an enormous amount of space to represent large graphs
 - Good for small graphs and dense graphs
- Graphs are mostly Sparse, so adjacency matrix is not a good implementation
- Graphs can be represented by an adjacency-list
 - Maintains a vertex-indexed array of lists
 - Cannot represent parallel edges
- In practice, adjacency-list graphs are used because real world graphs are sparse

- Sparse graphs contain a huge number of vertices, and small average vertex degree
- Bag data structure is used for the adjacency-list Graph
- Three algorithms to search through a graph:
 - Union find
 - Depth First Search (DFS)
 - Breadth First Search (BFS)
- DFS is recursive algorithm
- Recursive algorithms can be converted to Iterative algorithms via the stack data structure
- Undirected graphs have two edges: an inward edge and an outward edge
- Vertex connections can be direct or indirect
- DFS Pseudo Code:


```
DFS (to visit a vertex v)
```

 - Mark v as visited.
 - Recursively visit all unmarked vertices w adjacent to v.
- When DFS is called on a vertex, it gets marked
- Adjacency list is based off of adding edges/vertices to the graph based on input
 - ↳ Adjacency list is based on the initial input
- Output tree:


```

graph LR
    subgraph Graph [Graph]
        0((0)) --- 1((1))
        0((0)) --- 2((2))
        1((1)) --- 3((3))
        1((1)) --- 2((2))
        2((2)) --- 3((3))
        2((2)) --- 4((4))
        3((3)) --- 4((4))
        4((4)) --- 5((5))
    end
    subgraph Tree [Output Tree]
        0((0)) --- 2((2))
        0((0)) --- 1((1))
        0((0)) --- 3((3))
        0((0)) --- 4((4))
        0((0)) --- 5((5))
    end
  
```
- Generating the output tree requires the adjacency list
- Every marked vertex is connected to the source vertex
- Marking ensures that each vertex is visited once
- The maximum number of edges traversed is $2 \cdot E$

- If the source is connected to all edges, then the max. time is $V + 2E$

- DFS does not give the shortest path

↳ BFS does find shortest path

- Breadth-first Search pseudo-code:

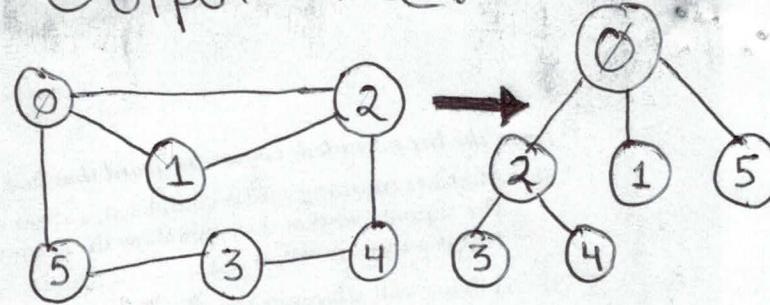
BFS (from a source vertex s)

- Put " s " onto a FIFO queue, and mark " s " as visited
- Repeat until the queue is empty:
 - Remove the least recently added vertex
 - Add each of " v 's" unvisited neighbours to the queue, and mark them as visited

- BFS uses a queue instead of a stack

↳ Not recursive like DFS

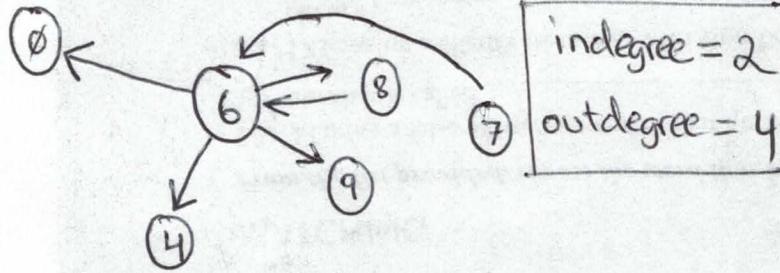
Output Tree:



- The "edgeTo[]" array is needed to generate the output tree
- BFS computes shortest paths from " s " to all other vertices in time proportional to $E + V$
- DFS uses a stack
- BFS uses a queue
- BFS takes the least recently added vertex
- DFS takes the most recently added vertex
- All components in a graph are reflexive
 - ↳ " v " is connected to " v "
 - ↳ self-loop
- DFS is used to find connected components in a graph
 - ↳ Takes $(V + E)$ time and space
- In practice, union find is faster than DFS

Graphs: Directed Graphs

- Directed graphs are also called digraphs
- In a digraph, edges are one way; contain a direction
- The indegree of a vertex is the number of directed edges that point to the vertex
- The outdegree of a vertex is the number of directed edges that point out from it



- Directed graphs are represented using adjacency lists
 - ↳ identical to undirected graphs
- DFS algorithm works for both undirected and directed graph
 - ↳ same with BFS
- Multiple shortest path is achieved via BFS by enqueueing all source vertices

- A directed acyclic graph (DAG) is a digraph with no directed cycles

↳ used to solve precedence constrained scheduling

- Topological sort, sorts the vertices of a digraph such that precedence constraints are maintained

↳ all directed edges point to a vertex later in the order, and there is no back-pointing

↳ is a linear ordering

- DFS is used to compute topological sort

- There are 3 vertex orderings:

↳ Preorder: Put the vertex on a queue before the recursive calls

↳ Postorder: Put the vertex on a queue after the recursive calls; when it's done

↳ Reverse Postorder
 Put the vertex on a stack after the recursive call

- The reverse postorder in a DAG is a topological sort
 - Topological sort does not work and is not possible if the graph has a cycle
 - In topological order:
 - The first vertex has an indegree of 0
 - The last vertex has an outdegree of 0
 - A digraph has topological order IFF it is acyclic
 - With DFS, topological order can be calculated in time proportional to $(V + E)$
 - There can be multiple topological sorts in a DAG
 - Second-to-last vertex in topological order can only point to the last vertex
 - Alternatively, second vertex in postorder can only point to first vertex
 - DFS is used to detect directed cycles in a digraph
-
- Vertices "v" and "w" are strongly connected if there is a directed path from "v" to "w" and a directed path from "w" to "v"
 - Note: This is a path, not an edge.
 - Each component is strongly connected to itself
 - "v" is strongly connected to "v"
 - The Kosaraju-Sharir algorithm is used to find the strongly connected components of a digraph
 - Step 1: Run DFS on G^R to compute reverse postorder
 - Step 2: Run DFS on G_I , considering vertices in order given by first DFS
 - Time Complexity: $(E + V)$
 - Kosaraju-Sharir uses DFS to find the Strongly connected components

Graphs: Minimum Spanning Trees

- MSTs are used for edge weighted graphs
- The purpose of an MST is to find a minimum weight set of edges that connects all of the vertices
- A spanning tree of a Graph, G_1 , is a subgraph, T , that is:
 - Connected
 - Acyclic
 - Includes all of the vertices in G_1
- An edge-weighted graph can have many spanning trees, but only one MST
 - However, in some cases you can have more than one MST; if some edge weights are identical
- The cut property is used to prove the correctness of an MST
- A cut in a graph is a partition of its vertices into two (non-empty) disjoint sets

- A crossing edge connects a vertex in one set with a vertex in the other set
- Cut Property: Given any cut, the crossing edge of minimum weight is in the MST
 - Out of all the crossing edges, the one with the minimum weight will be included in the MST
- Kruskal's Algorithm:
 - Step 1: Consider edges in ascending order of weight
 - Step 2: Add next edge to MST unless doing so would create a cycle
- In Kruskal's algorithm, the edges are sorted by weight from smallest to largest. Then, you keep on adding edges unless it creates a cycle or all edges are exhausted
- DFS can be used to identify cycles; in $O(n)$ time per cycle check
- Union find can be used to identify cycles; in $\log(n)$ time
 - works better than DFS

- CIPS
- Kruskal's algorithm
compute MST in $\Theta(|E| \log |E|)$ time
 - ↳ runtime is dominated by the sorting of edges
 - Prim's Algorithm:
 - Step 1: Start with vertex 0 and greedily grow tree, T
 - Step 2: Add to T the min weight edge with exactly one endpoint in T
 - Step 3: Repeat until $V - 1$ edges in MST
 - Two versions to implement Prim's Algorithm
 - Lazy
 - Both maintain a priority queue
 - Eager
 - In Prim's Algorithm, you maintain a list of all visited vertices, and you use that list each time you want to find the edge with a minimum weight
 - Lazy Prim's Algorithm computes the MST in time proportional to $(E \log E)$
 - Eager version of Prim's Algorithm uses extra space proportional to V and time proportional to $(E \log V)$ in the worst case to compute the MST
 - The running time of Kruskal's Algorithm is greater than the Eager implementation of Prim's Algorithm
 - ↳ However, there are instances where one might have a better performance than the other
 - ↳ i.e. Dense graph, sparse graph, etc.
 - A Minimum Spanning Forest (MSF) is the set of all MSTs
 - ↳ the collection of MSTs is an MSF
 - Lazy Prim's Algorithm does not check the edge weights, it adds them to the Priority Queue, regardless
 - ↳ Eager Prim's Algorithm does check the weights and only adds the next ones!

Graphs: Shortest Paths

C4P4-1

- The goal is to find the shortest paths in an edge-weighted digraph
- BFS finds the shortest paths in undirected graphs and digraphs that do not have edge-weights ~~edge-weights~~
- The shortest-paths tree (SPT) for an ~~edge-weighted~~ edge-weighted digraph is a subgraph containing the designated vertex, "s", and all the vertices reachable from "s" that forms a directed tree rooted at "s" such that every tree path is a shortest path in the subgraph
- Every tree has ~~an~~ SPT solution, no matter the case
 - even trees with a single node
- Minimum Spanning Forest (MSF) is used for undirected edges
- The SPT does not need to have a path to all nodes
 - only the nodes that are reachable

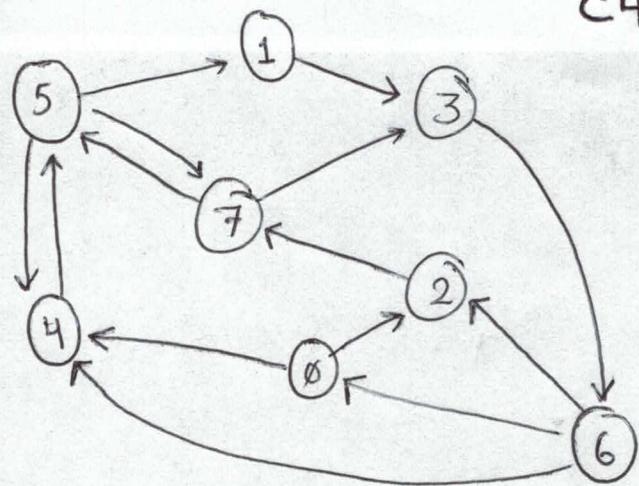


Figure: Shortest-Paths tree from 0

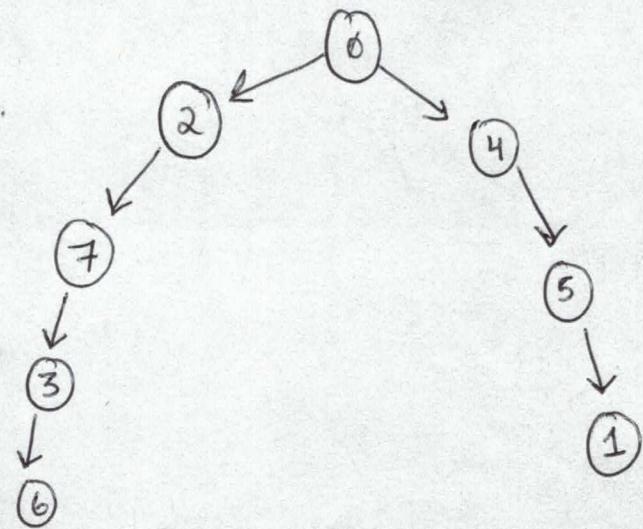
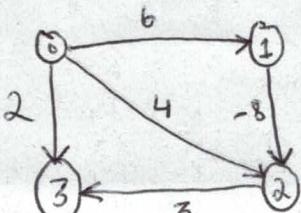


Figure: SPT drawn as tree rooted at 0 for figure above

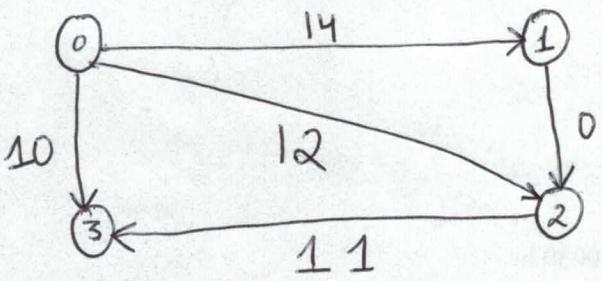
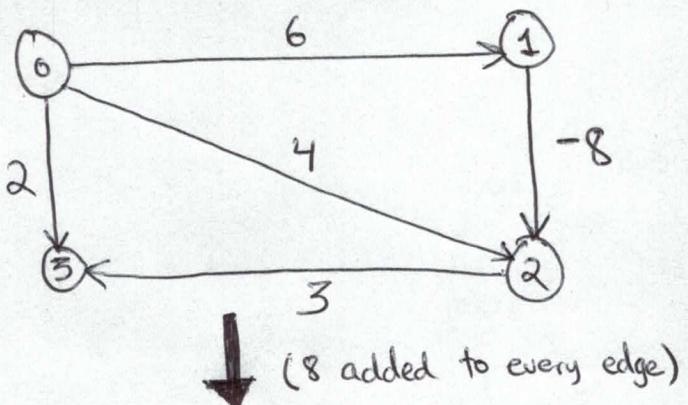
- Relaxing an edge means to update the "distTo[J]" and "edgeTo[J]" array
 - If an edge connecting vertex V and W gives a shorter path to W through V, then update both "distTo[W]" and "edgeTo[W]".
- In vertex relaxation, all the edges pointing from a given vertex are relaxed

- Generic Algorithm (to compute SPT from "s")
 - Initialize "distTo[s] = 0" and $\text{distTo}[v] = \infty$ for all other vertices
 - Repeat until optimality conditions are satisfied
 - ↳ Relax any edge
- If "v" is not reachable from "s", then "distTo[v]" remains at infinity
- The generic shortest paths algorithm does not specify in which order the edges are to be relaxed
 - ↳ 3 Efficient Implementations
 - Dijkstra's Algorithm
 - Topological Sort Algorithm
 - Bellman-Ford Algorithm
- Dijksta's is a greedy algorithm
- Dijksta's Algorithm cannot handle ~~negative weights~~ or negative cycles
- Topological sort cannot handle directed cycles
 - ↳ only works on DAGs
 - ↳ can handle negative weights
- Bellman-Ford algorithm cannot handle negative cycles
 - can handle negative weights
 - can handle directed cycles
- An SPT does not exist if there IS a negative cycle
- In Dijksta's Algorithm, the distance to all vertices ϵ is initialized to infinity
 - ↳ However, the distance to the source vertex is initialized to 0
- Dijksta's Algorithm uses extra space proportional to "V" and time proportional to $(E \log V)$
 - ↳ Only true when using a min. priority queue implemented as a binary heap
- Topological sort computes the SPT in edge-weighted DAG in time proportional to $(E + V)$
 - ↳ uses DFS
- Dijksta's Algorithm does not work for negative weights



• Dijksta selects vertex 3, immediately after 0. But, the shortest path from $0 \rightarrow 3$ is $0 \rightarrow 1 \rightarrow 2 \rightarrow 3$

- Adding a constant to every edge-weight, to remove the negatives, does not work



- Adding a constant to every edge-weight changes the properties of the graph

↳ Shortest path from $0 \rightarrow 3$ is $0 \rightarrow 3$

- A negative cycle is a directed cycle where the sum of the edge weights is negative

$$5 \rightarrow 4 \equiv -0.66$$

$$4 \rightarrow 7 \equiv 0.37$$

$$7 \rightarrow 5 \equiv 0.28$$

→ negative cycle: $(-0.66 + 0.37 + 0.28)$

- An SPT (Shortest path) exists IFF there are no negative cycles

Bellman - Ford Algorithm

- Initialize $\text{distTo}[s] = 0$ and $\text{distTo}[v] = \infty$ for all other vertices
- Repeat V times:
 - Relax each edge
- Running time: $O(EV)$
- ~~Relaxation is done to keep track of vertices~~

- Queue based Bellman - Ford computes the shortest path in any edge-weighted digraph with no ~~edge~~ negative cycles in time proportional to $(E \times V)$ in the worst case

↳ However, it is much faster in practice; takes $(E + V)$ time

- Cost Summary for single source shortest-path implementation

Algorithm	Restriction	Typical case	Worst Case	Extra Space
topological sort	no directed cycles	$E + V$	$E + V$	V
Dijkstra (binary heap)	no negative weights	$E \log V$	$E \log V$	V
Bellman - Ford	no negative cycles	EV	EV	V
Bellman - Ford (queue-based)	no negative cycles	$E + V$	EV	V