

2GA3 Tutorial #3

DATE: October 1st, 2021

TA: Jatin Chowdhary

Pay Attention

- Next slide is very important
- Listen to this carefully
 - This information might save your life one day
 - Seriously

FIPPA

- What is FIPPA?
 - Freedom of Information and Protection of Privacy Act
- Why care?
 - Use it for secure/confidential communication
 - Great for disclosing information you do not want anyone else but the intended party to see
 - i.e. Allergies (Food, chemicals, etc.)
 - i.e. Inform on person X to person Y
- How to properly use it?
 - By law it's already used to protect communication
 - But it's good practice to state it
 - i.e. On next slide

FIPPA Example

Dear Professor X,

I am writing this email to inform you about an incident that occurred on March 3rd, 2017, with a person who goes by the name of Logan. Mr. Logan consumed a peanut butter and jelly sandwich while being fully aware of my food allergies. Had I not taken the appropriate action to protect myself, serious harm could have occurred. Please take necessary action to alleviate this issue. Furthermore, I expect the contents of this email to remain confidential in accordance to **FIPPA**. Please do not disclose my identity, or any other identifying information to any party involved directly or indirectly, without my consent.

Regards,

Scott Summers

Participation

- The script
 - Everyone will be called upon, once
- Why?
 - This is the best way to learn
 - Repetition is the language of the brain
 - Speech > Writing
 - It's how our brain are wired
 - Millions of years of evolution
 - Participating > Taking notes
 - Everything will be posted
 - Last tutorial on Friday is recorded
 - Just pay attention and (immediately) ask questions!

I Don't Know

- If I call on you, and you don't know, just say, "I don't know"
 - Ain't nothing wrong with this
 - Everyone is here to learn
 - Proudly say, "I don't know"
 - You'll learn better this way; it takes away the pressure
 - It's the best answer
- If you still don't get it, ask for clarification!
 - No one gets *it* right away
 - It takes a few tries, but eventually you will get it, and that's all that matters

Clarification (1)

- Made a **mistake** about **sw** and **sd** during last week's tutorial
 - The mistake is underlined below:

64-Bit:

```
sub x30, x28, x29 // Compute [i - j]
slli x30, x30, 3 // Multiply by 8 to get byte offset
add x3, x3, x30 // Calculate location in memory
ld x30, 0(x3) // Load A[i - j]
sd x30, 64(x11) // Store in B[8]
```

32-Bit:

```
sub x30, x28, x29 // Compute [i - j]
slli x30, x30, 2 // Multiply by 4 to get byte offset
add x3, x3, x30 // Calculate location in memory
lw x30, 0(x3) // Load A[i - j]
sw x30, 32(x11) // Store in B[8]
```

Clarification (2)

- I can't remember what I said, but:
 - The difference in 64-bit and 32-bit RISC-V code (in the previous slide) is:
 - **slli**
 - 3 vs. 2
 - **ld** vs. **lw**
 - **sd** vs. **sw**
 - **64**(x11) vs. **32**(x11)
 - The number to the left of the brackets is the byte offset in memory
 - The data is stored memory that corresponds to the value in **x11** (base address), with an offset of 32 or 64
 - Why are 8 (64-bit) and 4 (32-bit) so special?
 - Explained on next slide

Mistake

- What I said last week was a bit misleading
 - Does everything make sense?
- Sorry for the mistake
 - If you make a mistake, come clean
 - Great way to learn, because you'll never forget
 - **Be honest**
- We are all here to learn

Word

- Question: How big is a *word*?
- Options:
 - 16-bits
 - 8-bits
 - 31-bits
 - 32-bits
 - 33-bits
 - None of the above
 - As big as you want

Double Word

- Question: How big is a *double word*?
- Options:
 - 60-bits
 - 61-bits
 - 62-bits
 - 63-bits
 - All of the above
 - None of the above

Byte Offset (1)

- In a 32-bit system, the byte offset is calculated by multiplying by **4**
 - For 64-bit, multiply by **8**
- Why these specific numbers?
 - First of all, what does 32-bit mean?
 - Question: If I say the architecture of this computer is 32-bits, what does that mean?
 - Question: What is 1 bit?
 - In the context of binary numbers
 - How many bits in a byte?
 - Calculation
 - Next slide

Byte Offset (2)

- For a 32-bit system, we divide 32 bits by 8 bits to get the size of the block, in bytes
 - $32 \text{ bits} / 8 \text{ bits} = 4 \text{ bytes}$
- What is the size of the block, in bytes, for a 64-bit system?
 - $\mathbf{X} \text{ bits} / \mathbf{Y} \text{ bits} = \mathbf{Z} \text{ bytes}$

128-Bit Architecture

- Question: To calculate the byte offset in a 128-bit computer, what number would you use to multiply?
- Options:
 - 12 bytes
 - 14 bytes
 - 16 bytes
 - 18 bytes
 - 20 bytes

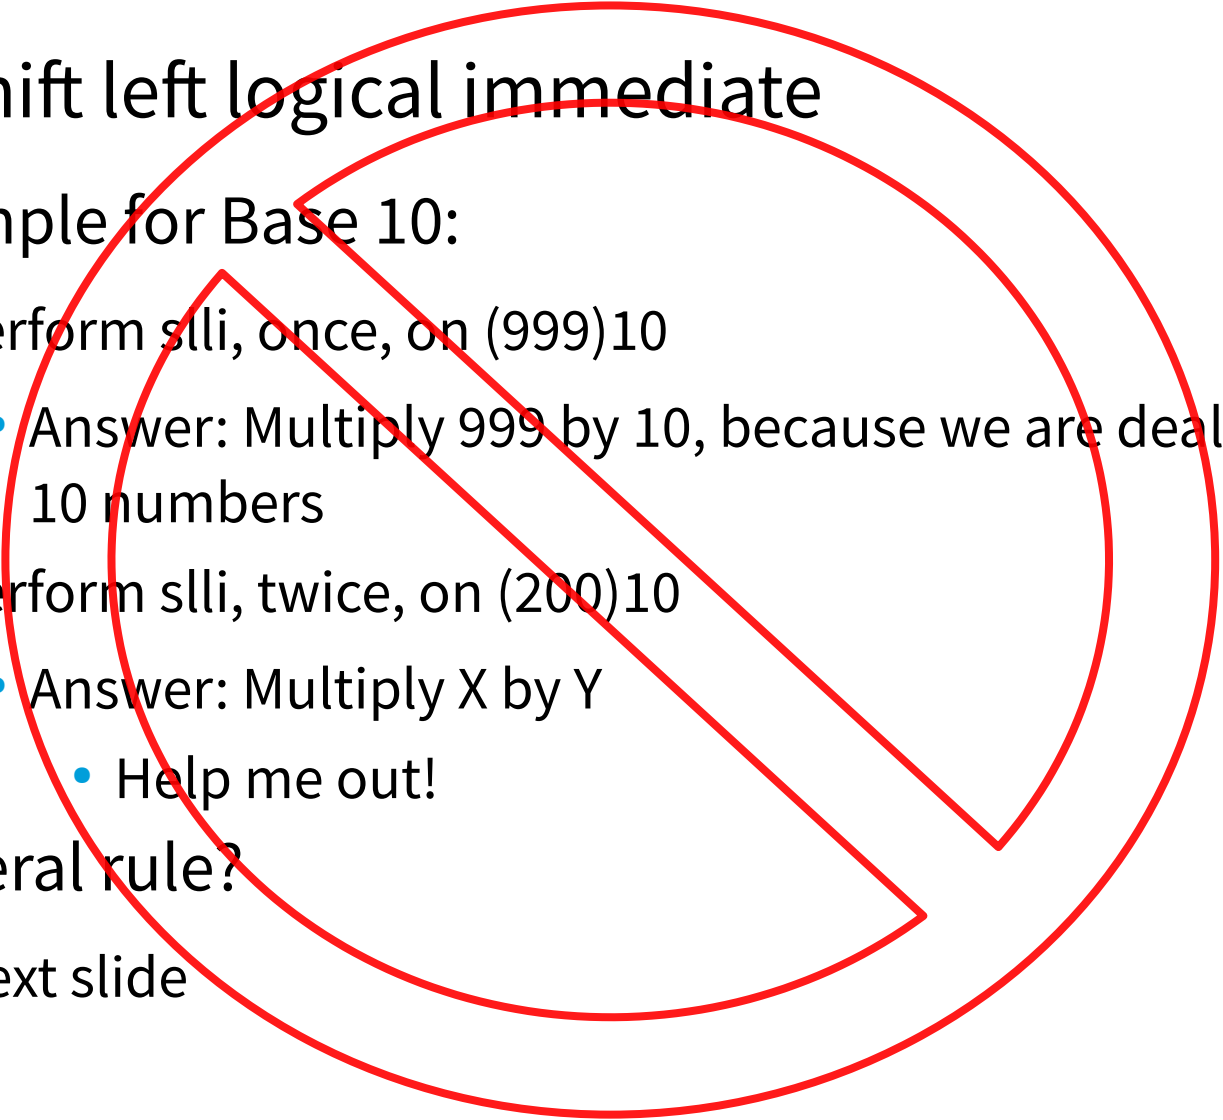
Any Questions?

- Does that make sense?
- All good?
- Redo it?

Food

- If you have food allergies, let me know!
 - Same with dietary restrictions
 - Use FIPPA
- Anyone hungry?
 - Snacks in my bag
 - Whoever asks the most questions, gets a piece

Recap SLLI (1)

- slli = shift left logical immediate
 - Example for Base 10:
 - Perform slli, once, on $(999)_{10}$
 - Answer: Multiply 999 by 10, because we are dealing with base 10 numbers
 - Perform slli, twice, on $(200)_{10}$
 - Answer: Multiply X by Y
 - Help me out!
 - General rule?
 - Next slide
- 

Recap SLLI (2)

- General rule is:
 - For every *slli*, add a 0
 - Translation: Multiply by **X**, where **X** is the base of the numbers you are working with
 - i.e. (*slli*, 1) on $(900)_{10} = 900 \times \mathbf{10} = 9000$
 - i.e. (*slli*, 2) on $(11)_2 = 11 \times \mathbf{2} \times \mathbf{2} = 01100$
 - Note: This is NOT "eleven"
 - Note: Only use this trick for $(\text{Base})_2$

More About SLLI

- Going overboard!
 - Assume you are working with a 4-bit architecture, where the maximum possible number is $(1111)_2$.
 - What is the result of: $(slli, 2)$ on $(1001)_2$?
 - Remember: The largest possible number is $(1111)_2$
 - Is the answer: $(100100)_2$
- Conclusion: Information is _ _ _ _ !
 - Same thing happens for *srli*
 - *srli* = Shift Right Logical Immediate

Any Questions?

- Does that make sense?
 - Any questions about anything I've said?
- Ask now or forever...

No Time

- I still gotta clarify:
 - CMOS/PMOS/NMOS
 - *Note: Was discussed during office hours*
 - 32-Bit vs. 64-Bit
 - 'strcpy' Example
- Always running out of time
 - Will try to cover next week
 - Or the week after that
 - Or after that
 - Or after that

Question #1 (2.13)

- For the following instruction, what is the instruction type? And what is the hexadecimal representation?

sw x5, 32(x30)

rs2 imm rs1

- Answer:
 - Next slide
 - Pay attention; this is **VERY** important

Answer #1 (2.13)

- What is the instruction type?
 - The instruction is **sw**, which stands for *store word*
 - Hence, it is a `_____` instruction, making it an `_` type
- Refer to lecture slides:



RISC-V Instruction Types

R-type (R for registers) instructions use 3 register operands (2 source registers and one destination register).

I-type (I for immediate) instructions use 2 register operands (1 source, 1 destination) and one 12-bit immediate.

S-type (S for stores) instructions use 2 register operands (2 source) and one 12-bit immediate.

SB-type (conditional Branch, fields like the S)

U-type (Upper immediate format)

UJ-type (Unconditional jump)

Answer #1 (2.13)

- What is the hexadecimal representation?
 - First, get the binary representation
 - Then, convert to hexadecimal
- Answer:
 - We've already established that we are dealing with an S-type instruction
 - So let's see what the format is for an S-type instruction
 - Next slide

RISC-V S-format Instructions

immediate[11:5]	rs2	rs1	funct3	immediate[4:0]	opcode
7 bits	5 bits	5 bits	3 bits	5 bits	7 bits

- Different **immediate format for store instructions**
 - **rs1**: base address register number
 - **rs2**: source operand register number
 - immediate: offset added to base address
 - Split so that rs1 and rs2 fields always in the same place

Answer #1 (2.13)

- So, we need:
 - *opcode* (7 bits)
 - *funct3* (3 bits)
 - *rs1* (5 bits)
 - *rs2* (5 bits)
 - *immediate* (5 + 7 bits)
- How do we get this information?
 - From the RISC-V instruction list!
 - $\text{opcode} = 0100011$
 - $\text{funct3} = 010$
 - $\text{rs1} = (30)_{10}$
 - $(0111110)_2 \rightarrow (11110)_2$
 - $\text{rs2} = (5)_{10}$
 - $(101)_2 \rightarrow (00101)_2$
 - $\text{Immediate} = (32)_{10}$
 - $(100000)_2$
- Now what?
 - Next slide

Answer #1 (2.13)

- Put it all together!
 - opcode = 0100011
 - funct3 = 010
 - rs1 = 11110
 - rs2 = 00101
 - immediate = 100000
 - Immediate[4:0] = 00000
 - [4:0] means: Take the first 5 least significant bits; from the 0th bit to the 4th bit
 - Immediate [5:11] = 0000001
 - [5:11] means: The range is from the 5th bit to the 11th bit
- The final answer is:
 - *immediate[5:11]* + *rs2* + *rs1* + *funct3* + *immediate[4:0]* + *opcode*
 - 0000001 00101 11110 010 00000 0100011

Answer #1 (2.13)

- How to convert to hexadecimal?

1. Take our answer and remove all spaces:

00000010010111110010000000100011

2. Add a space after every 4th bit:

0000 0010 0101 1111 0010 0000 0010 0011

3. Use the table, on the right, to convert to hex:

0x**025f2023**

0000	0	1000	8
0001	1	1001	9
0010	2	1010	A
0011	3	1011	B
0100	4	1100	C
0101	5	1101	D
0110	6	1110	E
0111	7	1111	F

Clarification For #1 (2.13)

- **Question:** Why is the immediate value separated into 2 parts?
 - **Answer:** To keep the format of the instructions consistent, which makes decoding easier
 - For more information: See lecture slides
- **Question:** What is *0x*?
 - **Answer:** It is convention to prefix hexadecimal numbers with *0x*. It tells the compiler (and you) that the number is in hexadecimal
 - i.e. *0xFFFF*
 - Hexadecimal
 - i.e. *0b11101*
 - Binary
 - i.e. *0o35*
 - Octal
 - We denote different bases via subscripts, the compiler/parser uses *0x*

The Big Picture!

- Everything in a computer's memory is just 1's and 0's
 - i.e. 00000010010111110010000000100011
 - Registers, Cache, Main memory, HDD, etc. contain 0's and 1's
- So when you send an instruction to the CPU, you're just sending a bunch of 0's and 1's like the example above
 - The 1's and 0's are generated by the *waves* (Draw This)
 - *Waves* = Low voltage and high voltage
 - The more waves there are, the more _____ we can send
 - The _____ are sent from main memory to the CPU through the _____
 - The _____ rate of a CPU measures the "_____"

Question #2 (2.17)

- Assume that the following registers contain:
 - **x5** = 0x0000AAAA
 - **x6** = 0x12345678
- For the register values shown above, what is the value of **x7** after the following sequence of instructions?
 - **slli** x7, x5, 4
 - **or** x7, x7, x6
- For the register values shown above, what is the value of **x7** after the following sequence of instructions?
 - **srli** x7, x5, 3
 - **andi** x7, x7, 0x1EF

Answer #2 (2.17)

- ***slli*** x7, x5, 4

= 0x0000AAAA × (2 × 2 × 2 × 2)

= 0x0000AAAA × (16)

= 0x000AAAA0

- ***or*** x7, x7, x6

0x**000**AAAA0 OR 0x**12345678**

000AAAA0: 0 0 0 0 0 0 0 0 0 0 0 0 0 0 1 0 1 0 1 0 1 0 1 0 1 0 1 0 0 0 0 0

12345678: 0 0 0 1 0 0 1 0 0 0 1 1 0 1 0 0 0 1 0 1 0 1 1 0 0 1 1 1 1 0 0 0

0123efef8: 0 0 0 1 0 0 1 0 0 0 1 1 1 1 1 0 1 1 1 1 1 1 1 0 1 1 1 1 1 0 0 0

Answer #2 (2.17)

- ***srli*** *x7, x5, 3*

= 0x0000AAAA ÷ (2 x 2 x 2)

= 0x0000AAAA ÷ (8)

= 0x00001555

- ***andi*** *x7, x7, 0x1EF*

= 0x00001555 AND 0x000001EF

= 1555 AND 01EF

= 0???

= 0145

Hexadecimal	Binary
555	0101 0101 0101
1EF	0001 1110 1111
145	0001 0100 0101

Clarification For #2 (2.17)

- **Question:** Do I have to convert to binary, perform the operation and then convert back?
 - **Answer:** There's a trick to performing these operations for base₁₆ numbers
 - Because base₂ numbers are 2^1 and base₁₆ numbers are 2^4
 - $2^4 = (2^1)^4$
 - You should try to figure it out, before asking me to spill the beans
 - Practice as much as you can!

Question #3 (2.23)

- Homework
 - You can find all questions/solutions on Teams
 - Check the *Files* tab
 - **Don't look at the answer before starting**

Question #4 (2.23)

- For the following C statement, write a minimal sequence of RISC-V assembly instructions that performs the identical operation.

Assume $x6 = A$, and $x17$ is the base address of C .

$A = C[0] \ll 4;$ // C Code

Question #4 (2.23)

- First of all:
 - What is **A** ?
 - What is **C[0]** ?
 - Where is **C[0]** stored?
 - Which instruction do we use?
 - What is **<<** ?
 - What is the equivalent instruction for this?

Answer #4 (2.23)

- Solution:

ld x6, 0(x17) // $x6 = C[0]$

slli x6, x6, 4 // $x6 = x6 * 16$