

2GA3 Tutorial #4

DATE: October 8th, 2021

TA: Jatin Chowdhary

Anonymous Feedback

- Spend ~4 Mins
 - Feedback on course/tutorial/content/etc.
- 100% Anonymous
 - I'm leaving the room
 - If you're paranoid about anonymity, write with your other hand
 - Then no one will be able to recognize your answers
- Be Honest
 - “The whole truth and nothing but the truth“
 - My cholesterol is high enough

Minor Stuff

- Expect Typos
 - Keyboard is breaking down
 - Keys are hard to press
 - The **S** & **T** keys
- Low Energy?
 - Let me know!
- Reading Week
 - Don't waste your time
 - Study for the midterm
 - It's gonna be long and hard

Major Stuff

- Note Taking?
 - No need
 - Everything is, or will be posted
 - Just participate
 - Because \mathbf{r} _____ is the language of the \mathbf{b} _____
 - *Participation* = \mathbf{R} _____
- When you go home, then take notes
 - You'll be surprised at your retention rate

Review (1)

- **Question:** What are registers?
- **Options:**
 - A) Measure of how fast a processor performs operations
 - B) Where CPUs temporarily store and process information
 - C) The bridge between the CPU and RAM
 - D) Used to store money at the store
 - E) Secondary storage in a computer

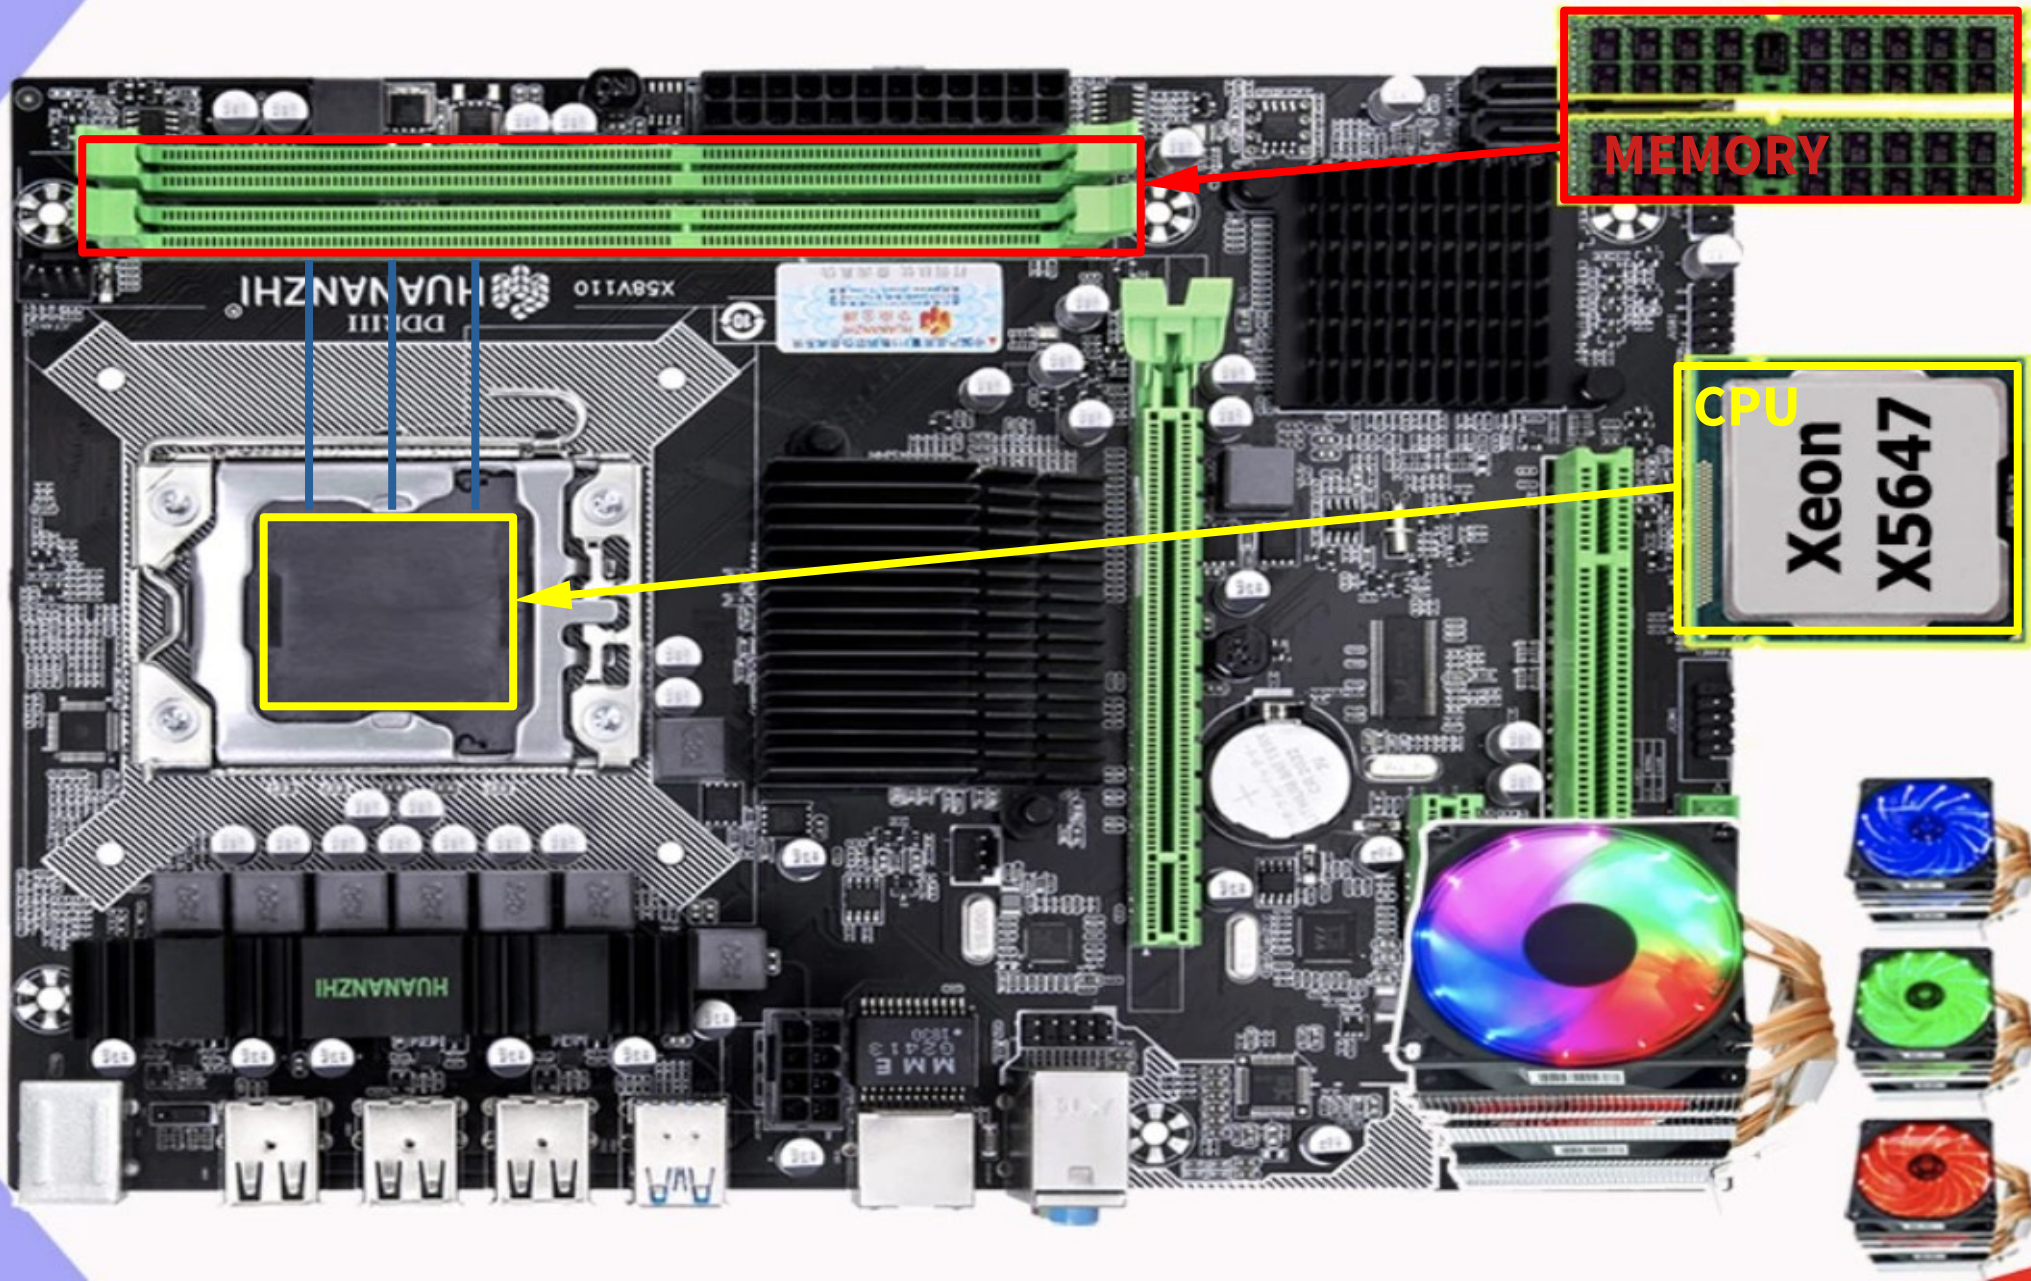
Review (2)

- **Question:** Which of the following are part of the main memory hierarchy?
- **Options:**
 - A) Registers
 - B) Cache
 - C) RAM
 - D) All of the above
 - E) None of the above

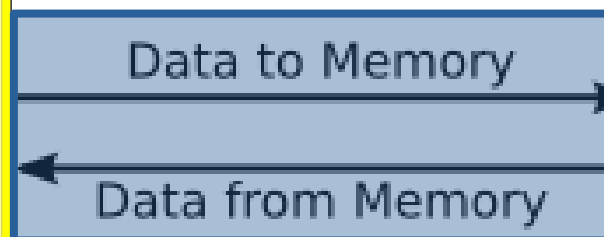
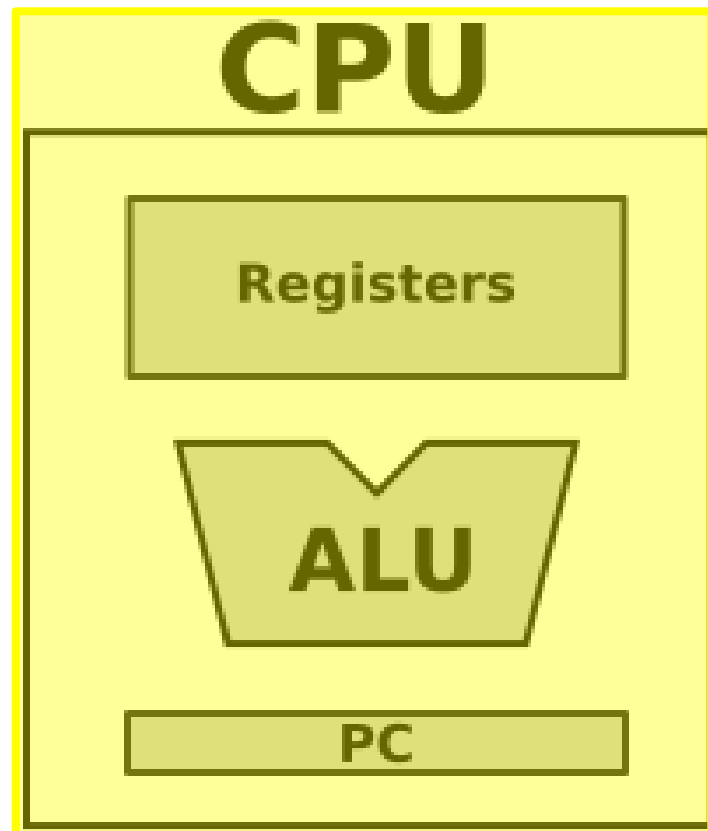
Review (3)

- **Question:** Main memory is...
- **Options:**
 - A) Fast and non-volatile
 - B) Slow and non-volatile
 - C) Fast and volatile
 - D) Slow and volatile
 - E) None of the above

CPU & Memory



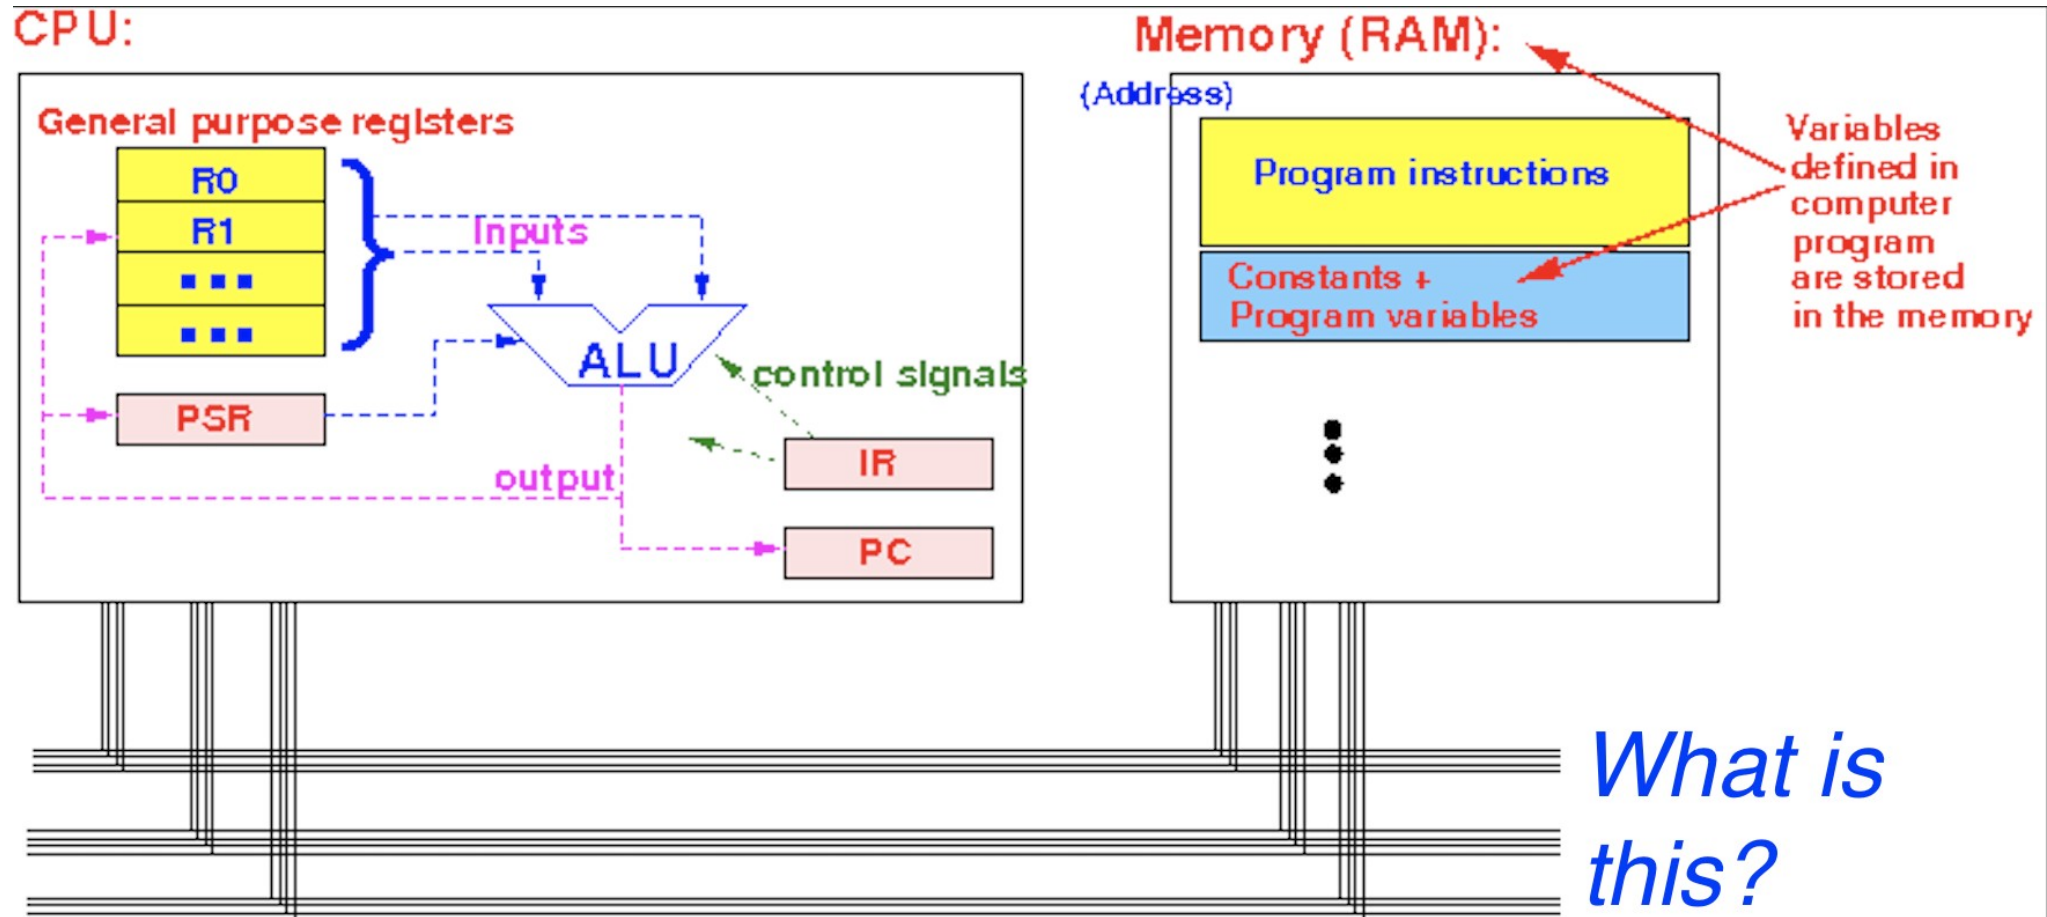
Abstract View



Memory

10001010	(Location 0)
00001100	(Location 1)
10111000	(Location 2)
01000001	(Location 3)
00001011	(Location 4)
11011101	(Location 5)
10110000	(Location 6)
01010010	(Location 7)
11111010	(Location 8)
01001100	(Location 9)
00100011	(Location 10)
00011010	(Location 11)
⋮	⋮

More Abstraction



- Everything in memory is either a 0 or a _
- Information is sent over the _ _ _
- R _ _ _ _ _ are found in the CPU
- ALU stands for...?
- RAM stands for...?
- PC stands for...?

Review (4)

- **Question:** Which of the following base-number systems are invalid?
- **Options:**
 - A) Base-4
 - B) Base-12
 - C) Base-20
 - D) Base-60
 - E) All of the above
 - F) None of the above

Review (5)

- **Question:** Can we use *base-3* or *base-4* number systems instead of *base-2*, in a computer? In other words, instead of a high/low voltage, there'll be 3, 4, etc. voltage(s)
- **Options:**
 - A) Yes
 - B) No
 - C) Maybe
 - D) I don't know
 - E) Can you repeat the question?
- *Next slide...*

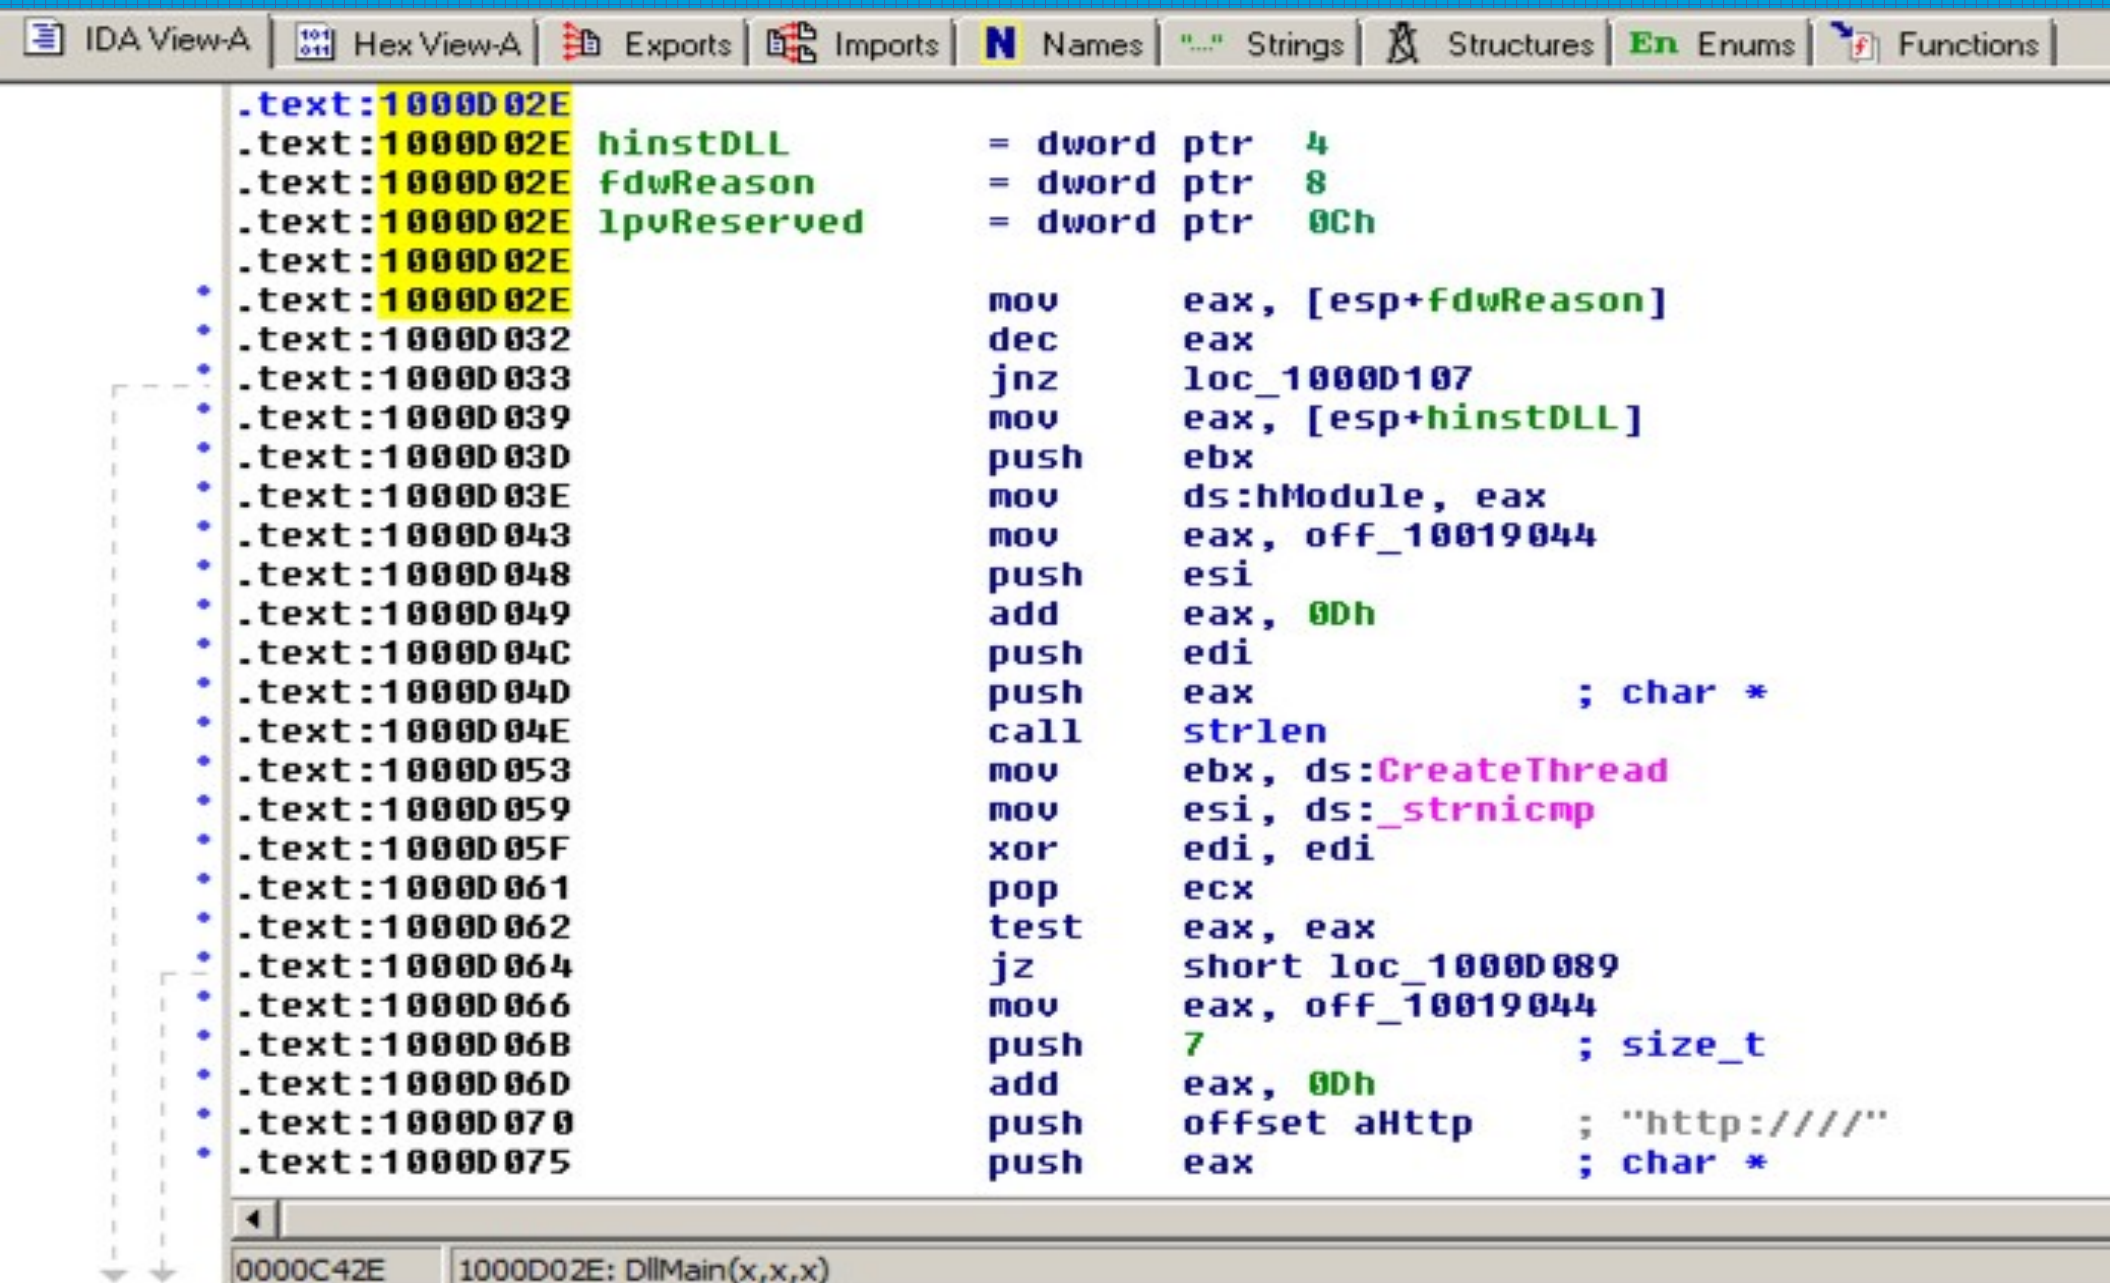
Base 'X'

- **Question:** Can we use *base-3* or *base-4* number systems instead of *base-2*, in a computer?
 - **Answer:** Yes we can.
 - But we don't, because:
 - Harder to implement
 - More complexity
 - Boolean logic was established before computers
 - Noise
 - Need a clear distinction between 0 (false) and 1 (true)

Why Bother With Hex?

- **Question:** What is the point of hexadecimal?
 - Asked by Emily
- **Answer:** It's a better way to represent binary-coded values
 - Each hexadecimal “character” represents 4 bits
 - *Bit* = **B**inary *D*igit
 - Reading hex numbers is easier than reading binary numbers
 - Used by computer scientists, chip designers, hackers, etc.
 - Hackers use disassemblers to study how a program works, find vulnerabilities, and recreate the original code (as much as they can)
 - i.e. IDA (*on next slide*)

IDA (Pro)



IDA View-A | 101 011 Hex View-A | Exports | Imports | Names | Strings | Structures | Enums | Functions

```
.text:1000D02E
.text:1000D02E hinstDLL      = dword ptr 4
.text:1000D02E fdwReason    = dword ptr 8
.text:1000D02E lpvReserved  = dword ptr 0Ch
.text:1000D02E
.text:1000D02E      mov     eax, [esp+fdwReason]
.text:1000D032      dec     eax
.text:1000D033      jnz     loc_1000D107
.text:1000D039      mov     eax, [esp+hinstDLL]
.text:1000D03D      push    ebx
.text:1000D03E      mov     ds:hModule, eax
.text:1000D043      mov     eax, off_10019044
.text:1000D048      push    esi
.text:1000D049      add     eax, 0Dh
.text:1000D04C      push    edi
.text:1000D04D      push    eax                ; char *
.text:1000D04E      call    strlen
.text:1000D053      mov     ebx, ds:CreateThread
.text:1000D059      mov     esi, ds:_strnicmp
.text:1000D05F      xor     edi, edi
.text:1000D061      pop     ecx
.text:1000D062      test    eax, eax
.text:1000D064      jz      short loc_1000D089
.text:1000D066      mov     eax, off_10019044
.text:1000D06B      push    7                  ; size_t
.text:1000D06D      add     eax, 0Dh
.text:1000D070      push    offset aHttp      ; "http:///"
.text:1000D075      push    eax                ; char *
```

0000C42E | 1000D02E: DllMain(x,x,x)

More Hexadecimal

- 4 bits (**binary digits**) can be represented by a *single* hexadecimal “character”
 - Much easier to read **0x000A**, than it is to read **0b00000000000001010**
 - Imagine how long a 32-bit or 64-bit number would be. Using hex makes sense
- **Question:** Which of the following is a hexadecimal number?
- **Options:**
 - 0x10101
 - 0b00AB
 - (00FF)₁₀
 - “ABC”
 - (54321)₆

Decimal	Binary	Hexadecimal
0	0000	0
1	0001	1
2	0010	2
3	0011	3
4	0100	4
5	0101	5
6	0110	6
7	0111	7
8	1000	8
9	1001	9
10	1010	A
11	1011	B
12	1100	C
13	1101	D
14	1110	E
15	1111	F

Review (6)

- **Question:** What is the difference between add and addi?
- **Options:**
 - A) There is no difference; it's the same thing
 - B) add is for large numbers; addi is for small numbers
 - C) add is for 32-bit; addi is for 64-bit
 - D) add adds the values in 2 registers; addi adds a constant value to a register
 - E) None of the above

Which Add?

- **Question:** Which instruction syntax corresponds to add & addi?
 - *rd, rs1, rs2* is the syntax for XXX
 - *rd, rs1, imm* is the syntax for XXX
- Note: XXX is the instruction

Registers Abbreviated

- rs = Register Source
- rd = Register Destination
- imm = Immediate
 - Also known as c _ _ _ _ _
 - Example?

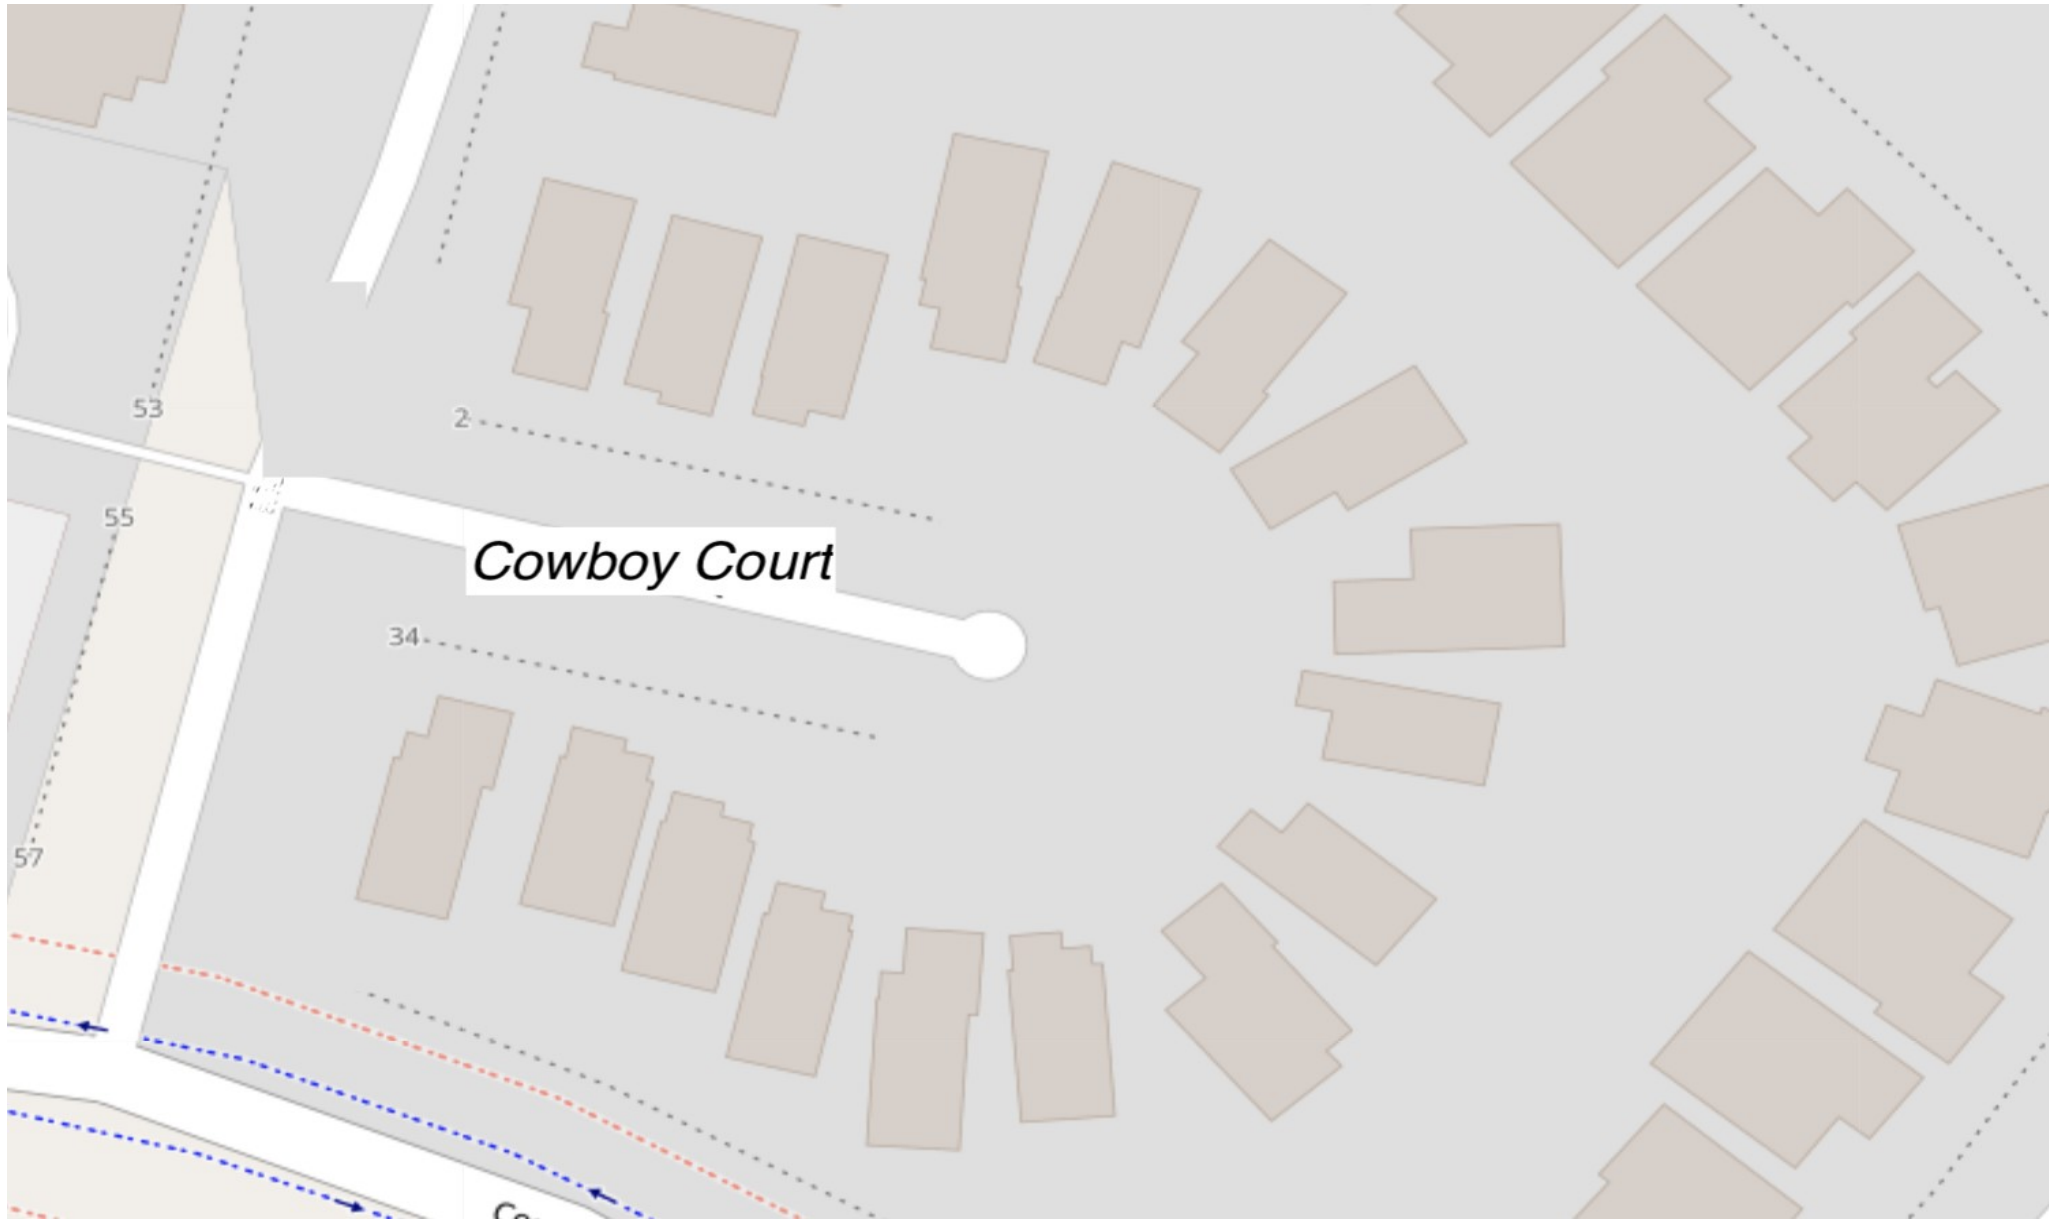
RISC-V Registers

- **Question:** Can the instruction list/set be used to determine which register is which?
 - Asked by Liam
- **Answer:** Nope, you should use the *GreenCard*, to determine which register is which. Then, use the *InstructionList* to convert RISC-V assembly to machine code
 - *Demo*

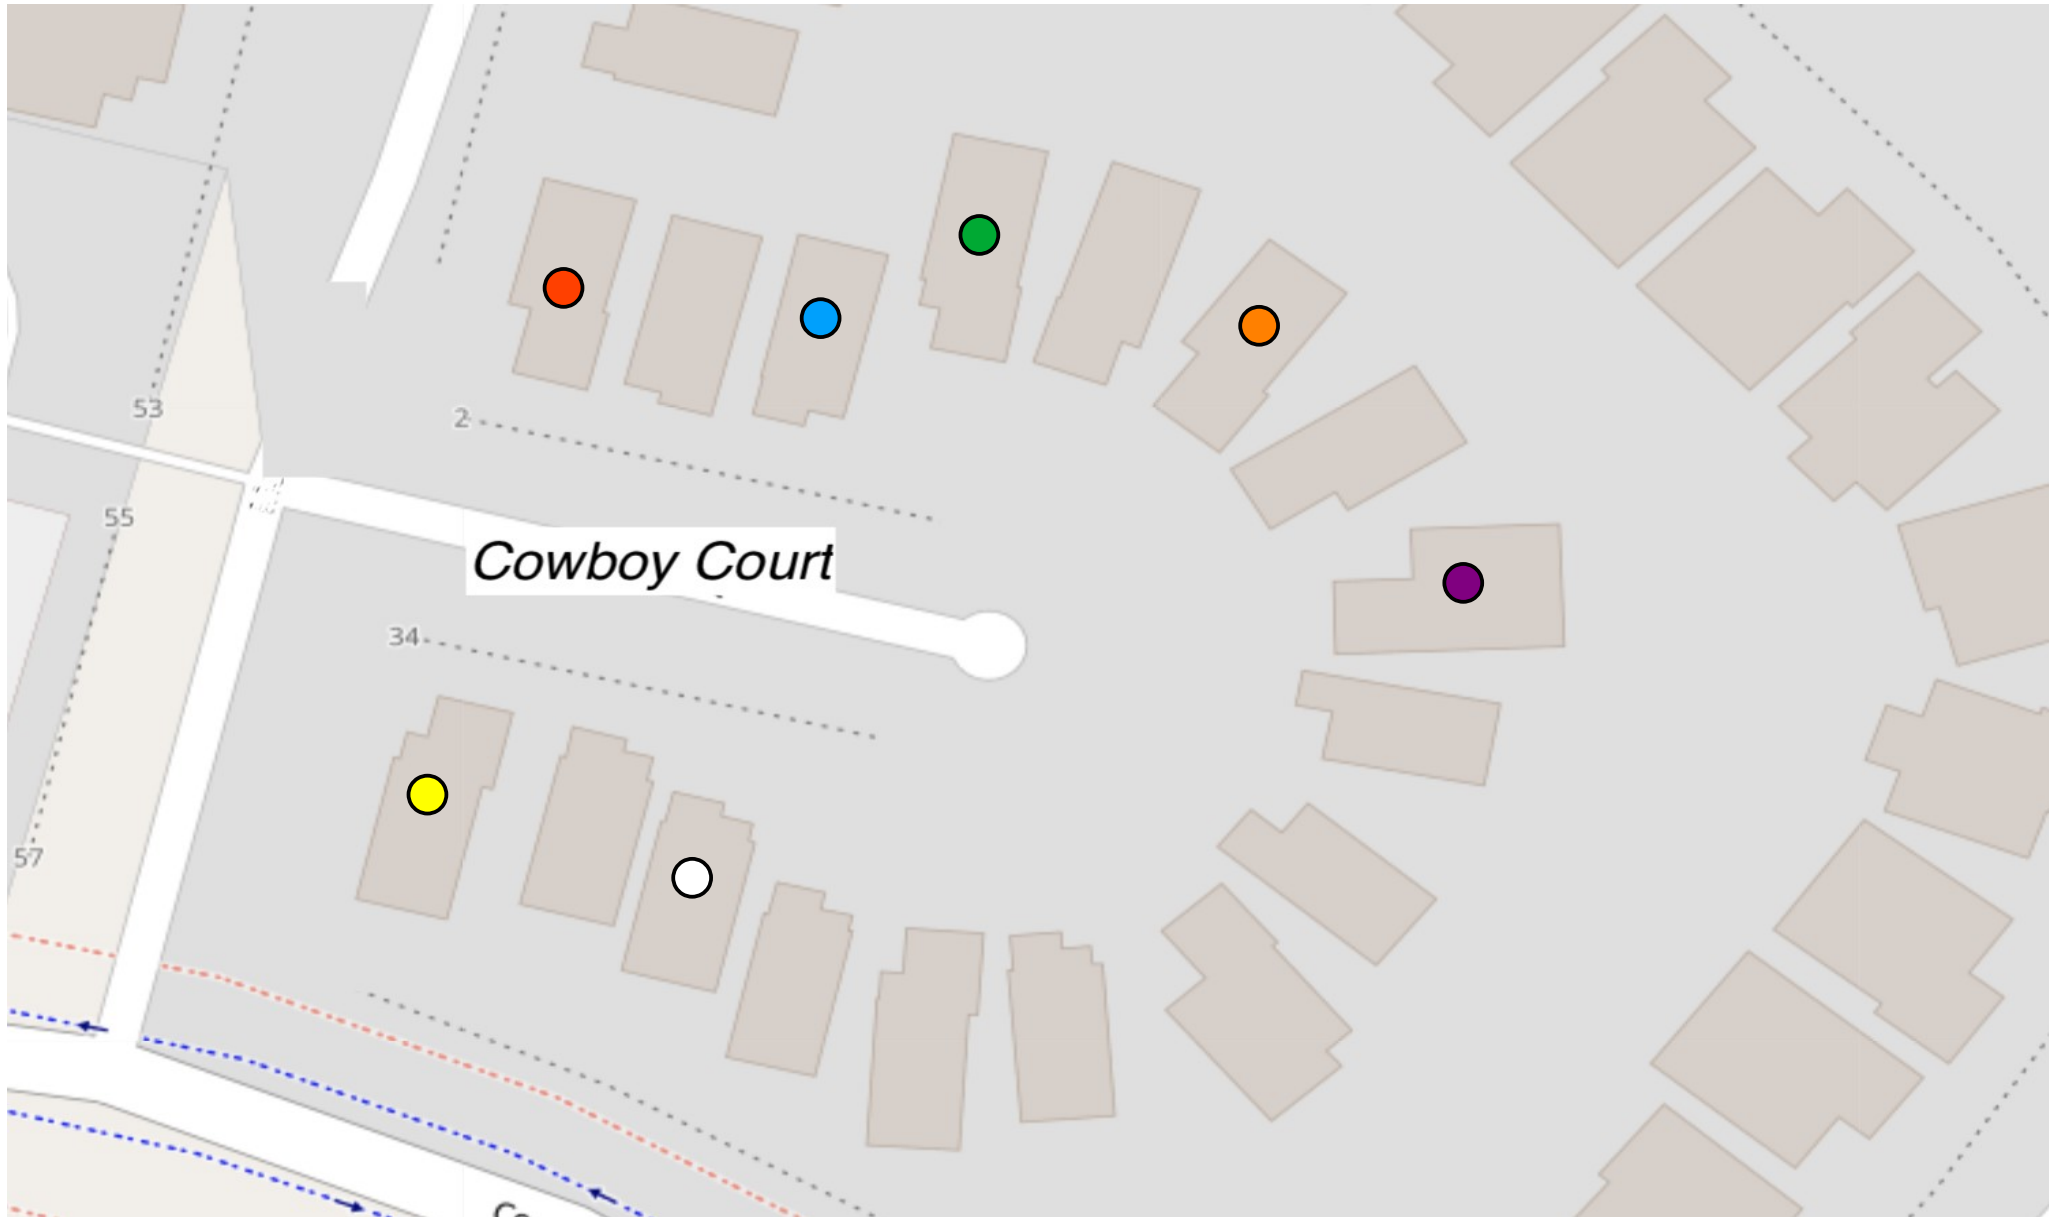
Base Address

- **Question:** *something something... Base address... something something?*
 - Asked by Liam
- **Answer:** Think about it like this
 - *<name>* takes my lunch money, and the only thing I know is that *<name>* lives at Cowboy Court, 5th house down, from where the Court starts
 - Where does *<name>* live?
 - *Next slide*

Cowboy Court



Cowboy Court (Labelled)



Abstraction → Application

- How does it work in practice?
 - i.e. *int.arrays.c*
 - i.e. *diagram.int.array.txt*
 - i.e. *double.arrays.c*
 - i.e. *diagram.double.array.txt*

Immediate(s)

- **Question:** Why is immediate split up?
 - Can't remember who asked this
- **Answer:** To keep the format/structure of the instruction the same
 - Makes decoding easier
 - Consistency is good
 - Size and place/location should be consistent

31	27	26	25	24	20	19	15	14	12	11	7	6	0	
funct7				rs2		rs1		funct3		rd		opcode		R-type
imm[11:0]						rs1		funct3		rd		opcode		I-type
imm[11:5]				rs2		rs1		funct3		imm[4:0]		opcode		S-type
imm[12 10:5]				rs2		rs1		funct3		imm[4:1 11]		opcode		B-type
imm[31:12]										rd		opcode		U-type
imm[20 10:1 11 19:12]										rd		opcode		J-type

Instruction Type

- **Question:** What type of instruction is **sw**?
 - What does **sw** stand for?
 - Is **sw** used for *32-bit* or *64-bit* architecture?
- **Options:**
 - A) R-type
 - B) I-type
 - C) B-type
 - D) U-type
 - E) S-type
 - F) J-type

Question #1 (2.27)

- Translate the following RISC-V assembly code into C. Assume that the variable *i* is stored in register *x6*, **result** is stored in *x5*, and *x10* holds the base address of an array called **MemArray**.
 - *Assembly code on next slide*

Question #1 (2.27)

.data

MemArray:

```
.word 0  
.word 1  
...  
.word 2  
.word 99
```

.text

```
addi x6, x0, 0  
addi x5, x0, 0  
addi x29, x0, 100
```

LOOP:

```
lw x7 0(x10)  
add x5, x5, x7  
addi x10, x10, 4  
addi x6, x6, 1  
blt x6, x29, LOOP
```

- Recall that:
 - **x6** = i
 - **x5** = *result*
 - **X10** = Base address of *MemArray*

Answer #1 (2.27)

- **Solution #1:**

```
int i;  
int result = 0;  
for (i = 0; i < 100; i++) {  
    result += *MemArray;  
    MemArray++;  
}  
return result;
```

- **Solution #2:**

```
int i;  
int result = 0;  
for (i = 0; i < 100; i++) {  
    result += MemArray[i];  
}  
return result;
```


Clarifications #1 (2.27)

- **Question:** Can a *while loop* be used instead of a *for loop*?
 - **Answer:** Yes! Semantically speaking, there is little difference between a *for loop* and a *while loop*.
 - Programmers prefer *for loops* because they are quicker to write, BUT not quicker to execute on the CPU
- **Question:** Why are there 2 solutions?
 - **Answer:** Because C code to assembly isn't a 1:1 mapping. When you compile C code, the output depends on what you've written AND the what the compiler does (i.e. The compiler may make optimizations)

Question #2 (2.31)

- Translate function ***f*** into RISC-V assembly language. Assume the function declaration for ***g*** is ***int g(int a, int b)***.

The code for function ***f*** is as follows:

```
int f(int a, int b, int c, int d) {  
    return g(g(a,b), c+d);  
}
```

- Note: For a question like this, you will have to specify registers yourself. (i.e. I will assume that variable “abc” is stored in register “yyz”).
 - Thus, you will need to reference the RISC-V Calling Convention

Calling Convention

RISC-V Calling Convention

Register	ABI Name	Saver	Description
x0	zero	---	Hard-wired zero
x1	ra	Caller	Return address
x2	sp	Callee	Stack pointer
x3	gp	---	Global pointer
x4	tp	---	Thread pointer
x5-7	t0-2	Caller	Temporaries
x8	s0/fp	Callee	Saved register/frame pointer
x9	s1	Callee	Saved register
x10-11	a0-1	Caller	Function arguments/return values
x12-17	a2-7	Caller	Function arguments
x18-27	s2-11	Callee	Saved registers
x28-31	t3-t6	Caller	Temporaries

Answer #2 (2.31)

f:

```
addi x2, x2, -8    // Allocate stack space for 2 words
sw x1, 0(x2)       // Save return address
add x5, x12, x13   // x5 = c + d
sw x5, 4(x2)       // Save x5, which is c + d, on the stack
jal x1, g          // Call x10 = g(a,b)
lw x11, 4(x2)      // Reload x11 = c + d, from the stack
jal x1, g          // Call x10 = g(g(a,b), (c + d))
lw x1, 0(x2)       // Restore return address
addi x2, x2, 8     // Restore stack pointer
jalr x0, 0(x1)     // Return to caller, f
```

Question #3 (2.37)

- **Question:** Write the RISC-V assembly code to implement the following C code as an atomic “set max” operation using the *lr.d* and *sc.d* instructions. Here, the argument *shvar* contains the address of a shared variable which should be replaced by *y* if *y* is greater than the value it points to. Assume *x10* is address of integer pointed by *shvar* and value of *y* is in *x11*.
- **Code:**
 - *Next slide*

Question #3 (2.37)

- **Code:**

```
void setmax(int* shvar, int y) {  
    // Begin critical section  
    if (y > *shvar)  
        *shvar = y;  
    // End critical section}  
}
```

Answer #3 (2.37)

- **Solution:**

setmax:

try:

```
lr.d x5, (x10)           // Load-reserve *shvar
```

```
bge x5, x11, release // Skip update if *shvar > y
```

```
addi x5, x11, 0          // *shvar = y;
```

release:

```
sc.d x7, x5, (x10)       // Check x5 before storing
```

```
bne x7, x0, try          // If store-conditional failed, try again
```

```
jalr x0, 0(x1)           // Jump to the calling function
```


Question #3 (2.37)

- The best way to understand critical sections is to understand:
 - Mutual exclusion (mutexes)
 - Race conditions
 - Deadlock
 - Etc.
- Why all this extra stuff?
 - Because in computer science, the *how* is more important than the *why*
- Side note: All original Q/A document can be found on Teams
 - Posted by Mingzhe Wang

HAVE AN AMAZING READING WEEK

GOD BLESS YOU ALL

BUT MOSTLY ME

& STAY SAFE

HAVE FUN!