

Threads

Bojan Nokovic

Based on: "Operating Systems Concepts", 10th Edition Silberschatz Et al.

Jan. 2021

What is a Thread?

A sequential execution **stream** (thread) within a process (also called **lightweight** process).

A process has at least **one** thread of control.

A process has two parts: threads (concurrency) and address spaces (protection).

Most modern applications are **multithreaded**.

A word processor may have a thread for


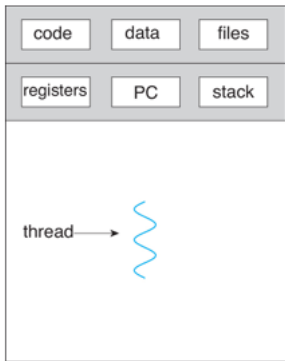
- Displaying graphics
- Responding to keystrokes from the user
- Spell and grammar checking

Multithreading can simplify code, increase efficiency.

Multiprocessing (multiprocessor systems): Split program into multiple threads to make it run faster by running on multiple processors.

Parallel programming.


Single and Multithreaded Processes



Single-threaded process and a multithreaded process.

Multithreaded Server Architecture

Web server accepts client requests for web pages, images, sound, and so forth.



Old approach: new process-creation to serve request.

- Time consuming and resource intensive.

New approach: **create a new thread** to service the request and resumes listening for additional requests.

Thread States

States **shared** by all threads in the same process/address space:

- global variables
- file system
- code

States **private** to each thread

- PC, registers
- execution stack contains parameters, temporary variables, return addresses

Responsiveness

- May **allow continued execution** if part of process is blocked, especially important for user interfaces.
- The time-consuming operation is performed in a separate, **asynchronous thread**, the application remains responsive to the user.

Resource Sharing

- Processes - shared memory and message passing. Must be explicitly arranged by the programmer.
- Threads share resources of process by **default**.

Economy

- Thread switching **lower overhead** than context switching.
- Thread creation consumes **less time and memory** than process creation.

Scalability

- Multithreading can take advantage of multicore architectures, where threads may be **running in parallel** on different processing cores.

Multicore or multiprocessor systems putting pressure on programmers, challenges include:

- Dividing activities
- Balance
- Data splitting
- Data dependency
- Testing and debugging

Parallelism implies a system can perform more than one task simultaneously.

Concurrency supports more than one task making progress

- Single processor/core, scheduler providing concurrency


Types of parallelism

- **Data parallelism** - distributes subsets of the same data across multiple cores, same operation on each
- **Task parallelism** - distributing threads across cores, each thread performing unique operation

Example: Summing an array of size N

- Single core: Sum the elements $[0] \dots [N - 1]$.
- Dual core: Thread A sum elements $[0] \dots [N/2 - 1]$, while thread B sum elements $[N/2] \dots [N - 1]$

Data and Task Parallelism



Amdahl's Law

Identifies **performance gains** from adding additional cores to an application that has both serial and parallel components.

S is serial portion %, N number of processing cores

$$speedup \leq \frac{1}{S + \frac{1-S}{N}}$$

What is *speedup* if application is 75% parallel and 25% serial when moving from 1 to 2 cores?

What happens when $N \rightarrow \infty$?

User Threads and Kernel Threads


User threads - management done by user-level threads library

Three primary thread libraries:

- POSIX **Pthreads**
- Windows threads
- Java threads


Kernel threads - Supported by the Kernel

- Windows, Linux, Mac OS X, iOS, Android




Many User Threads Mapped to One Kernel Thread

- One thread blocking causes all to block ☹️
- Multiple threads may not run in parallel on multicore system because only one may be in kernel at a time ☹️
- Few systems currently use this model: Solaris Green Threads, GNU Portable Threads




One to One

- Each user-level thread maps to kernel thread
- More concurrency than many-to-one 😊
- Number of threads per process sometimes restricted due to overhead
- Windows, Linux




Many-to-Many

- Allows the operating system to create a sufficient number of kernel threads - a smaller or equal number of user threads.
- Windows with the *ThreadFiber* package
- Not very common



Two-level Model

- Similar to Many-to-Many, except that it allows a user thread to be **bound** to kernel thread



Although the many-to-many (M:M) model appears to be the most flexible of the models discussed, in practice it is difficult to implement.

Most operating systems now use the one-to-one model.

Thread library provides programmer with API for **creating** and **managing** threads.

Two primary ways of implementing

- Library entirely in user space
- Kernel-level library supported by the OS

May be provided either as user-level or kernel-level

A POSIX [specification](#) standard (IEEE 1003.1c) API for thread creation and synchronization

`https:`

`//standards.ieee.org/standard/1003_1-2017.html`

API specifies [behavior](#) of the thread library, [implementation](#) is up to development of the library

Common in UNIX, Linux & Mac OS X

Pthreads Example sum.c

```
#include <pthread.h>
#include <stdio.h>
#include <stdlib.h>

int sum; /* this data is shared by the thread(s) */
void *runner(void *param); /* the thread */

int main(int argc, char *argv[])
{
    pthread_t tid; /* the thread identifier */
    pthread_attr_t attr; /* set of attributes for the thread */

    pthread_attr_init(&attr); /* get the default attributes */
    /* create the thread */
    pthread_create(&tid, &attr, runner, argv[1]);
    pthread_join(tid, NULL); /* now wait for the thread to exit */

    printf("sum = %d\n", sum);
}
```

Pthreads Example Continue

```
/**
 * The thread will begin control in this function
 */
void *runner(void *param)
{
    int i, upper = atoi(param);
    sum = 0;

    if (upper > 0) {
        for (i = 1; i <= upper; i++)
            sum += i;
    }

    pthread_exit(0);
}
```

On Linux compile with:

```
gcc -g -Wall -pthread sum.c -lpthread -o sum
```

Thread Programming Challenges

Write programs with multiple simultaneous points of execution, **synchronizing** through shared memory.

Global variables are **shared** among all the threads of the same process.

Threads can read and **write** the same memory locations.

The programmer is **responsible** for using the synchronization mechanisms of the thread facility to ensure that the shared memory is accessed in a manner that will give the correct answer!

Growing in popularity as numbers of threads increase, program correctness more difficult with explicit threads.

Creation and management of threads done by **compilers and run-time libraries** rather than programmers.

Five methods explored:

- 1 Thread Pools
- 2 Fork-Join
- 3 OpenMP
- 4 Grand Central Dispatch
- 5 Intel Threading Building Blocks


Create a number of threads in a pool where they await work.

Advantages

- Usually **slightly faster** to service a request with an existing thread than create a new thread
- Allows the number of threads in the application(s) to be bound to the size of the pool
- Separating task to be performed from mechanics of creating task allows different strategies for running task
 - **Tasks could be scheduled to run periodically**

Fork-Join Parallelism


Multiple threads (tasks) are forked, and then joined.



```
Task(problem)
  if problem is small enough
    solve the problem directly
  else
    s1 = fork(new Task(subset of problem))
    s2 = fork(new Task(subset of problem))

    result1 = join(s1)
    result2 = join(s2)

    return combined results
}
```



Open Multi-Processing (OpenMP)

- Set of compiler directives and an API for C, C++, FORTRAN
- Support for **parallel programming in shared-memory environments**
- Identifies parallel regions - blocks of code that can run in parallel

```
#include <omp.h>
#include <stdio.h>
int main(int argc, char *argv[]) {
    /* sequential code */

    #pragma omp parallel {
        printf("I am a parallel region.");
    }

    /* sequential code */

    return 0;
}
```

<https://www.openmp.org/>

Apple technology for macOS and iOS operating systems

Extensions to C, C++ and Objective-C languages, API, and run-time library

```
^ { printf("I am a block"); }
```

Blocks placed in dispatch queue - assigned to available thread in thread pool when removed from queue

Intel Threading Building Blocks (TBB)

Template library for designing parallel C++ programs

```
for (int i = 0; i < n; i++) {  
    apply(v[i]);  
}
```

```
parallel_for (size_t(0), n,  
    [=](size_t i) {apply(v[i]);}  
);
```

A serial version, and TBB version of a simple `for` loop.

Issues in Designing Multithreaded Programs

- 1 Semantics of `fork()` and `exec()` system calls
- 2 Signal handling - synchronous and asynchronous
- 3 Thread cancellation of target thread - asynchronous or deferred?
- 4 Thread-local storage
- 5 Scheduler Activations

Semantics of `fork()` and `exec()`

The semantics of the `fork()` and `exec()` system calls **change** in a multithreaded program.

If one thread in a program calls `fork()`, does the new process duplicate **all threads**, or is the new process **single-threaded**?

If a thread invokes the `exec()` system call, the program specified in the parameter to `exec()` will replace the entire process-including all threads.

Which of the two versions of `fork()` to use depends on the application.

A signal is used in UNIX systems to notify a process that a particular event has occurred.

A **signal handler** is used to process signals

- 1 A signal is generated by the occurrence of a particular event.
- 2 The signal is delivered to a process.
- 3 Once delivered, the signal must be handled either by **default** or **user-defined** signal handler.

Every signal has default handler that kernel runs when handling signal

- User-defined signal handler can override default

Synchronous signal example

- Illegal memory access
- Division by 0

Delivered to the same process that performed the operation that caused the signal (that is the reason they are considered synchronous).

When a signal is generated by an event external to a running process, that process receives the signal **asynchronously**.

- Terminating a process with specific keystrokes (such as `<control><C>`)
- Having a timer expire

For single-threaded, signal is delivered to process.

Where should a signal be delivered for multi-threaded?

- Deliver the signal to the thread to which the signal applies
- Deliver the signal to every thread in the process
- Deliver the signal to certain threads in the process
- Assign a specific thread to receive all signals for the process

The method for delivering a signal depends on the type of signal generated - i.e. (`<control><C>`, for example) should be sent to all threads.

Thread Cancellation

Terminating a thread before it has finished.

Thread to be canceled is **target thread**

- **Asynchronous cancellation** terminates the target thread immediately
- **Deferred cancellation allows** the target thread to periodically check if it should be cancelled (default type)

```
pthread_t tid;  
/* create the thread */  
pthread_create(&tid, 0, worker, NULL);  
...  
/* cancel the thread */  
pthread_cancel(tid);  
  
/* wait for the thread to terminate */  
pthread_join(tid, NULL);
```

Thread Cancellation

Invoking `pthread_cancel()` indicates only **a request** to cancel the target thread. However actual cancellation depends on how the target thread is set up to handle the request.

When the target thread is finally canceled, the call to `pthread_join()` in the cancelling thread returns.

A thread may set its cancellation state and type using an API.

Mode	State	Type
Off	Disabled	–
Deferred	Enabled	Deferred
Asynchronous	Enabled	Asynchronous

The cancellation only occurs when thread reaches cancellation point i.e. `pthread_testcancel()` then **cleanup handler** is invoked.

Thread-Local Storage (TLS)

Allows each thread to have its own copy of data.

Useful when you do not have control over the thread creation process (i.e., when using a thread pool)

Different from local variables

- Local variables visible only during single function invocation
- TLS visible across function invocations


Similar to `static` data - unique to each thread

Scheduler Activations

Both M:M and Two-level models **require communication to maintain the appropriate number of kernel threads** allocated to the application.

Typically use an intermediate data structure between user and kernel threads - lightweight process (LWP)

- Appears to be a virtual processor on which process can schedule user thread to run
- Each LWP attached to kernel thread - how many to create?



Scheduler activations provide **upcalls** - a communication mechanism from the kernel to the **upcall handler** in the thread library.

Thank you !

Operating Systems are among the most complex pieces of software ever developed !