

# Deadlocks

Bojan Nokovic

Based on: "Operating Systems Concepts", 10th Edition Silberschatz Et al.  
"Slides 3SH3 '12" - Sanzheng Qiao

Feb. 2021

# System Model

System consists of a **finite** number of resources of different types  $R_1, R_2, \dots, R_m$ .

Resource types: CPU cycles, memory space, I/O devices.

Each resource type  $R_i$  has  $W_i$  identical instances - if system has 2 CPUs, then resource type CPU has two instances.

Thread requests an instance of a resource type, the allocation of **any** instance of the type should satisfy the request.

Each process utilizes a resource as follows:

- 1 Request the resource
- 2 Use the resource
- 3 Release the resource

# Kernel-managed Resource

OS checks to make sure that the thread has requested and has been allocated the resource.

**A system table** - for each resource that is allocated, the table records the thread to which it is allocated.

Queue of waiting threads.

A set of threads is in a **deadlocked** state when every thread in the set is waiting for an event that can be caused only by another thread in the set.

**Livelock** occurs when a thread continuously attempts an action that fails - less common than deadlock.

# Example: Deadlock in Multithreaded Application

```
/* two mutex locks are created and initialized */
pthread_mutex_t first_mutex;
pthread_mutex_t second_mutex;

pthread_mutex_init(&first_mutex, NULL);
pthread_mutex_init(&second_mutex, NULL);


/* thread one runs in this function */
void *do_work_one(void *param) {
    pthread_mutex_lock(&first_mutex);
    pthread_mutex_lock(&second_mutex);
    /* Do some work */
    pthread_mutex_unlock(&second_mutex);
    pthread_mutex_unlock(&first_mutex);
    pthread_exit(0);
}

/* thread two runs in this function */
void *do_work_two(void *param) {
    pthread_mutex_lock(&second_mutex);
    pthread_mutex_lock(&first_mutex);
    /* Do some work */
    pthread_mutex_unlock(&first_mutex);
    pthread_mutex_unlock(&second_mutex);
    pthread_exit(0);
}
```

# A Resource Allocation Graph

Deadlock is possible if thread 1 acquires `first_mutex` and thread 2 acquires `second_mutex`.

Thread 1 then waits for `second_mutex` and thread 2 waits for `first_mutex`.



Difficult to identify and test deadlocks that may occur only under certain scheduling circumstances.

# Deadlock Characterization

Deadlock **can** arise if four conditions hold simultaneously.

- **Mutual exclusion**: only one process at a time can use a resource.
- **Hold and wait**: a process holding at least one resource is waiting to acquire additional resources held by other processes.
- **No preemption**: a resource can be released only **voluntarily** by the process holding it, after that process has completed its task.
- **Circular wait**: there exists a set  $\{P_0, P_1, \dots, P_n\}$  of waiting processes such that  $P_0$  is waiting for a resource that is held by  $P_1$ ,  $P_1$  is waiting for a resource that is held by  $P_2, \dots, P_{n-1}$  is waiting for a resource that is held by  $P_n$ , and  $P_n$  is waiting for a resource that is held by  $P_0$ .

# Resource-Allocation Graph

Deadlocks described more precisely by a **system resource-allocation graph**.

A set of vertices  $V$  and a set of edges  $E$ .

$V$  is partitioned into two types:

- $P = \{P_1, P_2, \dots, P_n\}$ , the set consisting of all the processes in the system
- $R = \{R_1, R_2, \dots, R_m\}$ , the set consisting of all resource types in the system


**request edge** - directed edge  $P_i \rightarrow R_j$ ,  $i \in \{1, \dots, n\}$ ,  $j \in \{1, \dots, m\}$

**assignment edge** - directed edge  $R_j \rightarrow P_i$

# Resource Allocation Graph Example

We represent thread  $T_i$  as a circle resource type  $R_i$  as a rectangle.


- One instance of  $R_1$ , two of  $R_2$ , one of  $R_3$ , three of  $R_4$
- $T_1$  holds one instance of  $R_2$  and is waiting for an instance of  $R_1$
- $T_2$  ?
- $T_3$  ?



$$E = \{T_1 \rightarrow R_1, T_2 \rightarrow R_3, R_1 \rightarrow T_2, \dots\}$$



# Graph With a Deadlock




No cycles  $\Rightarrow$  no deadlock.

How many cycles exist in the system?

$T_1, T_2, T_3$  are deadlocked.

# Graph With a Cycle But No Deadlock



Cycle:  $T_1 \rightarrow R_1 \rightarrow T_3 \rightarrow R_2 \rightarrow T_1$

No deadlock! Why?

If there is a cycle, then the system **may** or **may not** be in a deadlocked state.

# Methods for Handling Deadlocks

Ensure that the system will **never** enter a deadlock state:

- Deadlock prevention
- Deadlock avoidance

**Deadlock-prevention** algorithms prevent deadlocks by limiting how requests can be made.

- ▷ low device utilization
- ▷ reduced system throughput

**Avoiding deadlocks** is to require additional information about how resources are to be requested.

## Algorithm

What is an algorithm?

# Methods for Handling Deadlocks

Ensure that the system will **never** enter a deadlock state:

- Deadlock prevention
- Deadlock avoidance

**Deadlock-prevention** algorithms prevent deadlocks by limiting how requests can be made.

- ▷ low device utilization
- ▷ reduced system throughput

**Avoiding deadlocks** is to require additional information about how resources are to be requested.

## Algorithm

What is an algorithm?

A predetermined set of instructions for solving a specific problem in a limited number of steps.

**Mutual Exclusion** - **must hold** for non-sharable resources.

- ▷ not required for sharable resources (e.g., read-only files)

**Hold and Wait** - must guarantee that whenever a process requests a resource, it does not hold any other resources.

- Require process to request and be allocated all its resources before it begins execution.
- ▷ low resource utilization - resources allocated but not used
- ▷ starvation possible - waiting to resource forever

**No Preemption** - if a process that is holding some resources requests another resource that cannot be immediately allocated to it, then **all resources currently being held are released**.

- Preempted resources are added to the list of resources for which the process is waiting

**Circular Wait** - impose a **total ordering** of all resource types, and require that each process requests resources in an increasing order of enumeration.

Invalidating the circular wait condition is most common.

Simply assign each resource  $R = \{R_1, R_2, \dots, R_n\}$  a unique number,  $F : R \rightarrow \mathbb{N}$ , injective (one-to-one).

Resources must be acquired in **order**.

If the lock ordering in the Pthread program shown in page 4

`F(first_mutex) = 1`

`F(second_mutex) = 5`

A thread that wants to use both `first_mutex` and `second_mutex` at the same time **must first request** `first_mutex` and then `second_mutex`.

# Deadlock Avoidance

Requires that the system has some additional **a priori** information available.

► In a system with resources  $R_1$  and  $R_2$ , the system might need to know that thread **P** will request first  $R_1$  and then  $R_2$  before releasing both resources, whereas thread **Q** will request  $R_2$  and then  $R_1$ .

- The deadlock-avoidance algorithm **dynamically examines** the resource-allocation state to ensure that there can never be a circular-wait condition.
- Resource-allocation **state** is defined by the number of available and allocated resources, and the maximum demands of the processes.



A state is **safe** if the system can allocate resources to each thread and still avoid a deadlock.

System is in safe state if there exists a sequence  $\langle P_1, P_2, \dots, P_n \rangle$  of all the processes in the systems such that for each  $P_i$ , the resources that  $P_i$  can still request can be satisfied by:


currently available res. + res. held by all the  $P_j$ , with  $j < i$

# Basic Facts

If a system is in safe state  $\Rightarrow$  no deadlock

If a system is in unsafe state  $\Rightarrow$  possibility of deadlock

Avoidance  $\Rightarrow$  ensure that a system will never enter an unsafe state.



# Example

System with **twelve** resources and three threads.

▷  $T_0$  - requires ten resources,  $T_1$  - four resources,  $T_2$  - nine resources

Suppose that, at time  $t_0$ , thread  $T_0$  is holding five resources,  $T_1$  two resources,  $T_2$  two resources.

	<u>Max. Needs</u>	<u>Current Need</u>
$T_0$	10	5
$T_1$	4	2
$T_2$	9	7

The sequence  $\langle T_1, T_0, T_2 \rangle$  satisfies the safety condition - the system in a **safe state**.

## Exercise

Suppose that, at time  $t_1$ , thread  $T_2$  requests and is allocated one more resource. Is the system in a safe state?

# Avoidance Algorithms

If we have a resource-allocation system with only one instance of each resource type we use

- Resource-allocation graph

For multiple instances of a resource type

- Banker's Algorithm

# Resource-Allocation Graph Scheme

**Claim edge**  $P_i \rightarrow R_j$  indicated that process  $P_j$  **may** request resource  $R_j$



**Claim edge** converts to **request edge** when a process requests a resource.

**Request edge** converted to an **assignment edge** when the resource is allocated to the process.

Resources must be claimed **a priori** in the system - before thread  $T_i$  starts executing, all its claim edges must already appear in the resource-allocation graph.

# Resource-Allocation Graph

Claim edge is represented by a dashed line.



Unsafe State

- ▶  $T_2$  requests  $R_2$
- ▶ Although  $R_2$  is currently free, we cannot allocate it to  $T_2$ , since this action will create a cycle in the graph

# Banker's Algorithm

Multiple instances of resources.

Each process must a priori claim maximum use.

When a process requests a resource it may have to wait.

When a process gets all its resources it must return them in a finite amount of time.

# Data Structures for the Banker's Algorithm

- ▶  $n$  = number of processes
- ▶  $m$  = number of resources types
- **Available:** Vector of length  $m$ . If  $available[j] = k$ , there are  $k$  instances of resource type  $R_j$  available
- **Max:**  $n \times m$  matrix. If  $Max[i, j] = k$ , then process  $P_i$  may request at most  $k$  instances of resource type  $R_j$
- **Allocation:**  $n \times m$  matrix. If  $Allocation[i, j] = k$  then  $P_i$  is currently allocated  $k$  instances of  $R_j$
- **Need:**  $n \times m$  matrix. If  $Need[i, j] = k$ , then  $P_i$  may need  $k$  more instances of  $R_j$  to complete its task

$$Need[i, j] = Max[i, j] - Allocation[i, j]$$



# Safety Algorithm

Let  $X, Y$  be vectors of length  $n$ .

$$X \leq Y \Leftrightarrow X[i] \leq Y[i], \quad i = 1, 2, \dots, n$$

We treat row in the matrices  $Allocation_i$  and  $Need_i$  as vectors.

- 1 Let  $Work$  and  $Finish$  be vectors of length  $m$  and  $n$ , respectively. Initialize:  
     $Work = Available$   
     $Finish[i] = \text{false}$  for  $i = 0, 1, \dots, n - 1$
- 2 Find an  $i$  such that:  
     $Finish[i] == \text{false} \wedge Need_i \leq Work$   
if no such  $i$  exists, go to step 4
- 3      $Work = Work + Allocation_i$   
     $Finish[i] = \text{true}$   
go to step 2
- 4 If  $Finish[i] == \text{true}$  for all  $i$ , the system is in a **safe state**

# Resource-Request Algorithm for Process $P_i$

Determine whether request can be safely granted.

$Request_i[j] = k$ ,  $P_i$  wants  $k$  instances of  $R_j$ .

- 1 If  $Request_i \leq Need_i$  go to step 2. Otherwise, raise error condition, since process has **exceeded its maximum claim**.
- 2 If  $Request_i \leq Available$ , go to step 3. Otherwise  $P_i$  must wait, since **resources are not available**.
- 3 Pretend to allocate requested resources to  $P_i$  by modifying the state as follows:  
 $Available = Available - Request_i$ ;  
 $Allocation_i = Allocation_i + Request_i$ ;  
 $Need_i = Need_i - Request_i$ ;
  - If **safe**  $\Rightarrow$  the resources are allocated to  $P_i$
  - If **unsafe**  $\Rightarrow P_i$  must wait, and the old resource-allocation state is restored

# Example of Banker's Algorithm

5 processes  $P_0$  through  $P_4$ ;

3 resource types: A (10 instances), B (5), and C (7)

Snapshot at time  $t_0$ :

	<u>Allocation</u>	<u>Max</u>	<u>Available</u>
	ABC	ABC	ABC
$P_0$	0 1 0	7 5 3	3 3 2
$P_1$	2 0 0	3 2 2	
$P_2$	3 0 2	9 0 2	
$P_3$	2 1 1	2 2 2	
$P_4$	0 0 2	4 3 3	

# Example of Banker's Algorithm

▷  $Need = Max - Allocation$

	<u>Need</u>
	ABC
$P_0$	7 4 3
$P_1$	1 2 2
$P_2$	6 0 0
$P_3$	0 1 1
$P_4$	4 3 1

The system is in a safe state since the sequence  $\langle P_1, P_3, P_4, P_2, P_0 \rangle$  satisfies safety criteria.

A state is safe if the system can allocate resources to each thread (up to its maximum) in some order and still avoid a deadlock.

## Example: $P_1$ Request (1,0,2)

$Request_1 = (1, 0, 2)$

Check  $Request_1 \leq Available$  that is,  $(1,0,2) \leq (3,3,2) \Rightarrow \text{true}$

	<u>Allocation</u>	<u>Need</u>	<u>Available</u>
	ABC	ABC	ABC
$P_0$	0 1 0	7 4 3	2 3 0
$P_1$	3 0 2	0 2 0	
$P_2$	3 0 2	6 0 0	
$P_3$	2 1 1	0 1 1	
$P_4$	0 0 2	4 3 1	

Executing safety algorithm shows that sequence  $\langle P_1, P_3, P_4, P_0, P_2 \rangle$  satisfies safety requirement

Can request for (3,3,0) by  $P_4$  be granted?

If a system does not employ either a deadlock-prevention or a deadlock avoidance algorithm, then a deadlock situation may occur.

- **Detection algorithm** - determine whether a deadlock has occurred.
- **Recovery scheme** - an algorithm to recover from the deadlock

# Single Instance of Each Resource Type


Maintain **wait-for** graph

- Nodes are processes
- $P_i \rightarrow P_j$  if  $P_i$  is waiting for  $P_j$


Periodically invoke an algorithm that **searches for a cycle in the graph**. If there is a cycle, there exists a deadlock.

An algorithm to detect a cycle in a graph requires an order of  $O(n^2)$  operations, where  $n$  is the number of vertices in the graph.

# Resource-Allocation Graph and Wait-for Graph



Resource-allocation graph



Corresponding wait-for graph

We obtain **wait-for** graph from the resource-allocation graph by removing the resource nodes and collapsing the appropriate edges.



# Several Instances of a Resource Type

The wait-for graph scheme is not applicable to a resource-allocation system with **multiple instances** of each resource type.

Data structures:

- **Available:** A vector of length  $m$  indicates the **number of available resources** of each type.
- **Allocation:** An  $n \times m$  matrix defines the number of resources of each type currently allocated to each process.
- **Request:** An  $n \times m$  matrix indicates the current request of each process. If **Request**  $[i][j] = k$ , then process  $P_i$  is requesting  $k$  more instances of resource type  $R_j$ .

# Detection Algorithm

- ① Let *Work* and *Finish* be vectors of length  $m$  and  $n$ , respectively.
  - (a)  $Work = Available$
  - (b) for  $i = 1, 2, \dots, n$  if  $Allocation_i \neq 0$ , then  
 $Finish[i] = \text{false}$  otherwise  $Finish[i] = \text{true}$
- ② Find an index  $i$  such that both
  - (a)  $Finish[i] == \text{false}$
  - (b)  $Request_i \leq Work$

If no such  $i$  exists, go to step 4 (next page)

# Detection Algorithm (Cont.)

- ③  $Work = Work + Allocation_i$   
 $Finish[i] = \text{true}$   
go to step 2
- ④ If  $Finish[i] == \text{false}$ , for some  $i, i \leq n$ , then the system is in deadlock state.  
Moreover, if  $Finish[i] == \text{false}$ , then  $P_i$  is deadlocked.

Algorithm requires an order of  $O(m \times n^2)$  operations to detect whether the system is in deadlocked state

# Example of Detection Algorithm

5 processes:  $P_0$  through  $P_4$ ;

3 resource types: A(10 instances), B(2), and C(7)

Snapshot at time  $T_0$ :

	<u>Allocation</u>	<u>Request</u>	<u>Available</u>
	ABC	ABC	ABC
$P_0$	0 1 0	0 0 0	0 0 0
$P_1$	2 0 0	2 0 2	
$P_2$	3 0 3	0 0 0	
$P_3$	2 1 1	1 0 0	
$P_4$	0 0 2	0 0 2	

Sequence  $\langle P_0, P_2, P_3, P_1, P_4 \rangle$  will result in  $Finish[i] = \text{true}$  for all  $i$ .

The system is not in a deadlocked state.

## Example (Cont.)

$P_2$  requests an additional instance of type C

	<u>Request</u>
	ABC
$P_0$	0 0 0
$P_1$	2 0 2
$P_2$	0 0 1
$P_3$	1 0 0
$P_4$	0 0 2

System can reclaim resources held by process  $P_0$ , but insufficient resources to fulfill other processes.

**Deadlock exists**, consisting of processes  $P_1$ ,  $P_2$ ,  $P_3$ , and  $P_4$

# Detection-Algorithm Usage

When, and how often, to invoke depends on:

- **How often** a deadlock is likely to occur?
- **How many** processes will be affected by deadlock when it happens?

We can invoke the deadlock detection algorithm **every time** a request for allocation cannot be granted immediately.

▷ Considerable overhead in computation time, but we can identify process that "caused" the deadlock

Invoked **arbitrarily** (i.e. once per hour)

▷ There may be many cycles in the resource graph and so we would not be able to tell which of the many deadlocked processes "caused" the deadlock.

# Recovery from Deadlock: Process Termination

Abort all deadlocked processes or abort one process at a time until the deadlock cycle is eliminated?

In which order should we choose to abort?

- 1 Priority of the process
- 2 How long process has computed, and how much longer to completion
- 3 Resources the process has used
- 4 Resources process needs to complete
- 5 How many processes will need to be terminated
- 6 Is process interactive or batch

# Recovery from Deadlock: Resource Preemption

To eliminate deadlocks using **resource preemption**, we successively preempt some resources from processes and give these resources to other processes until the deadlock cycle is broken.

**Selecting a victim** - which resources and which processes are to be preempted to minimize cost.

**Rollback** - return to some safe state, restart process for that state.

**Starvation** - same process may always be picked as victim, include number of rollback in cost factor.



# Thank you !

Operating Systems are among the  
most complex pieces of software ever developed !