

Lecture.2&3.OS.Structures.txt

- Operating System Services
 - Provide functions that are helpful to the user. For example:
 - User interface
 - i.e. CLI, GUI, Touch-screen, Batch, Voice-controlled, etc.
 - Program execution
 - Load program into memory, run program, end execution normally or abnormally
 - I/O operation
 - May involve a file or an I/O device
 - File system (manipulation)
 - Read, write, create, delete, search files/directories
 - List information and manage permission for a file or directory
 - Communications
 - Via shared memory or through message passing
 - Packets moved by the OS
 - Error detection
 - Identify errors on CPU, memory, hardware, I/O devices, user programs, etc.
 - Identify each type of error and be able to take the appropriate action
 - The action is done by the operating system
 - Resource allocation
 - Resources must be allocated to each user or job that is running concurrently
 - i.e. In DOS, only one application was allowed to run at a time, but now, modern OS's can run multiple programs at once
 - Accounting/Logging
 - Keep track of how much computer resource each user uses
 - A record of how much CPU time and storage each user uses can be used to bill them in the future
 - i.e. Pay per use data, Dropbox, etc.
 - Protection & Security
 - Ensures all access to system resources is controlled, and requires user authentication
 - i.e. Permission to use webcam, mic, etc.
 - i.e. Permission to read or write a file/directory
- A View Of Operating System Services
 - An accounting service is a record that keeps information about the resources used by a user
 - This can be used for billing purposes
 - The kernel is part of the operating system
 - Drivers are also part of the operating system
 - Although some drivers may be part of the kernel
 - System calls provide an interface to the services made available

- by an operating system
- Operating Systems are written in a mix of Assembly and C
 - Time critical structures are written in Assembly
- System Calls
 - The programming interface to the services provided by the OS
 - System calls are mostly accessed by programs via a high-level Application Programming Interface (API) rather than direct system call use
 - The three most common APIs are:
 - Win32 API for Windows
 - POSIX API for POSIX-based systems
 - Includes all versions of UNIX, Linux, and Mac OS X
 - Java API for the Java virtual machine (JVM)
 - i.e.
 - The Unix command: ``cp in.txt out.txt``
 - On Linux and other Unix-like operating systems, ``man`` is the interface used to view the system's reference manuals
- System Call Implementation
 - Typically, a number is associated with each system call
 - The system call interface maintains a table indexed according to these numbers
 - It invokes the intended system call in the OS kernel and returns the status of the system call and any return values
 - The caller does not need to know anything about how the system call is implemented
 - All it needs to do is obey the API and understand what the OS will do as a result of making the system call
 - The details of the OS interface are hidden from programmers by an API
 - This is abstraction; a core principle of (software) engineering
 - Abstraction is purposeful suppression, or hiding, of some details of a process or artifact, in order to bring out more clearly other aspects, details, or structure
 - The API is managed by run-time support libraries
 - This is a set of functions built into libraries and included with the compiler
- System Call Parameter Passing
 - There are three general methods to pass parameters to the OS
 1. Register
 - Passes the parameters in the registers
 - Usually the fastest way to pass parameters
 - Number of parameters that can be passed is limited to number of registers
 2. Block
 - Parameters are stored in a block of memory, and the

address of the block is passed as a parameter in a register

3. Stack

- Parameters are pushed onto the stack by a program, and then popped off the stack by the operating system
- Note: Block and Stack methods do not limit the number or length of parameters being passed

- Types Of System Calls (0)

1. Process control
2. File management
3. Device management
4. Information maintenance
5. Communications
6. Protection

- Types Of System Calls (1)

- Process Control
 - Create process, terminate process
 - End, abort, load, execute
 - Get process attributes, set process attributes
 - Wait for time, wait event, signal event
 - Allocate and free memory, dump memory if error
 - Debugger for determining bugs, single step execution
 - Locks for managing access to shared data between processes
- => Example: Unix system calls: fork(), exit(), wait()

- Types Of System Calls (2, 3)

- File Management
 - Create, delete, open, close file
 - Read, write, reposition
 - Get and set file attributes
- => Example: Unix system calls: open(), read(), write(), close()
- Device Management
 - Request device, release device
 - Read, write, reposition
 - Get device attributes, set device attributes, logically attach or detach devices
- => Example: Unix system calls: ioctl(), read(), write()

- Types Of System Calls (4, 5)

- Information Maintenance
 - Get time or date, set time or date, get system data, set system data
 - Get and set process, file, or device attributes
- => Example: Unix system calls: getpid(), alarm(), sleep()
- Communications
 - Create, delete communication connection, transfer status information, attach and detach remote devices

- Send, receive messages if message passing model; create and gain access to memory regions in shared-memory model
 - => Example: Unix system calls: pipe(), shm_open(), mmap()
- Types Of System Calls (6)
 - Protection
 - Control access to resources
 - Get and set permissions
 - Allow and deny user access
 - => Example: Unix system calls: chmod(), umask(), chown()
- System Call Interface
 - The standard C library provides a portion of the system-call interface
 - i.e. A C program invoking the 'printf()' function, located in the "stdio" library, which invokes write(), a system call
 - i.e.


```
#include <stdio.h>
int main () {
    // code ...
    printf("Hello World\n");
    // code ...
    return 0;
}
/*
 * The 'printf()' function in 'stdio' calls the standard C
 * library and tells it to print. The standard C library
 * invokes the 'write()' system call and performs the
 * task. Then it returns to the standard C library, which
 * returns back to the 'printf()' function at the point it
 * was called in the C program.
 *
 * The 'printf()' function is in user mode.
 * The 'write()' function is in kernel mode.
 */
```
- Example: Arduino
 - Single tasking, single memory space
 - No operating system, boot loader loads program (sketch) loaded via USB into flash memory
 - 1KB = 1024 bytes
- Example: FreeBSD
 - FreeBSD is a Unix-like OS via Berkeley Software Distribution
 - Commonly abbreviated as BSD
 - The Shell executes the 'fork()' command to create process(es)
 - Then, it executes 'exec()' to load programs into process(es)
 - Finally, the Shell waits for process(es) to terminate or continues to carry out user commands

- A process can exit with:
 1. An exit code of 0
 - Process finished with no errors
 2. An exit code greater than 0
 - Represents an error code
 3. An exit code less than 0
 - The process was not created
 - Running processes in parallel is/was not possible in DOS operating system
- Why Applications Are OS Specific?
 - Each operating system provides its own unique system calls
 - i.e. Own file formats, etc.
 - This is why an application built for Windows will not work on a Mac
 - The system calls for Windows are different from a Mac
 - i.e. Both have different ways of creating processes, threads, etc.
 - Apps can run on different operating systems, If:
 - Written in a high level language like Python, Ruby, etc., and if there is an interpreter available for that language on multiple operating systems
 - The app is written in a language that includes a virtual machine (VM), that the app can run on
 - i.e. Java's motto is "Write Once, Run Anywhere", and this is due to the Java Virtual Machine, that runs on multiple operating systems
 - A standard language, like C, needs to be compiled separately on each operating system in order to run
 - The Application Binary Interface (ABI) is the architecture equivalent of the API
 - The ABI defines how different components of binary code can interface for a given OS on a given architecture, CPU, etc.
- Operating System Design Challenges
 - Even though the internal structure of different operating systems can vary widely, their goals are pretty much the same
 - These goals can be divided into:
 - User goals
 - The OS should be convenient to use, easy to learn, reliable, safe, and fast
 - There may be a trade-off between these goals
 - i.e. More reliability at the cost of speed
 - System goals
 - The OS should be easy to design, implement, and maintain, as well as flexible, reliable, error-free, and efficient
 - There can be a trade-off between these goals
 - i.e. More flexibility at the cost of making the system more prone to errors; less

error-free

- These goals are affected by choice of hardware, type of system, customer base, etc.
- Implementation
 - Operating systems are made using a mix of programming languages
 - i.e.
 - The lowest levels are in Assembly
 - The main body is in C
 - System programs are in C, C++, scripting languages like PERL, Python, shell scripts, and other high-level languages
 - Even though high-level languages are easier to code and port to other hardware, they come at the cost of slow(er) performance
 - C and Assembly are still used to write major components of operating systems
 - Emulation can allow an OS to run on non-native hardware
- Operating System Monolithic Structure
 - The original Unix OS consists of two separable parts:
 1. System Programs
 2. Kernel
 - Consists of everything below the system call interface, and above the physical hardware
 - Provides the file system, CPU scheduling, memory management, virtual memory, etc., and other operating system functions
 - This is a large number of functions for one level
 - The advantage to this is:
 - Little overhead in the system call interface
 - Communication within the kernel is fast
 - The disadvantage is that:
 - It is difficult to implement and extend
 - Changes in one component can affect changes to another component
 - i.e. Changing how the file-system works can introduce bugs into the memory-management component
- Traditional Unix System Structure
 - The original Unix structure was monolithic
 - This makes it difficult to implement and extend
 - This is very bad
 - There is little overhead in the system call interface
 - This is good
 - Communication within the kernel is fast
 - This is good
- Linux System Structure

- The Linux OS relies on a Monolithic plus modular design
 - In a Monolithic plus modular design, each component is placed into its own block of code and these blocks comprise the system call interface
 - The benefit is that changes in one component/block, only affect that block/component
 - i.e. Changing the CPU scheduler won't alter the file system or memory management blocks/components
 - Remember: the Linux system structure is modular based
- Layered Approach
 - The operating system is divided into a number of layers/levels
 - Each subsequent layer is built on top of lower layers
 - The bottom layer, layer 0, is the hardware
 - The highest layer, layer N, is the user interface
 - Each layer uses functions/operations and services of only lower level layers
 - i.e. Layer 3 uses Layer 2, 1 and 0, but Layer 1 does not use Layer 2 or 3
 - Advantages are
 - System verification and debugging are simplified
 - Each layer hides the existence of certain data structures, operations, and hardware from higher level layers
 - Disadvantage(s)
 - Performance of this system is poor
 - The Layered Approach is borrowed from telecommunications
- Microkernels
 - A microkernel design moves as much code from the kernel into user space as possible
 - Mach is an example of microkernel
 - Mac OS X and Darwin, their kernels, are partly based on Mach
 - OS X is based on Darwin
 - Communication takes place between user modules using message passing
 - In other words, when a function makes a call, that message gets passed from user-space into kernel-space, then back into user-space, where the correct module is called and executes the required function, and then once it finishes, it passes a message from user-space into kernel-space, then back into user-space where the original calling function is given an exit code
 - Benefits are
 - Kernel is more reliable and faster
 - Easier to extend a microkernel
 - Easier to port operating systems to new architectures
 - Less code is running in kernel mode, and more in user-space
 - This makes the kernel more secure
 - See below

- Microkernels are more secure than other structures
 - This is because only the core structures are in the kernel, and everything else is in user-space
 - This makes it easier to debug, test all possible combinations, and decreases the attack vector because there are less possible holes in the shrunken (kernel) code base
 - Remember: The kernel controls everything
 - A malfunctioning kernel is a malfunctioning system
- Detriments are
 - The performance overhead of user-space to kernel-space communication
 - Running back and forth between user-space and kernel-space is costly
 - Passing messages in the kernel-space is quicker
- Microkernel System Structure
 - Assigns only a few essential functions to the kernel
 - Non-essential components from the kernel are implemented as user level programs
- Modules
 - Many modern operating systems implement loadable kernel modules; referred to as LKM
 - Each LKM...
 - Uses an object-oriented approach
 - Each core component is separate
 - Talks to other modules over known interfaces
 - Is loadable as needed within the kernel
 - You use only what you need to use
 - This is similar to layers, but it is much more flexible
 - This type of design is common in modern implementations of UNIX, such as Linux, macOS, and Solaris, as well as Windows, which is not part of the UNIX family
- Hybrid Systems
 - Most modern operating systems do not follow one model; instead, they are a hybrid
 - The hybrid model combines multiple approaches to address performance, security, usability needs, etc.
 - The major operating systems have adopted a hybrid model
 - i.e.
 - Linux and Solaris use a monolithic kernel model, but they incorporate modular aspects for dynamic loading of functionality
 - Windows is mostly monolithic, but incorporates microkernels for different subsystem personalities
 - This model was chosen for speed
 - Apple's Mac OS X is hybrid, and layered
 - The kernel consists of Mach-based Darwin microkernel

and BSD Unix parts, plus I/O kit and dynamically loaded modules

- The dynamically loaded modules are called kernel extensions
 - This is where the term "kexts" comes from
 - Programming environment is Aqua UI plus Cocoa
 - The application framework "Cocoa" provides an API for the Objective-C and Swift programming languages
 - Core frameworks support graphics and media
 - i.e. Quicktime, OpenGL, etc.
 - In iOS, the environment (home-screen) is called SpringBoard
- Darwin
- Darwin consists primarily of the Mach microkernel and the UNIX kernel
 - There are two system call interfaces:
 1. Mach system calls (aka Mach traps)
 - Note: Do not confuse Mach traps with trap exceptions
 2. BSD Unix system calls
 - POSIX
 - IPC = Inter-process communication
- Android (1)
- Developed by Open Handset Alliance (mostly Google)
 - Android is open source
 - Anybody can take the code and develop their own version
 - i.e. GrapheneOS, LineageOS, Pixel ROM, AOKP, etc.
 - Linux is also open source
 - There are more than 600 variants of Linux
 - i.e. Manjaro, Arch, Ubuntu, Kali, Etc.
 - The stack is similar to Apple's iOS; which was built for iPhone and iPad
 - Note: iPads now run iPadOS, which is a spin-off of iOS
 - Android is based on the Linux kernel
 - This provides it with process, memory, and device-driver management
 - And other features/services
 - Adds power management
 - This was added later because Android is built to primarily run on mobile devices, which have a limited capacity battery
- Android (2)
- Android was originally running on a virtual machine called Dalvik
 - Since 2015, Dalvik runtime is no longer maintained or available to the public
 - Dalvik has been replaced by the Android RunTime (ART) VM

- In addition to the VM, Android ships with essential core set of libraries
- Google switched Android from Dalvik to ART to fully control everything
 - ART was built to be better, faster, and more flexible than Dalvik
- Apps are developed in Java plus the Android API
 - Java class files are compiled to Java bytecode, and then translated to executable that runs in ART VM
- Android libraries include frameworks for web browser (webkit), database (SQLite), multimedia, smaller libc, OpenGL, etc.
- Android (3)
 - Hardware Abstraction Layer (HAL)
 - The HAL is a map between generic hardware commands and responses and those unique to a specific platform
 - By abstracting all hardware, the HAL provides applications with a consistent view independent of specific hardware
 - Example of hardware: GPS chip, camera, microphone, etc.
 - Bionic is a C library developed by Google to avoid using GNU's C
 - Bionic was developed because of:
 - Issues with the general public license (GPL)
 - Google wanted more control and the ability to fine tune things
- Android Runtime
 - JNI stands for Java Native Interface
 - JNI allows us to access hardware functions without going through the virtual machine
 - Virtual machines allow developers to run their code on different platforms with little to no additional effort
 - Every Android application runs in its own process with its own instance of the ART VM
 - The ART VM executes files in the Dalvik Executable (.dex) format
 - Android includes a set of core libraries that provides most of the functionality available in Java's core libraries
 - In other words, existing Java programmers can easily write apps for Android devices
 - To execute an operation, the ART VM calls on the corresponding C/C++ library using the Java Native Interface (JNI)
- Power Management
 - Power management is important because Android runs on mobile devices, which run on batteries
 - Being able to control how much power is used is essential in conserving battery power and capacity
 - Android adds two features to the Linux kernel to enhance the

ability to perform power management

1. Alarms

- Implemented in the Linux kernel and is visible to the app developer through the AlarmManager in the Runtime core libraries
 - The implementation is in the kernel, so an alarm can ring even if the system is in sleep mode

2. Wakelocks

- Prevents an Android system from entering sleep mode
 - i.e. Keep phone on while user watches video or plays game
- Wakelocks are requested through the API whenever an application requires an I/O to remain powered on

- End Of OS Structures

- Operating Systems are among the most complex pieces of software ever developed
 - When you change something in the OS, don't change too many things or else you will break something