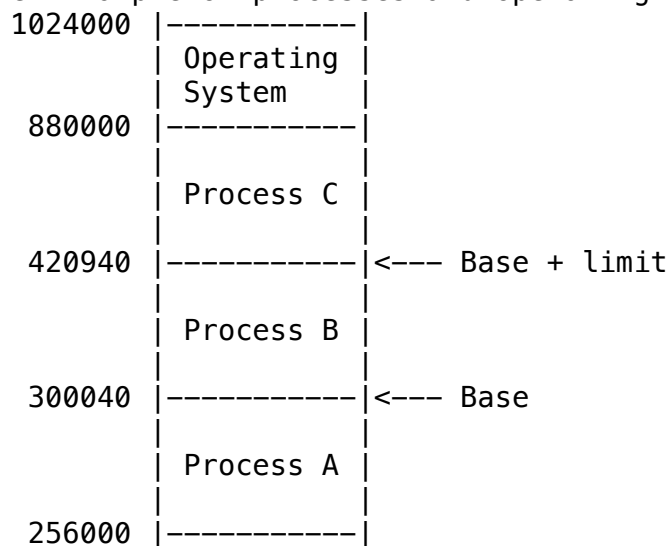



Lecture.9.Memory.Management.txt

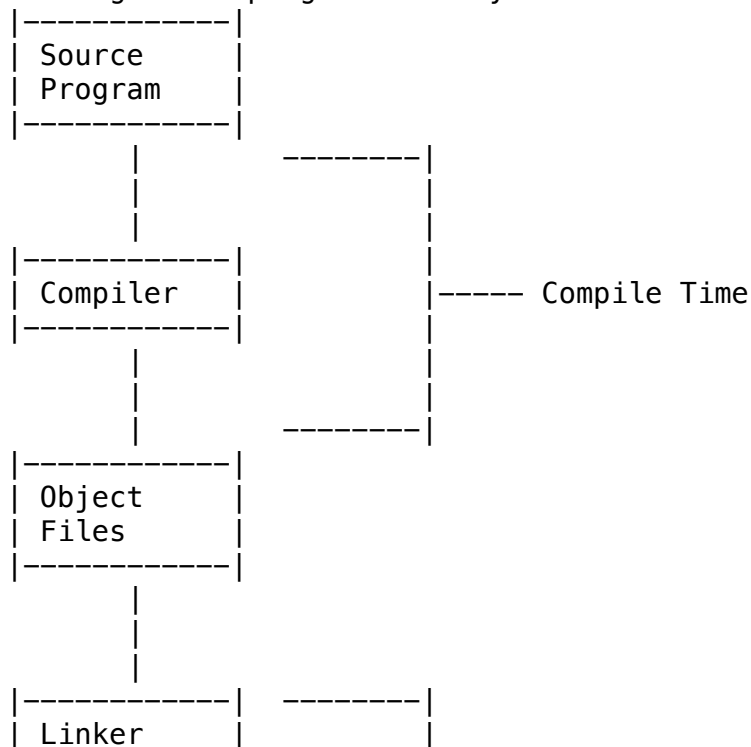
- Background
 - Role in programming language
 - Instructions
 - Are usually read-only
 - In normal programming, instructions are read-only
 - We do not allow code to rewrite itself
 - Some embedded systems allow rewriting of instructions
 - Variables
 - Read and write, only
 - Constants
 - Read only
 - To access memory, we need the address
 - When writing to memory, we need an address and data
 - Memory unit only sees a stream of:
 - Addresses + Read requests
 - Addresses + Data + Write requests
 - Main memory and registers are the only storage mediums that the CPU can access directly
 - The CPU cannot write into secondary storage
 - The CPU can only read and write data into registers and memory that is closest to the CPU, which is main memory
 - Main memory is cache and RAM
 - Protection of memory is required to ensure correct operation
 - We do not want different processes to overlap in memory space
- Protection
 - OS has to be protected from access by user processes, as well as protect user processes from one another
 - We can provide this protection by base and limit registers
 - i.e. Example of processes and operating system inside memory

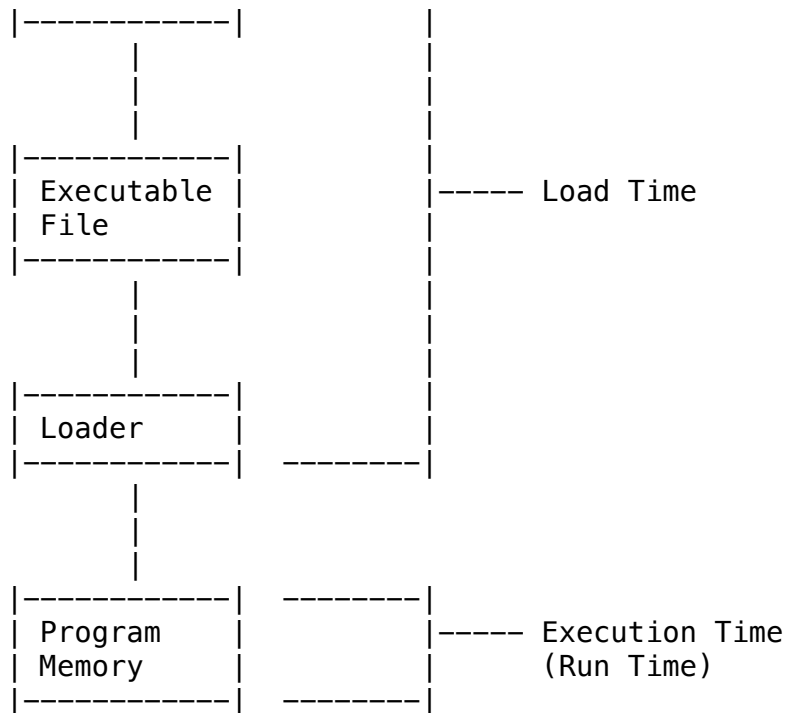




- The operating system is protected, because no process can access its memory
 - No process can read or write to the operating system's memory
 - The operating system's memory is at the top of the memory address space
- Each process has its own space in memory
- The memory space of 'Process B' is defined by the base register, and the top, base plus limit
 - All addresses pertaining to 'Process B' are in the range 300040 and 420940
- A base and a limit register, pair, define a logical address space
 - The logical address space is on the CPU side
 - Real memory is on the physical address space
 - In a system, we need both physical and logical address space
- Hardware Address Protection
 - Is used to ensure that the CPU does not access memory it is not allowed to access
 - This is achieved through base and base + limit
 - CPU must check every memory access generated in user mode to be sure it is between base and limit for that user
 - If the generated memory address is lower than base, then an error is raised
 - If the address is not within base and base + limit, then memory access is denied
 - Protection of memory space is accomplished by having the CPU hardware
 - The instructions to loading the base and limit registers are privileged
 - Can be used only by the OS
- Address Binding
 - Addresses are created during the stages of program development
 - Program resides on a disk as a binary executable file
 - Must be loaded into memory address 0000
 - Inconvenient
 - Addresses are represented in different ways at different stages of a program's life
 - Addresses in the source program are generally symbolic
 - i.e. The variable 'count'
 - 'count' references some space in address in memory
 - This notation is easier for humans to understand
 - Compiled code addresses bind to relocatable addresses
 - i.e. 14 bytes from beginning of this module
 - Linker or loader will bind relocatable addresses to absolute addresses

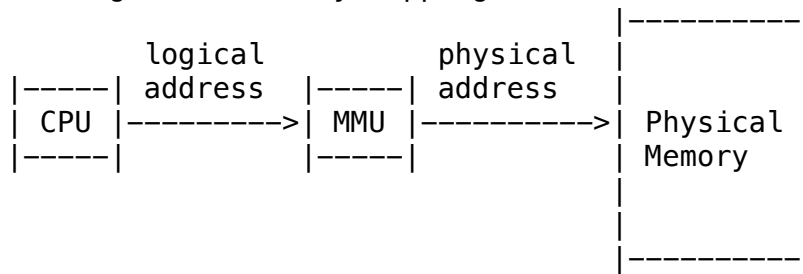
- i.e. 74014
 - This is the address inside memory
- Each binding maps one address space to another
 - Symbolic address -> Compiled code address -> Absolute address
- Binding Of Instructions & Data To Memory
 - Address binding of instructions and data to memory addresses can happen at three different stages:
 - Compile time: If memory location is known in advanced, then absolute address(es) can be generated
 - Load time: Must generate relocatable code if memory location is not known at compile time
 - Execution time: If the process can be moved during its execution from one memory segment to another, binding must be delayed until run time
 - Most (99%) of the address binding takes place on execution time
 - Need hardware support for address mapping
- Multistep Processing Of A User Program
 - Compile time
 - Absolute code
 - Load time
 - Relocatable code
 - Execution time
 - Run time binding
 - i.e. Diagram of program life cycle



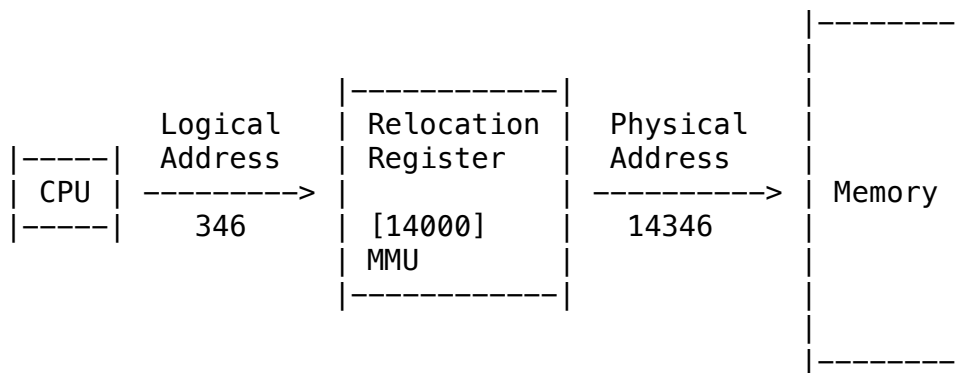


- The first step is to compile source code to create object file
 - Multiple source files result in multiple object files
 - Then, the Linker takes all the object files and creates relocatable code
 - The job of the Loader is to load the code into memory
 - Execution time is also referred to as run time
-
- Logical VS. Physical Address Space
 - The concept of logical address space that is bound to a separate physical address space is central to proper memory management
 - Logical address: Generated by the CPU
 - Also referred to as virtual address
 - Physical address: Address seen by the memory unit
 - Directly dealing with physical addresses is not an ideal solution
 - For instance, moving data in memory changes the physical address, but not the logical address
 - This is abstraction
 - Logical and physical addresses are the same in compile-time and load-time address-binding schemes
 - Logical (virtual) and physical addresses differ in execution-time address-binding scheme
-
- Memory Management Unit (MMU) (1)
 - Is used to translate logical addresses to physical addresses
 - The runtime mapping from virtual to physical addresses is done by a hardware device called memory-management unit (MMU)

- We can choose from many different methods to accomplish runtime mapping
- i.e. Diagram Of Memory Mapping



- Memory Management Unit (MMU) (2)
 - One way to implement a MMU is to use a relocation register
 - A generalization of the base-register scheme
 - The base register now called relocaion register
 - The value in the relocation register is added to every address generated by a user process at the time it is sent to memory
 - i.e.



- The value of the relocation register is added to the value of the logical address
 - Logical address + Relocation register = Physical address
 - $346 + 14000 = 14346$
 - If the base is at 14000, then an attempt by the user to address location 346 is dynamically relocaed to location 14346
 - Refer to diagram above
 - The MMU is responsible for generating the relocation register address value
- Dynamic Loading
 - Initially we should have entire program and all data of a process to be in physical memory for the process to execute
 - Size of a process has been limited to the size of physical memory
 - If the size of the process is bigger than physical memory, then it cannot be executed
 - For better memory-space utilization unused routine is never loaded

- Only routines that are used are loaded into memory
 - Unused routines are kept in secondary memory
- All routines kept on disk in relocatable load format
- Useful when large amounts of code are needed to handle infrequently
 - For instance, error handling is a lot of code, and is infrequently utilized. Therefore, it is better to keep it on secondary memory, instead of clogging up main memory with code that may never be used
- Dynamic loading requires no special support from the OS
 - It is implemented through program design
 - The OS can help by providing libraries to implement dynamic loading
- Dynamic Linking
 - Dynamically linked libraries (DLLs) are system libraries that are linked to user programs when the programs are run
 - DLLs are added during execution time
 - Static linking: System libraries and program code combined by the loader into the binary program image
 - In other words, the libraries are linked in advance
 - Dynamic linking: Linking postponed until execution time
 - This is useful because the compiled program occupies less space in memory
 - With dynamic linking, multiple programs can share libraries
 - This is not possible with static linking
 - Without this facility, each program on a system must include a copy of its language library – or at least the routines referenced by the program – in the executable image
 - Decreases size of executable image
 - Libraries can be shared by multiple processes, only one instance of DLL in main memory
 - This saves space in memory
 - Due to limitation of resources, we need to optimize the use of these resources and prevent wasting
- Contiguous Memory Allocation
 - Main memory must support/store both OS and user processes
 - Main memory needs to be allocated in the most efficient way possible
 - Contiguous allocation is one way to achieve this
 - Main memory is usually divided into two partitions:
 - Resident OS usually held in low memory with interrupt vector
 - User processes then held in high memory
 - Each process is contained in a single contiguous section of memory
 - Each process has its own range of memory
- Contiguous Allocation
 - Relocation registers used to protect user processes from each


- other, and from changing operating-system code and data
 - Base register contains value of smallest physical address
 - Limit register contains range of logical addresses
 - Each logical address must be less than the limit register
 - MMU maps logical address(es) dynamically
 - Can then allow actions such as kernel code being transient and kernel changing size

- Hardware Support For Relocation & Limit Registers

- The logical address must be between the 'limit register' and the 'relocation register'
 - Used for memory protection
- The relocation register contains the value of the smallest physical address
 - i.e. Relocation = 100040
 - If the logical address is greater than this value, then it will generate an error
- The limit register contains the range of logical addresses
 - i.e. Limit = 74960
 - This is the maximum value of the address the CPU can use to access the memory

- Variable Partition

- The relocation-register scheme provides an effective way to allow the operating system's size to change dynamically
- Operating system maintains information about:
 - Allocated partitions
 - Free partitions (hole)
- i.e. Process Allocation Diagram



- The operating system's memory is on top of all other memory
- In the first block, there are 3 processes in memory: 5, 8, and 2
- This method of allocation creates fragmentation
 - If process 5 completes, there are 2 free blocks of memory, but process 8 cannot run
 - Contiguous allocation does not work with fragmentation
 - Paging solves this problem

- Allocation
 - What happens when there isn't sufficient memory to satisfy the demands of an arriving process?
 - You can try to rearrange the processes in memory and free up some space; apart from this, there isn't much that can be done
- Dynamic Storage Allocation Problem
 - How to satisfy a request of size n from a list of free holes?
 - First fit: Allocate the first hole that is big enough
 - Best fit: Allocate the smallest hole that is big enough, must search entire list, unless ordered by size
 - Produces the smallest leftover hole
 - Takes more time than the other methods
 - Worst fit: Allocate the largest hole, must also search entire list
 - Produces the largest leftover hole
 - First fit and Best fit better than Worst fit in terms of speed and storage utilization
 - Best fit is better than first fit and worst fit in terms of storage utilization
 - Best fit and worst fit take more time than first fit
 - First fit is the best in terms of speed
- Fragmentation (1)
 - Both the first fit and best fit strategies for memory allocation suffer from external fragmentation
 - External fragmentation: Total memory space exists to satisfy a request, but it is not contiguous
 - Internal fragmentation: Allocated memory may be slightly larger than requested memory
 - Leads to loss of memory
 - Is a result of paging
 - Usually, external fragmentation is worse than internal fragmentation
 - First fit analysis reveals that given N blocks allocated, another $0.5N$ blocks lost to fragmentation
 - 1/3 may be unusable \rightarrow 50% rule
 - i.e. Given 3 blocks, 1 may be lost due to fragmentation
- Fragmentation (2)
 - Reduce external fragmentation by compaction
 - Shuffle memory contents to place all free memory together in one large block
 - Compaction is possible only if relocation is dynamic, and is done at execution time
 - The simplest compaction algorithm is to move all processes toward one end of memory
 - Another possible solution is to permit the logical address space of processes to be noncontiguous thus allowing a process to

- allocate physical memory wherever such memory is available
 - This is called Paging
 - Paging allows us to utilize all free space in memory


- Paging

- Physical address space of a process can be noncontiguous
 - This means that a process can be divided into parts, and each part can be stored in memory at any location
 - i.e. For a given program, some instructions can be stored at address `N`, and the remaining instructions can be stored at address `N+1`
 - Avoids external fragmentation
 - Avoids problem of varying sized memory chunks
- Paging refers to the logical address
- Frames refer to the physical memory address
- Paging is done by dividing physical memory into fixed-sized blocks called frames
 - Size is power of 2
 - Between 512 bytes and 16 Megabytes
- The steps involved in paging are:
 - Divide logical memory into blocks of same size called pages
 - To run a program of size `N` pages, need to find `N` free frames
 - Pages are associated with frames, and both have the same size
 - Set up a page table to translate logical to physical addresses
 - This is the simplest implementation of the memory management unit (MMU) for non-contiguous processes
 - Backing store likewise split into pages

- Address Translation Scheme


- Address generated by CPU is divided into two parts:
 - Page number (p)
 - Used as an index into a page table which contains base address of each page in physical memory
 - Page offset (d)
 - Combined with base address to define the physical memory address that is sent to the memory unit

- i.e. Diagram



- For given logical address space 2^m and page size 2^n
 - The page number is `m - n`
 - The page offset is `n`
- Logical address space is now toally separated from the physical address space

- The CPU generates a logical address that consists of a page number and offset
- The page table translates the logical address into a physical address
- Paging Hardware
 - The page number is used as an index into a per-process page table
 - The page table contains the base address of each frame in physical memory
 - Offset `d` does not change
- Paging Model Of Logical & Physical Memory
 - i.e. Diagram of Logical Memory



- i.e. Diagram of Page Table

Frame Number	
0	1
1	4
2	3
3	7

Page Table

- i.e. Diagram of Physical Memory

Frame Number	
0	-----
1	page 0
2	-----
3	page 2

4	page 1
5	-----
6	-----
7	page 3

Physical Memory

- According to the diagrams above, logical address 0 is page 0, offset 0. Indexing into the page table, we find that page 0 is in frame 1

- Paging is a form of dynamic relocation

- Paging Example

- i.e. Diagram of Logical Memory

0	a
1	b
2	c
3	d
4	e
5	f
6	g
7	h
8	i
9	j
10	k
11	l
12	m
13	n
14	o
15	p

Logical Memory

- $2^m = 16$; $m = 4$
- Each page has 4 elements
 - Thus, $n = 4$
- These (logical) addresses are generated by a CPU
- On logical memory, each unit is called a page

- i.e. Diagram of Page Table

0	5
1	6
2	1

3	2
---	---

Page Table

- Size: 2^n

- $n = 2$

- i.e. Diagram of Physical Memory

0	
4	i j k l
8	m n o p
12	
16	
20	a b c d
24	e f g h
28	

Physical Memory

- On physical memory, each unit is called a frame

- Referring to the diagrams above:
 - Page 0 on the logical memory is associated to frame 5 in the page table
 - Frame 5 is at index number 20 in physical memory
 - Logical address: $n = 2$ and $m = 4$
 - Page size of 4 bytes, physical memory of 32 bytes (8 pages)
 - Logical address 0 maps to physical address 20
 - $[(5 \times 4) + 0] = 20$
 - 3 (page 0, offset 3) \rightarrow 23
 - $[(5 \times 4) + 3] = 23$
 - 4 (page 1, offset 0) \rightarrow 24
 - $[(6 \times 4) + 0] = 24$
- A page table is used to associate a page to a particular frame
 - i.e. Page 0 on logical memory is translated to Frame 5 on physical memory
- Paging Calculating Internal Fragmentation
 - External fragmentation is when the system has enough free memory to satisfy the demands of a process, but the free memory is not contiguous
 - The problem of external fragmentation is solved by introducing pages
 - Pages allows us to store the content of a process across different locations inside physical memory
 - Paging causes internal fragmentation
 - Calculating internal fragmentation is shown below
 - Internal Fragmentation Example:
 - Given page size = 2048 bytes
 - For process size = 72766 bytes, we need 36 pages
 - Internal fragmentation = $36 \times 2048 - 72766 = 962$ bytes
 - Worst case fragmentation = frame - 1 byte
 - Average fragmentation = $1/2$ frame size
 - Pages are typically either 4 KB or 8 KB in size
 - Some CPUs and operating systems even support multiple page sizes
 - Windows 10 = 4KB and 2MB
 - Linux = 4KB and huge pages
 - Note: The combination of multiple page sizes works best
- Example
 - Frequently, on a 32-bit CPU, each page-table entry is 4 bytes long
 - A 32-bit entry can point to one of 2^{32} physical page frames
 - If the frame size is 4KB (2^{12}), then a system can address 16TB (2^{44}) bytes of physical memory
- Free Frames
 - A list of free frames needs to be maintained
 - Once a frame is used, it cannot be associated to a page anymore
 - When a process arrives in the system to be executed, its size,

- expressed in pages, is examined
- Each page of the process needs one frame
 - i.e. Free Frame Diagram (Before)

Free-Frame

List:

14	13	
13		
18	14	
20		
15	15	
	16	
	17	
	18	
	19	
	20	
	21	


- i.e. Free Frame Diagram (After)

Free-Frame

List:

	13	page 1
15	14	page 0
	15	
	16	
	17	
	18	page 2
	19	
	20	page 3
	21	

0	14
1	13
2	18



- The two diagrams above depict free frames before and after 4 pages allocation
- Implementation Of Page Table
 - A page table is a mapping from pages to frames
 - The hardware implementation of the page table can be done in several ways
 - Page table is kept in main memory
 - Page table base register (PTBR) points to the page table
 - Page table length register (PTLR) indicates size of the page table
 - Problems with storing the page table in main memory:
 - If the page is very big, it will take up a lot of space in memory; contiguously as one portion
 - There may be 2 steps to access the page, and each access requires one step for the page table and another step for the actual data
 - In this scheme every access requires two steps
 - One for the page table, and one for the actual data
 - This approach is reused in many different circumstances
 - The two step memory access problem can be solved by the use of a "special fast lookup hardware cache" called translation look-aside buffers (TLBs)
 - This method puts the most used addresses in a separate buffer
 - As a result, performance is greatly increased
- Translation Look Aside Buffer
 - The TLB is a special hardware cache
 - Some TLBs store address space identifiers (ASIDs) in each TLB entry
 - Uniquely identifies each process used to provide address space protection for that process
 - TLBs are typically small
 - i.e. 64 - 1024 entries of all addresses
 - On a TLB miss, the regular page table is consulted, and the value is loaded into the TLB for faster access next time the same value is used
 - Replacement policies must be considered (if full)
 - Addresses that are not used for an extended period of time are replaced with newer addresses
 - Some entries can be wired down for permanent fast access
- Paging Hardware With TLB
 - When the frame number is obtained, we can use it to assess memory
 - We add the page number and frame number to the TLB, so that they

- will be found quickly on the next reference
 - A system can, and will, work without a TLB
 - However, it will be slower
 - The TLB can only store a subset of the page table
- Effective Access Time
 - Percentage of times that a page number is found in the TLB is called hit ratio
 - An 80% hit ratio means that we find the desired page number in the TLB 80% of the time
 - This is a relatively good hit ratio
 - Example
 - Question: Suppose that takes 10ns to access memory if we find the desired page in TLB, otherwise we need two memory access so it is 20 ns
 - Answer: Effective Access Time (EAT) =

$$(0.80 \times 10) + (0.20 \times 20) = 12\text{ns}$$
 Therefore, 20% slowdown in access time
 - Question: Calculate slowdown in access time for hit ratio of 99%
 - The EAT will be much lower than the hit ratio for 80%
- Memory Protection
 - Some pages in memory need to be dedicated to the operating system, and some pages need to be protected so other processes cannot read them
 - Each process should only be able to access its own memory
 - Memory protection in paged environment is accomplished by protection bits associated with each frame
 - These bits are kept in the page table
 - One bit can define a page to be read-write or read-only
 - Valid-invalid bit attached to each entry in the page table:
 - Valid indicates that the associated page is in the process logical address space, and is thus a legal page
 - Invalid indicates that the page is not in the process' logical address space
 - Or use page-table length register (PTLR)
 - Any violations result in a trap to the kernel
- Valid (V) Or Invalid (I) Bit In A Page Table
 - i.e. Diagram of Page Table With Bit

00000	-----			Valid-Invalid Bit	0	-----
	page 0					
	-----					-----
	page 1				1	
	-----					-----
	page 2	0	2	v	2	page 0
	-----					-----
	page 3	1	3	v	3	page 1
	-----					-----
	page 4	2	4	v	4	page 2
	-----					-----

10,468	-----	3	7	v	5	-----
12,287	page 5	4	8	v	6	-----
	-----	5	9	v	7	page 3
		6	0	i	8	page 4
		7	0	i	9	page 5

						*
						*
						*

						page n

Page
Table


- Any attempt to generate an address in pages 6 or 7, however, will find that the valid-invalid bit is set to invalid, and the computer will trap to the OS
 - The only valid pages are from 0 to 5
 - If the process has a page 6, then it will not get the corresponding frame associated to page 6
 - Because there is nothing in the frame
 - This is kind of like a page fault
 - A page fault results in trying to access a page that is not valid
- Shared Pages
 - Shared code
 - One copy of read-only (re-entrant) code shared among processes
 - i.e. Text editors, Compilers, Window systems, Etc.
 - Similar to multiple threads sharing the same process space
 - Also useful for interprocess communication if sharing of read-write pages is allowed
 - i.e. Multiple processes generate their own logical address space, and all of them are mapped to the same physical address space
 - Private code & data
 - Each process keeps a separate copy of the code and data
 - Nothing is shared
 - The pages for the private code and data can appear anywhere in the logical address space
- Shared Pages Example
 - Shared code helps processes reuse code
 - i.e. Multiple processes using the same library, like the standard C library
 - Three processes sharing the pages for the standard C library 'libc'

- Only one copy of the standard C library need to be kept in physical memory
 - The page table for each user process maps onto the same physical copy of 'libc'
 - Only reading is allowed; no writing
- Note: On a theoretical level, physical memory can refer to either cache or RAM
- Structure Of The Page Table
 - Memory structures for paging can get huge using straight-forward methods
 - Consider a 32-bit logical address space as on modern computers
 - Page size = 4KB
 - $2^{12} = 4096 \text{ B}$
 - Page table would have 1 million entries ($2^{20} = 2^{32} / 2^{12}$)
 - If each entry is 4 bytes -> Each process 4MB of physical address space for the page table alone
 - We don't want to allocate the page table contiguously in main memory, because it takes a lot of space
 - One simple solution is to divide the page table into smaller units
 - i.e.
 - Hierarchical paging
 - Hashed page Tables
 - Inverted page tables
- Hierarchical Page Tables
 - Break up the logical address space into multiple page tables
 - A simple technique is a two-level page table
 - i.e. Outer page table, and inner page table
 - The outer page table points to the inner page table
 - Inner page table points to location in physical memory
 - The page tables do not have to be contiguous and can be placed anywhere in memory
 - We then page the page table
 - Note: This does not save any space in memory
- Two Level Paging Example
 - A logical address (on 32-bit machine with 4K page size) is divided into two numbers:
 - A page number consisting of 20 bits
 - 'p1' and 'p2' occupy 10 bits
 - A page offset consisting of 12 bits
 - $2^{12} = 4096 \text{ B} = 4 \text{ KB}$
 - Since the page table is paged, the page number is further divided into:
 - A 10-bit page number
 - A 12-bit page offset
 - Thus, a logical address is as follows:

Page number		Page offset
p1	p2	d
10	10	12

- Where 'p1' is an index into the outer page table, and 'p2' is the displacement within the page of the inner page table
 - Known as forward-mapped page table
- Note: Hierarchical paging helps the system to better utilize memory space

- Address Translation Scheme
 - i.e. Diagram of Address Translation



- Address translation for a two level 32-bit paging architecture
 - 'p1' is used to go to the outer table
 - Outer table points to the inner table
 - 'p2' gives offset to the elements
 - 'p1' is used as an offset to find the pointer to the outer table
- 64 Bit Logical Address Space
 - Even two level paging scheme is not sufficient!
 - If page size is 4KB (2^{12}), then page table has 2^{52} entries
 - If two level scheme, inner page tables could be 2^{10} with 4-byte entries
 - Outer page table has 2^{42} entries
 - One solution is to add a 2nd outer page table

- Three Level Paging Scheme

- In the following example, the 2nd outer page table is still 2^{32} bytes in size

- i.e. Diagram of Two Level Paging

outer page	inner page	offset
p1	p2	d
42	10	12

- This is the previous paging scheme

- i.e. Diagram Three Level Paging

2nd outer page	outer page	inner page	offset
p1	p2	p3	d
32	10	10	12

- This is the new paging scheme for 64-bit logical addresses

- The '2nd outer page' can be divided into many pages

- In modern microprocessors, it goes up to 4 tables

- The next step would be a four-level paging scheme, where the second-level outer page table itself is also paged, and so forth

- Hashed Page Tables

- Common in address spaces with more than 32 bits
- The virtual page number is hashed into a page table
- This page table contains a chain of elements hashing to the same location
- Each element contains:
 1. Virtual page number
 2. Value of the mapped page frame
 3. A pointer to the next element
- Virtual page numbers are compared in this chain searching for a match

- Hashed Page Table

- The hash function gives the chain of the page frames
- If a match is found, the corresponding physical frame is extracted
- Variation for 64-bit addresses is clustered page tables
- Similar to hashed but each entry refers to several pages, such as 16, rather than 1
- Especially useful for sparse address spaces
 - Where memory references are non-contiguous and scattered

- Inverted Page Table

- Rather than each process having a page table and keeping track of all possible logical pages, track all physical pages
- One entry for each real page of memory
 - An inverted page table has one entry for each real page, or frame, of memory


- Although this scheme decreases the amount of memory needed to store each page table, it increases the amount of time needed to search the table when a page reference occurs
- Use hash table to limit the search to one or at most a few page-table entries
- TLB can accelerate access
- Inverted Page Table Architecture
 - < pid, pagenumber > is presented to the memory subsystem
 - The inverted page table is then searched for a match. If a match is found at entry 'i' then the physical address < i, offset > is generated
 - The 'offset' in the page table is the frame
 - In the inverted page table architecture, we look for the content, and the 'offset' is the frame
 - In hashed page table, we look for the offset and the content is the frame
 - Observation
 - How to implement shared memory?
 - This is not possible, because inverted page tables always point to the particular page
 - This is a limitation of inverted page table architecture
 - Map multiple virtual addresses to the same physical address
- Swapping
 - There are two types of addresses:
 - Virtual Address
 - Created by the CPU
 - Physical Address
 - The address on which the data is stored in memory
 - The job of the memory management unit (MMU) is to translate virtual addresses to physical addresses, and vice versa
 - A process can be swapped temporarily out of memory to a backing store, and then brought back into memory for continued execution
 - The backing store is usually a hard drive (HDD)
 - Backing store: Fast disk large enough to accommodate copies of all memory images for all users
 - Roll out, roll in: Swapping variant used for priority-based scheduling algorithms
 - Major part of swap time is transfer time
 - Total transfer time is directly proportional to the amount of memory swapped
- Schematic View Of Swapping
 - A process' memory can be swapped out, from the main memory into the backing store, for another process
 - i.e. Process P1 is swapped for Process P2
 - The backing store is greater in size than main memory
 - But processes need to be in main memory, in order for them

- to be executed
 - Swapping makes it possible for the total physical address space of all processes to exceed the real physical memory of the system!
 - Swapping can take a lot of time
- Context Switch Time Including Swapping
 - If next processes to be put on CPU is not in memory, need to swap out a process and swap in target process
 - Context switch time can be very high
 - Example: 100MB process swapping to hard disk with transfer rate of 50MB/sec
 - Swap out time of 2000ms
 - Plus swap in of same sized process
 - Total context switch swapping component time of 4000ms
 - 4 seconds
 - Can reduce if reduce size of memory swapped
 - By knowing how much memory really being used
 - Context switching with swapping increases the time by a huge factor
 - Normally, context switching is very fast and takes micro-seconds. But swapping makes the entire thing very slow
- Swapping In Modern OS
 - Modified versions of swapping are found on many systems
 - i.e. UNIX, Linux, and Windows
 - Normally, swapping is disabled
 - But, it is started if more than threshold amount of memory is allocated
 - And then, disabled again once memory demand reduced below threshold
 - i.e. If 80% or more of main memory has been allocated, then swapping is enabled
- Schematic View Of Swapping With Paging
 - A subset of pages for processes A and B are being paged-out and paged-in respectively
 - Pages of a process rather than an entire process can be swapped
 - Unused parts of a process are moved out
 - Most systems, including Linux and Windows, now use swapping with paging
- Swapping On Mobile Systems
 - Not typically supported
 - Due to lack of huge hard drive (HDD)
 - Flash memory based: Small amount of space, limited number of write cycles, poor throughput between flash memory and CPU on mobile platform
 - Thus, swapping is not a good option
 - iOS asks apps to voluntarily relinquish allocated memory:

- Read-only data thrown out and reloaded from flash
 - Failure to free can result in termination
- Android terminates apps if low free memory, but first writes application state to flash for fast restart
- Both iOS and Android support paging
- Implementation
 - The Intel 32 and 64-bit architectures
 - These are the current most popular architectures
 - Intel x86-64
 - ARM architecture
 - Mostly used on mobile devices
- Example: The Intel 32 & 64 Bit Architectures (1)
 - The dominant industry chips are:
 - Pentium CPUs are 32-bit
 - Called IA-32 architecture
 - Current Intel CPUs are 64-bit
 - Called IA-64 architecture
 - Many variations in the chips, cover the main ideas here
- Example: The Intel 32 & 64 Bit Architectures (2)
 - The pages may be different sizes when translating from logical to physical addresses
 - This is how it works on most CPUs
 - CPU generates logical address
 - Selector given to segmentation unit which produces linear addresses
 - Linear address given to paging unit
 - Which generates physical address in main memory
 - Paging units form equivalent of MMU
 - Pages sizes can be 4KB or 4MB
 - These are typical page sizes
 - 4KB = small
 - 4MB = large

- Logical To Physical Address Translation In IA-32

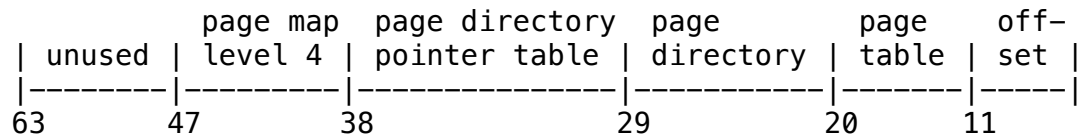
- i.e. Diagram of Address Translation



- Note: PA = Physical Address
- Between logical and physical address conversion, there are two steps:
 - Segmentation
 - Creates linear address
 - Paging

- Creates physical address
 - Memory management divided into two components:
 - Segmentation
 - Paging
 - The CPU generates logical addresses, which are given to the segmentation unit
 - The segmentation unit produces a linear address for each logical address
 - The linear address is then given to the paging unit, which in turn generates the physical address in main memory
- Intel IA-32 Segmentation (1)
- The logical address comprises of two parts
 - Selector
 - Offset
 - The logical address space of a process is divided into two partitions
 - Up to 8K segments that are private to that process
 - Information kept in the local descriptor table (LDT)
 - Up to 8K segments that are shared among all the processes
 - Global descriptor table (GDT)
- Intel IA-32 Segmentation (2)
- The logical address is a pair (selector, offset)
 - The selector is a 16-bit number:

-----	-----	-----
s	g	p
-----	-----	-----
13	1	2
 - Where:
 - `s`: The segment number
 - `g`: Indicates whether the segment is in the GDT or LDT
 - `p`: Deals with protection
 - The offset is a 32-bit number specifying the location of the byte within the segment in question
- Intel IA-32 Paging Architecture
- The address translation scheme for this architecture is similar to the scheme shown in Figure on Page 39
- Intel x86-64
- Current generation Intel x86 architecture
 - 64 bits is ginormous (> 16 exabytes)
 - 1 EB = 1,000,000,000,000,000,000 B
 - In practice only implement 48-bit addressing
 - Page sizes of 4 KB, 2 MB, 1 GB
 - Four levels of paging hierarchy
 - Can also use page address extension (PAE) so virtual addresses are 48-bits and physical addresses are 52-bits
 - i.e. Diagram of Current Paging Table



- There are more inner tables in today's CPU architecture

- Example: ARM Architecture
 - Dominant mobile platform chip
 - i.e. Apple iOS, Google Android, Etc.
 - Modern, energy efficient 32-bit CPU
 - 4 KB and 16 KB pages
 - An architecture should be able to address different sizes
 - 1 MB and 16 MB pages
 - Termed sections
 - One level paging for sections, two level for smaller pages
- End
 - Operating Systems are among the most complex pieces of software ever developed