

Operating Systems: Synchronization Tools and examples – Part II

Neerja Mhaskar

Department of Computing and Software, McMaster University, Canada

Acknowledgements: Material based on the textbook Operating Systems Concepts (Chapters 6 & 7)

Semaphore

- Synchronization tool that provides more sophisticated ways (other than Mutex locks) for processes to synchronize their activities.
- Semaphore **S** – integer variable (usually initialized with a +ve integer)
- Can only be accessed via two indivisible (**atomic**) operations
 - **wait()** and **signal()**
 - Originally called **P()** and **V()**

- **wait() operation:**

```
wait(S) {  
    while (S <= 0)  
        ; // busy wait  
    S--;  
}
```

- **signal() operation:**

```
signal(S) {  
    S++;  
}
```

Semaphore Usage

- **Counting semaphore** – integer value can range over an unrestricted domain
- **Binary semaphore** – integer value can range only between 0 and 1
 - Behaves like a mutex lock
- Semaphores can solve various synchronization problems
 - Used as locks (binary semaphores)
 - To solve synchronization problems (may or may not access shared data)
 - Used to control access to a given resource consisting of a finite number of instances (counting semaphore)

Semaphores – synchronization problems

- Consider two concurrently running processes P_1 and P_2 that require S_1 to happen before S_2 (here P_1 and P_2 may or may not access/modify shared variables.)
- **Solution:** Create a semaphore “**synch**” initialized to 0

P1 :

$S_1 ;$

signal (synch) ;

P2 :

wait (synch) ;

$S_2 ;$

Semaphore - Control access to resources

- Counting semaphores can be used to control access to a given resource consisting of a finite number of instances.
- The semaphore is initialized to the number of resources available.
- Each process that wishes to use a resource performs a **wait()** operation on the semaphore (thereby decrementing the count).
- When a process releases a resource, it performs a **signal()** operation (incrementing the count).
- When the count for the semaphore goes to 0, all resources are being used. After that, processes that wish to use a resource will block until the count becomes greater than 0.

Semaphores – No busy waiting

- Semaphores described so far suffer from *busy waiting*.
- **Solution:**
 - The process waiting on a semaphore is blocked, and removed from execution on the CPU
 - Later awakened when the semaphore is available.
 - As a result, each semaphore needs to maintain a list of processes that are blocked and waiting for it.
 - When the semaphore becomes available, one of the processes can be woken up and scheduled to execute on the CPU.

Semaphore Implementation with no Busy waiting

- With each semaphore there is an associated waiting queue
- Each semaphore has two data items:
 - value (of type integer)
 - pointer to the list of processes waiting
- `wait()` and `signal()` operations are atomic.
- Two additional operations: `block()` and `wakeup()`

- **Semaphore Structure**

```
typedef struct{  
  
    int value;  
  
    struct process *list;  
  
} semaphore;
```

Implementation with no Busy waiting (Cont.)

Wait Operation

```
wait(semaphore *S) {  
    S->value--;  
    if (S->value < 0) {  
        add this process to S->list;  
        block();  
    }  
}
```

block – places the process invoking the operation on the appropriate waiting queue



Signal Operation

```
signal(semaphore *S) {  
    S->value++;  
    if (S->value <= 0) {  
        remove a process P from S->list;  
        wakeup(P);  
    }  
}
```

wakeup – removes a processes in the waiting queue and places it in the ready queue



Semaphore Implementation with no Busy waiting

- Note that in this implementation, semaphore values may be negative
- If a semaphore value is negative, its magnitude is the number of processes waiting on that semaphore.
- Implementation of semaphores using waiting queues may result in **Deadlock and/or starvation**.

Deadlock and Starvation

- **Deadlock** – two or more processes are waiting indefinitely for an event that can be caused by only one of the waiting processes
- Let S and Q be two semaphores initialized to 1

P_0	P_1
<code>wait(S);</code>	<code>wait(Q);</code>
<code>wait(Q);</code>	<code>wait(S);</code>
<code>...</code>	<code>...</code>
<code>signal(S);</code>	<code>signal(Q);</code>
<code>signal(Q);</code>	<code>signal(S);</code>

- **Starvation – indefinite blocking**
 - A process may never be removed from the semaphore queue in which it is suspended

Semaphores in C

- Semaphores are not part of Pthread standards and instead they belong to the POSIX SEM extension.
- To use semaphores in C on a Linux machine include 'semaphore.h' header file:
 - `#include <semaphore.h>`
 - `sem_t` data type for semaphores.
 - semaphores created with `sem_init()` function.
 - `sem_init()` - Takes 3 arguments:
 1. Pointer to the semaphore
 2. Flag indicating the level of sharing
 3. The semaphore's initial value
 - `sem_wait()` ;
 - `sem_post()` ;

Semaphores in C

```
/*Declaring Semaphore*/
```

```
Sem_t sem;
```

```
/*Initialize Semaphore*/
```

```
if (sem_init(&sem, 0, n) !=0) {
```

```
    printf("Error in initializing empty semaphore \n"
```

```
}
```

```
/* acquire the semaphore */
```

```
sem_wait(&sem);
```

```
/* critical section */
```

```
/* release the semaphore */
```

```
sem_post(&sem);
```

sem_init arguments

1. pointer to the semaphore
2. flag indicating the level of sharing
3. The semaphore's initial value

Returns 0 when semaphore created with no errors, otherwise returns non zero value.

Classical Problems of Synchronization

- Bounded-Buffer Problem (Producer - Consumer Problem) (discussed in part I)
- Readers and Writers Problem
- Dining-Philosophers Problem

Bounded-Buffer Problem

- n buffers, each can hold one item
- Semaphore **mutex** initialized to the value 1
 - `mutex = 1` => indicates that `mutex` lock is available
 - `mutex = 0` => indicates that `mutex` lock is unavailable
- Semaphore **full** initialized to the value 0, as initially the number of filled slots in the buffers is zero.
- Semaphore **empty** initialized to the value n , as initially the number of empty slots in the buffer is n .

Bounded Buffer Problem (Cont.)

- The structure of the producer process

```
do {  
    ...  
    /* produce an item in next_produced */  
    ...  
    wait(empty);  
    wait(mutex);  
    ...  
    /* add next produced to the buffer */  
    ...  
    signal(mutex);  
    signal(full);  
} while (true);
```


Bounded Buffer Problem (Cont.)

- The structure of the consumer process

```
Do {  
  
    wait(full);  
  
    wait(mutex);  
  
    ...  
  
    /* remove an item from buffer to next_consumed */  
  
    ...  
  
    signal(mutex);  
  
    signal(empty);  
  
    ...  
  
    /* consume the item in next consumed */  
  
    ...  
  
} while (true);
```


Dining-Philosophers Problem

- Philosophers spend their lives alternating between **thinking** and **eating**
- Don't interact with their neighbors, occasionally try to pick up 2 chopsticks (one at a time) to eat from a bowl
 - Need both chopsticks to eat, then release both when done
- In the case of 5 philosophers
 - Shared data
 - Bowl of rice (data set) – assumed infinite
 - Semaphore **chopstick [5]** initialized to 1 (indicates its available)



Dining-Philosophers Problem Algorithm

- The structure of Philosopher i :

```
do {  
    wait (chopstick[i] );  
    wait (chopstick[ (i + 1) % 5])  
  
    // eat  
  
    signal (chopstick[i] );  
    signal (chopstick[ (i + 1) % 5] );  
  
    // think  
} while (TRUE);
```



- What is the problem with this algorithm?

Dining-Philosophers Problem Algorithm

- Consider the following scenario in which **all** the philosophers are hungry and try to eat at the same time.
- Each philosopher grabs the chopstick to their left right (`chopstick[i]`).
- Each semaphore `chopstick[i]=0`, where $1 \leq i \leq 5$.
- When each philosopher tries to grab her right left chopstick (`chopstick[(i+1) % 5]`), she will be delayed forever. Thus resulting in a ***deadlock***.

Dining-Philosophers Problem Algorithm (Cont.)

- Deadlock can be handled by
 - Allowing at most 4 philosophers to be sitting simultaneously at the table.
 - Allow a philosopher to pick up the forks only if both are available (picking must be done in a critical section).
 - Use an asymmetric solution -- an odd-numbered philosopher picks up first the left chopstick and then the right chopstick. Even-numbered philosopher picks up first the right chopstick and then the left chopstick.

Problems with Semaphores

- Incorrect use of semaphore operations could result in mutual exclusion being violated, deadlock and/or starvation.
- Consider the following situation:
 - Processes share a Semaphore `mutex`, which is initialized to 1.
- Suppose that a process interchanges the order in which the `wait()` and `signal()` operations on the semaphore `mutex` are executed. That is, it executes

```
signal(mutex);  
  
...  
  
critical section  
  
...  
  
wait(mutex);
```

- **What is the problem with the above code?**

Problems with Semaphores

- Suppose that a process replaces `signal(mutex)` with `wait(mutex)`.

That is, it executes

```
wait(mutex);
```

```
...
```

```
critical section
```

```
...
```

```
wait(mutex);
```

What is the problem here?

- Omitting of `wait(mutex)` or `signal(mutex)` (or both)
 - Deadlock and starvation are possible.

Monitors

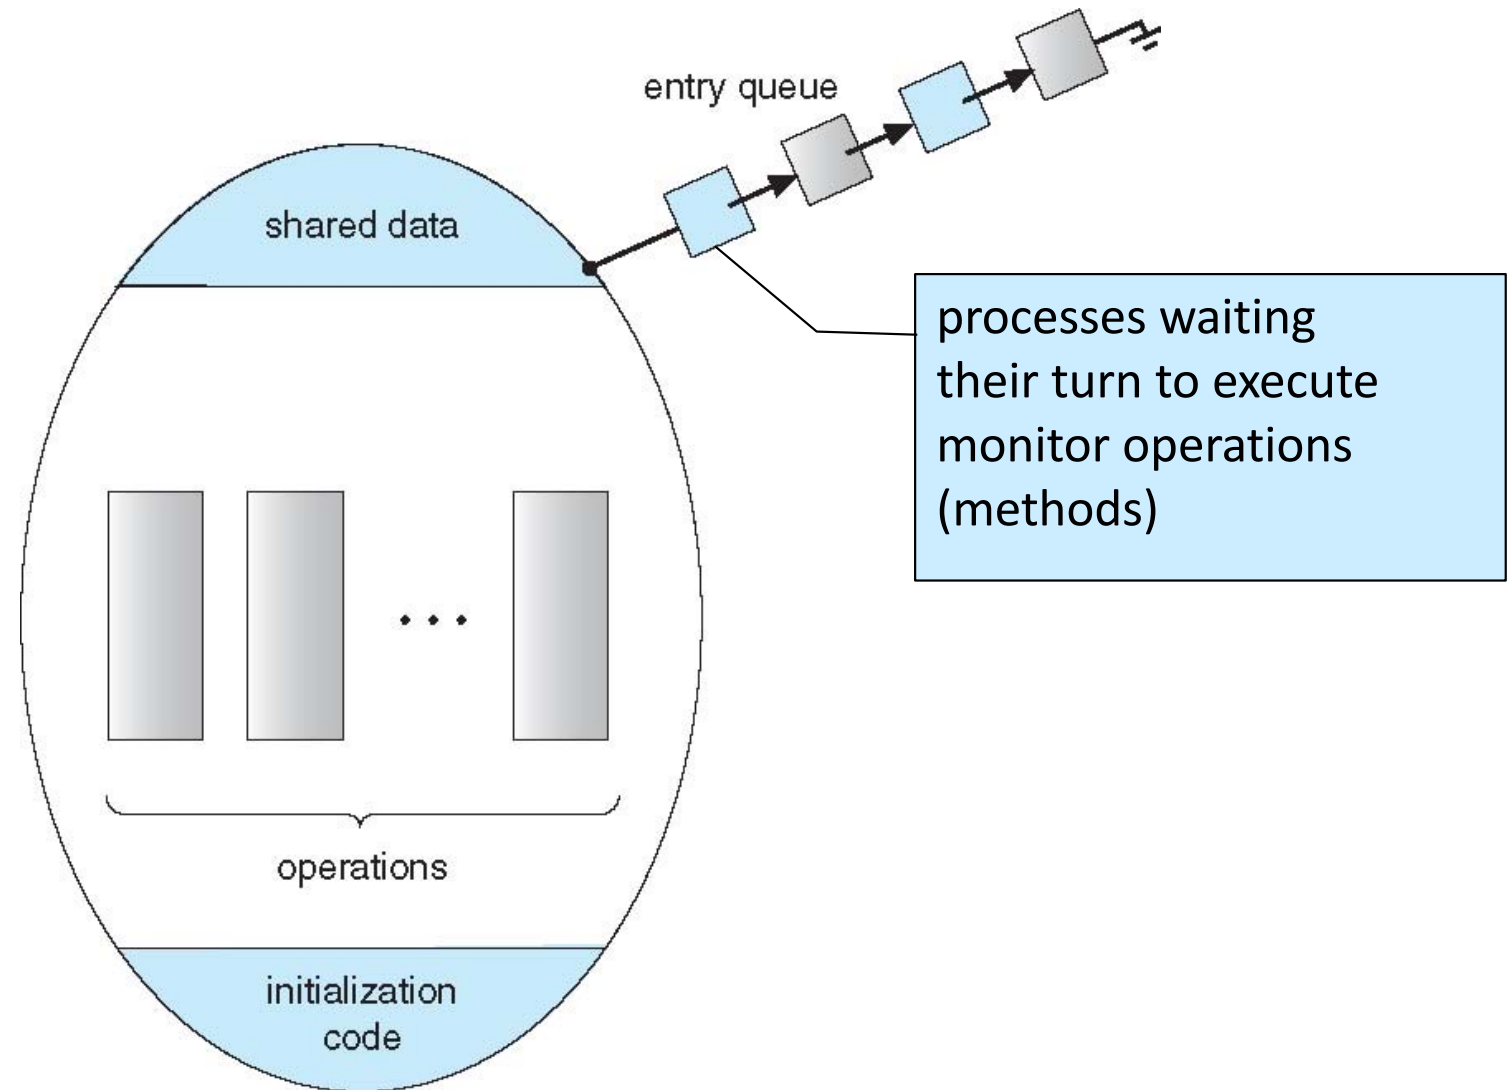
- A high-level abstraction that provides a convenient and effective mechanism for process synchronization.
- A monitor ADT has
 - Private data members
 - Public functions that are implicitly executed with mutual exclusion.
 - For each monitor instance there exists a mutual exclusion lock.
 - To enter the monitor, a process acquires the mutual exclusion lock
 - While exiting the monitor, a process releases the lock, and therefore the monitor, for other threads.
 - In addition, monitors may define wait conditions (condition variables) that can be used inside the monitor to synchronize the member functions.

Syntax of Monitors

```
monitor monitor-name
{
    // shared variable declarations
    procedure P1 (...) { ... }
    .
    .
    .
    procedure Pn (...) {.....}

    Initialization code (...) { ... }
}
}
```


Schematic view of a Monitor



Monitor Example

```
monitor BankAccount {  
  
    void BankAccount_init() {  
        amount = 0;  
    }  
    void withdraw(int value) {  
        amount = amount - value;  
    }  
    void deposit(int value) {  
        amount = amount + value;  
    }  
    private: int amount;  
};
```

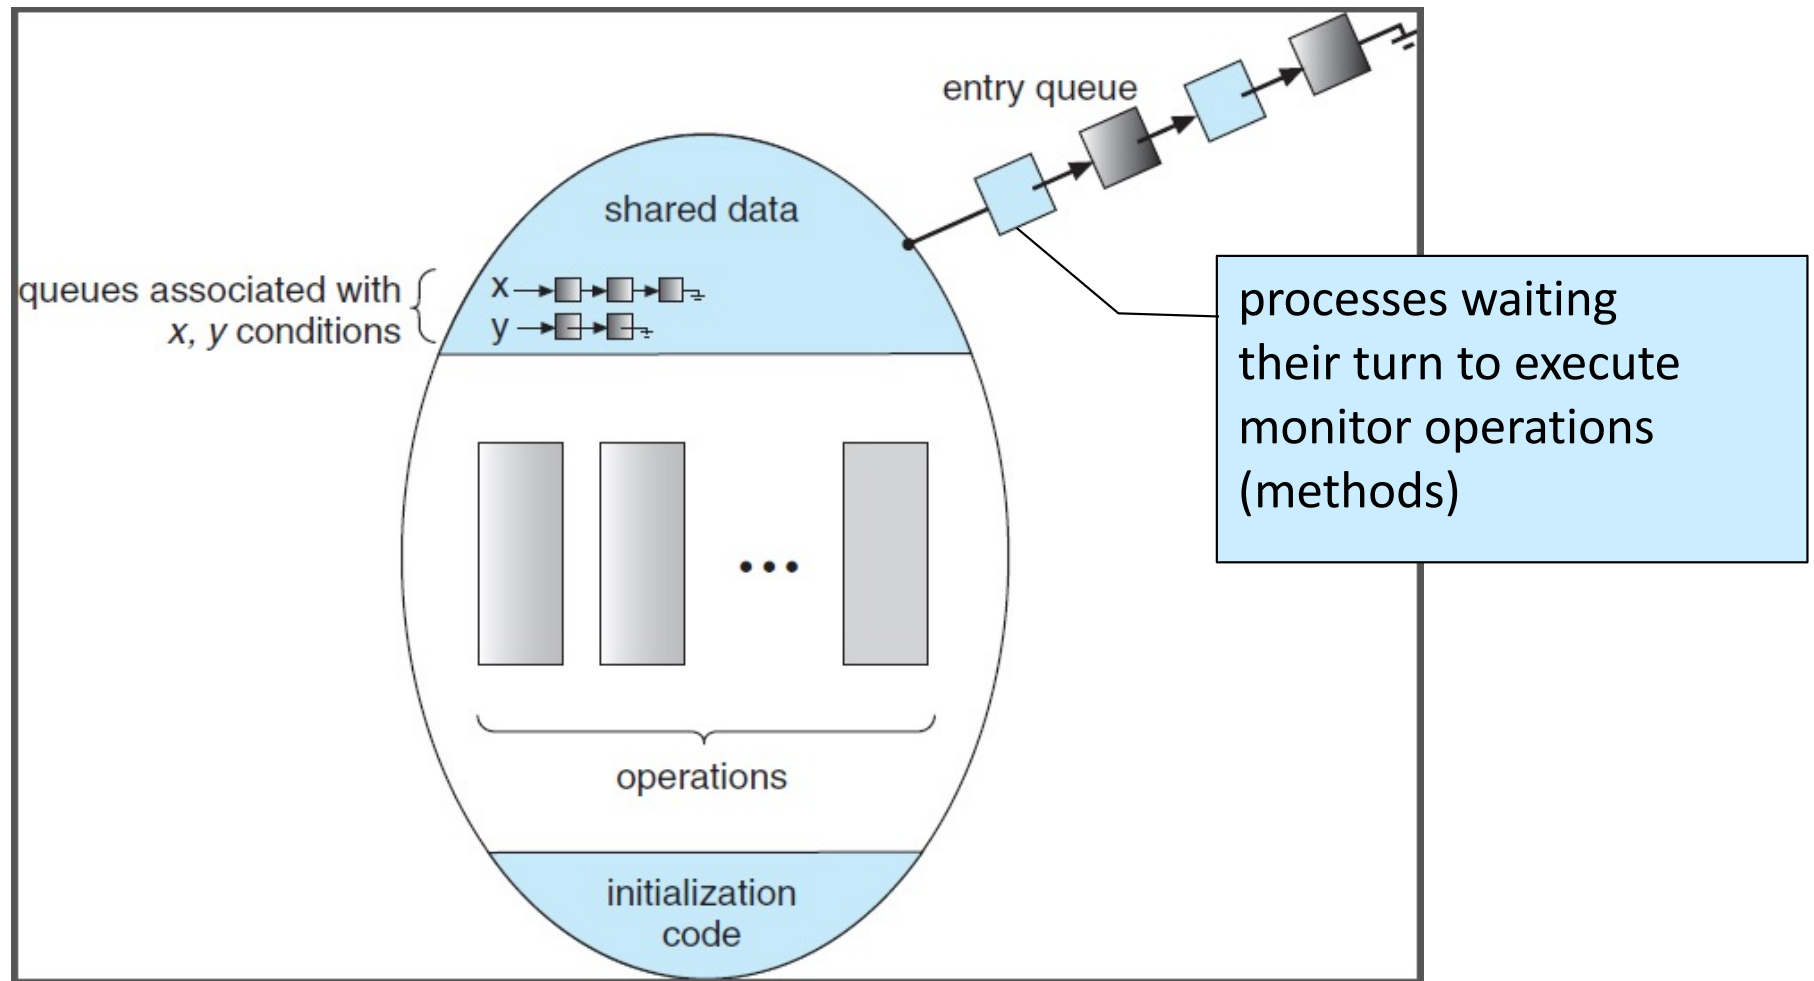
Monitor Example cont...

- How to ensure the following two conditions are met:
 - Deposits do not take place if amount ≥ 500
 - Withdrawals do not take place if amount ≤ 0

Condition Variables

- Condition variables defined as: `condition x, y;`
- Condition variables have two legal atomic operations: `wait()` and `signal()`
- If `x` is a condition variable then the two operations can be accessed as: `x.wait()`, `x.signal()`
- `wait()` – blocks a process (until some other process calls `signal` on it), and adds it to a queue associated with that condition variable.
- `signal()` – wakes up exactly one process from the condition variable's queue of waiting processes. If no processes waiting then does nothing.

Monitor with Condition Variables



Monitor Example with condition variables

```
monitor BankAccount {  
  
    void BankAccount_init() {  
        amount = 0;  
    }  
  
    void withdraw(int value) {  
        if (amount <= 0) { c1.wait(); }  
        amount = amount - value;  
        c2.signal();  
    }  
  
    void deposit(int value) {  
        if (amount >= 500) { c2.wait(); }  
        amount = amount + value;  
        c1.signal();  
    }  
  
    int amount;  
    condition c1, c2;  
};
```

In this example we assume that the value deposited/withdrawn is the same, similar to the problem discussed in Lab 6

Condition Variables Choices

- If process P invokes `x.signal()`, and process Q is suspended in `x.wait()`, what should happen next?
 - Both Q and P **cannot** execute in parallel. If Q is resumed, then P must wait
- Options include
 - **Signal and wait** – P waits until Q either leaves the monitor or it waits for another condition
 - **Signal and continue** – Q waits until P either leaves the monitor or it waits for another condition
 - Both have pros and cons – language implementer can decide

Resuming Processes within a Monitor

- If several processes queued on condition `x`, and `x.signal()` executed, which should be resumed?
 - FCFS (first come first serve) frequently not adequate
- **conditional-wait** construct of the form `x.wait(c)`
 - where `c` is **priority number**
 - process with ***lowest number*** (highest priority) is scheduled next

Monitor Solution to Dining Philosophers

```
monitor DiningPhilosophers
{
    enum { THINKING; HUNGRY, EATING}
    state [5] ;
    condition self [5];

    void pickup (int i) {
        state[i] = HUNGRY;
        test(i);
        if (state[i] != EATING) self[i].wait;
    }

    void putdown (int i) {
        state[i] = THINKING;
        // test left and right neighbors
        test((i + 4) % 5);
        test((i + 1) % 5);
    }
}
```

Solution to Dining Philosophers (Cont.)

```
void test (int i) {
    if ((state[(i + 4) % 5] != EATING) &&
        (state[i] == HUNGRY) &&
        (state[(i + 1) % 5] != EATING) ) {
        state[i] = EATING ;
        self[i].signal () ;
    }
}

initialization_code() {
    for (int i = 0; i < 5; i++)
        state[i] = THINKING;
}
}
```

Solution to Dining Philosophers (Cont.)

- Each philosopher i invokes the operations `pickup()` and `putdown()` in the following sequence:

`DiningPhilosophers.pickup(i) ;`

EAT

`DiningPhilosophers.putdown(i) ;`

- No deadlock, but starvation is possible