

# Operating Systems: Deadlocks

Neerja Mhaskar

Department of Computing and Software, McMaster University, Canada

**Acknowledgements:** Material based on the textbook Operating Systems Concepts (Chapters 5)

# System Model

- System consists of resources
- Resource types  $R_1, R_2, \dots, R_m$ 
  - *CPU cycles, memory space, I/O devices*
- Each resource type  $R_i$  has  $W_i$  instances.
  - Eg: You can have two instances of a printer available
- Each process utilizes a resource as follows:
  - **request**
  - **use**
  - **Release**
- if the resources not available, the process enters a *waiting state*.

# Deadlock

- **Deadlock** – when a waiting process has requested a resource held by other waiting processes.
- In a deadlock processes never finish executing
  - System resources are tied up.
  - Thus, preventing other processes from starting.


# Deadlock Characterization

Deadlock  $\Leftrightarrow$  below **four conditions hold** at the same time.

- **Mutual exclusion:** only one process at a time can use a resource.
- **Hold and wait:** a process holding at least one resource is **waiting to acquire additional resources** held by other processes.
- **No preemption:** a resource can be **released only voluntarily** by the process holding it, after that process finishes its task.

# Deadlock Characterization contd..


- **Circular wait:** there exists. a set  $\{P_0, P_1, \dots, P_n\}$  of **waiting processes** such that  $P_0$  is waiting for a resource that is held by  $P_1$ ,  $P_1$  is waiting for a resource that is held by  $P_2$ , ...,  $P_{n-1}$  is waiting for a resource that is held by  $P_n$ , and  $P_n$  is waiting for a resource that is held by  $P_0$ .



circular-wait condition  $\Rightarrow$  the hold-and-wait condition.


# Resource-Allocation Graph

- **System resource-allocation graph** (directed graph  $G = (V, E)$ ) used to describe deadlocks.
- $V$  is partitioned into two sets:
  - $P = \{P_1, P_2, \dots, P_n\}$ , the set consisting of all the processes in the system
  - $R = \{R_1, R_2, \dots, R_m\}$ , the set consisting of all resource types in the system
- $E$  has two types of edges:
  - **Request edge** – directed edge  $P_i \rightarrow R_j$ : Process  $P_i$  requested an instance of resource  $R_j$  and is waiting.
  - **Assignment edge** – directed edge  $R_j \rightarrow P_i$ : Resource  $R_j$  assigned to process  $P_i$




# Resource Allocation Graph Examples

Resource allocation graph with a cycle but no Deadlock



Resource allocation graph with a cycle and a Deadlock



Cycles with circular wait condition satisfied

- $P_1 \rightarrow R_1 \rightarrow P_2 \rightarrow R_3 \rightarrow P_3 \rightarrow R_2 \rightarrow P_1$
- $P_2 \rightarrow R_3 \rightarrow P_3 \rightarrow R_2 \rightarrow P_2$

# Basic Facts

- If a graph contains no cycles  $\Rightarrow$  no deadlock
- If graph contains a cycle  $\Rightarrow$ 
  - if only one instance per resource type, then deadlock
  - if several instances per resource type, possibility of deadlock



# Methods for Handling Deadlocks

Ensure that the system will **never** enter a deadlock state:

- **Deadlock prevention:** ensure that atleast one of the necessary conditions (stated in slide #4) cannot hold.
  - Only the **Circular Wait condition is a practical** option for deadlock prevention.
- **Deadlock avoidance:** requires that OS be given **additional information in advance** about the type and no. of resources a process will request. Based on this data OS **decides the waiting strategy for the process.**
- **Deadlock detection:** Allow the system to enter a deadlock state and then recover.
- **IGNORE!** the problem and pretend that deadlocks never occur; used by most operating systems such as UNIX, Linux, and Windows!

# Deadlock Avoidance Strategy

- Requires that the OS be given **additional information in advance** about the type and no. of resources a process will request.
- Each process declares its ***maximum need = number of resources*** of each type that the process may need.
- The deadlock-avoidance algorithm dynamically examines the ***resource-allocation state*** to ensure that there can never be a **circular-wait condition**.

# Resource allocation state

- **Resource-allocation state** is defined by
  - Number of available resources
  - Number of allocated resources
  - the maximum need of all the processes in the system.

**Example:** Consider a system with a single resource type – magnetic tapes= 12.

Resource allocation state of the system at time  $T_0$ :

	<u>Maximum Need</u>	<u>Allocated resources</u>	<u>Available</u>
P0	10	5	12-9 = 3
P1	4	2	
P2	9	2	

# Safe State

- System is in **safe state** if there exists a sequence  $\langle P_1, P_2, \dots, P_n \rangle$  of ALL the processes in the systems such that for each  $P_i$ :
  - Resources  $P_i$  still needs  $\leq$  total available resources + resources held by/allocated to all  $P_1, P_2, \dots, P_{i-1}$ .
- To assess whether the system is in a safe state, we compute the needs of a process  $P_i$  as below

***$P_i$ 's Need = Max. Need of Process  $P_i$  – Allocated resources for process  $P_i$***

# Safe State Example

**Is the below system at time  $t_0$  in safe state?**

- Consider a system with 12 magnetic tapes.
- **Resource allocation state** of the system at time  $t_0$ :

	<u>Maximum Needs</u>	<u>Allocated resources</u>	<u>Available</u>
$P_0$	10	5	12-9 = 3
$P_1$	4	2	
$P_2$	9	2	

# Safe State Example – with Need vector

- *$P_i$ 's Need = Max. Need of Process  $P_i$  – Allocated resources for process  $P_i$*
- **Resource allocation state** of the system at time  $t_0$  with need vector:


	<u>Maximum Needs</u>	<u>Allocated resources</u>	<u>Need</u>	<u>Available</u>
$P_0$	10	5	5	3
$P_1$	4	2	2	
$P_2$	9	2	7	

Does there exists a sequence  $\langle P_1, P_2, \dots, P_n \rangle$  of ALL the processes in the systems such that for each  $P_i$  satisfying the state state requirement?

YES! The sequence is  $\langle P_1, P_0, P_2 \rangle$

# Basic Facts

- If a system is in safe state  $\Rightarrow$  no deadlocks
- If a system is in unsafe state  $\Rightarrow$  possibility of deadlock
- Avoidance  $\Rightarrow$  ensure that a system will never enter an unsafe state.



# Deadlock Avoidance algorithm Outline

- Given the resource allocation state of a system,
  - The algorithm checks if the system is in a *safe state*.
  - If the system is in a safe state, whenever a process requests an instance of a resource type,
    - It check to see if allocating the resources continues to have the system in a safe state.
      - If yes -- allocate the resources.
      - If no -- have the process wait.




# Deadlock Avoidance Algorithms

- Single instance of a resource type
  - Use a **resource-allocation graph**
- Multiple instances of a resource type
  - Use the **Banker's algorithm**


# Resource-Allocation Graph Scheme

- Add claim edges to existing resource allocation graph.
- **Claim edge**  $P_i \rightarrow R_j$  indicates that process  $P_i$  may request resource  $R_j$ ; *represented by a dashed line*
- Claim edge converts to request edge when a process requests a resource
- Request edge converted to an assignment edge when the resource is allocated to the process
- When a **resource is released by a process**, assignment edge reconverts to a claim edge
- Resources must be claimed *a priori* in the system



# Resource-Allocation Graph Algorithm

- Suppose that process  $P_i$  requests a resource  $R_j$
- The request can be granted **only if** converting the request edge to an assignment edge **does not result in the formation of a cycle** in the resource allocation graph
- We check for safety by using a cycle-detection algorithm.



## Unsafe State In Resource-Allocation Graph

Although  $R_2$  is free, we cannot allocate it to  $P_2$  since this will create a cycle!

## Banker's Algorithm Outline for *multiple instances* of a resource type

- The algorithm consists of two parts
  - **PART 1 - Safety Algorithm** – checks whether a system is in a safe state or not.
  - **PART 2 - Resource-Request Algorithm** – checks to see if resources requested by a process can be satisfied or not.
- Each process must a priori claim maximum use
- When a process gets all its resources it must return them in a *finite amount of time*
- When a process requests a resource, it *may have to wait*

# Data Structures for the Banker's Algorithm

Let  $n$  = number of processes, and  $m$  = number of resources types.

- **Available**: Vector of length  $m$ . If  $available[j] = k$ , there are  $k$  instances of resource type  $R_j$  available
- **Max**:  $n \times m$  matrix. If  $Max[i,j] = k$ , then process  $P_i$  may request at most  $k$  instances of resource type  $R_j$
- **Allocation**:  $n \times m$  matrix. If  $Allocation[i,j] = k$  then  $P_i$  is currently allocated  $k$  instances of  $R_j$
- **Need**:  $n \times m$  matrix. If  $Need[i,j] = k$ , then  $P_i$  may need  $k$  more instances of  $R_j$  to complete its task

$$Need[i,j] = Max[i,j] - Allocation[i,j]$$

## Example of Data Structures for Banker's Algorithm

- 5 processes  $P_0$  through  $P_4$ ;

3 resource types:

$A$  (10 instances),  $B$  (5 instances), and  $C$  (7 instances)

**Resource allocation state at time  $T_0$ :**

	Max. Need			Allocation			Available		
	A	B	C	A	B	C	A	B	C
$P_0$	7	5	3	0	1	0	3	3	2
$P_1$	3	2	2	2	0	0			
$P_2$	9	0	2	3	0	2			
$P_3$	2	2	2	2	1	1			
$P_4$	4	3	3	0	0	2			

# Need matrix for Banker's Algorithm

The content of the matrix **Need** is defined to be **Max. Need – Allocation**

- 5 processes  $P_0$  through  $P_4$ ;

3 resource types:

$A$  (10 instances),  $B$  (5 instances), and  $C$  (7 instances)

**Snapshot at time  $T_0$ :**

	Max. Need			Allocation			Need			Available		
	A	B	C	A	B	C	A	B	C	A	B	C
$P_0$	7	5	3	0	1	0	7	4	3	3	3	2
$P_1$	3	2	2	2	0	0	1	2	2			
$P_2$	9	0	2	3	0	2	6	0	0			
$P_3$	2	2	2	2	1	1	0	1	1			
$P_4$	4	3	3	0	0	2	4	3	1			

## Part 1 - Safety Algorithm outline explained with an example

**Step 1:** Maintain **Work** and **Finish** vectors of length  $m$  and  $n$ , respectively.

Initially:

**Work** = **Available** = (3 3 2)

**Finish** [ $i$ ] = *false* for  $i = 0, 1, \dots, n-1$  =

False	False	False	False	False
-------	-------	-------	-------	-------

**Step 2:** Find a process  $P_i$  such that both:

(a) **Finish** [ $i$ ] = *false*

(b) **Need** $_i \leq$  **Work**

If no such  $i$  exists, go to step 4

➤  $P_1$  satisfies conditions (a) and (b).



## Part 1 - Safety Algorithm outline explained with an example – contd...

**Step 3:**

**$Work = Work + Allocation_i$**

**$Finish[i] = true$**

$Work = Work + Allocation_1 = (3\ 3\ 2) + (2\ 0\ 0) = (5\ 3\ 2)$

$Finish[1] = true$  and  $Finish =$

False	True	False	False	False
-------	------	-------	-------	-------

go to step 2

Next, we see that process  $P_3$ ,  $P_4$ ,  $P_2$ , and  $P_0$  all satisfy the conditions in step 2.

**Step 4:** If  **$Finish[i] == true$**  for all  $i$ , then the system is in a safe state

**$Finish =$**

True	True	True	True	True
------	------	------	------	------

Therefore, the system is in a safe state and the sequence of processes satisfying the safety requirement is -

**$\langle P_1, P_3, P_4, P_2, P_0 \rangle$**

## Part – 2 Resource-Request Algorithm for Process $P_i$

- **$Request_i$**  = request vector for process  $P_i$ . If  **$Request_i[j] = k$**  then process  $P_i$  wants  $k$  instances of resource type  $R_j$
- If  **$Request_i \leq Need_i$**  go to step 2. *Otherwise, raise error condition*, since process has exceeded its maximum claim
- If  **$Request_i \leq Available$** , go to step 3. *Otherwise,  $P_i$  must wait*, since resources are not available
- Pretend to allocate requested resources to  $P_i$  and update the system state as follows:

$$Available = Available - Request_i;$$

$$Allocation_i = Allocation_i + Request_i;$$

$$Need_i = Need_i - Request_i;$$

- If safe  $\Rightarrow$  the resources are allocated to  $P_i$
- If unsafe  $\Rightarrow P_i$  must wait, and the old resource-allocation state is restored

## PART 2 - Resource-Request Algorithm Explained with an Example

- **P<sub>1</sub> requests resources (1 0 2)**
- Check if **Request<sub>1</sub> ≤ Need<sub>1</sub>**
  - $(1\ 0\ 2) \leq (1\ 2\ 2) \Rightarrow \text{true}$
- Check if **Request<sub>1</sub> ≤ Available**
  - $(1\ 0\ 2) \leq (3\ 2\ 2) \Rightarrow \text{true}$
- Pretend that resources requested have been granted.
- Update system state. Max need, Allocation<sub>1</sub> and Need<sub>1</sub> data structures
  - $\text{Available}_1 = (3\ 2\ 2) - (1\ 0\ 2) = (2\ 2\ 0)$
  - $\text{Allocation}_1 = (2\ 0\ 0) + (1\ 0\ 2) = (3\ 0\ 2)$
  - $\text{Need}_1 = (1\ 2\ 2) - (1\ 0\ 2) = (0\ 2\ 0)$

## PART 2 - Resource-Request Algorithm Explained with an Example contd...

### Updated resource allocation state:

	Max. Need			Allocation			Need			Available		
	A	B	C	A	B	C	A	B	C	A	B	C
P <sub>0</sub>	7	5	3	0	1	0	7	4	3	2	2	0
P <sub>1</sub>	3	2	2	3	0	2	0	2	0			
P <sub>2</sub>	9	0	2	3	0	2	6	0	0			
P <sub>3</sub>	2	2	2	2	1	1	0	1	1			
P <sub>4</sub>	4	3	3	0	0	2	4	3	1			

## PART 2 - Resource-Request Algorithm Explained with an Example contd...

- Run safety algorithm on the **updated resource allocation state**.
- System is in safe state and the sequence of processes satisfying the safety requirement is  $\langle P_1, P_3, P_4, P_2, P_0 \rangle$

### Updated Resource allocation state after request has been granted for $P_1$

	Max. Need			Allocation			Need			Available		
	A	B	C	A	B	C	A	B	C	A	B	C
$P_0$	7	5	3	0	1	0	7	4	3	2	2	0
$P_1$	3	2	2	3	0	2	0	2	0			
$P_2$	9	0	2	3	0	2	6	0	0			
$P_3$	2	2	2	2	1	1	0	1	1			
$P_4$	4	3	3	0	0	2	4	3	1			

## Resource-Request Algorithm Example – Cont...

- When system in this state, can request for (3 3 0) by  $P_4$  be granted?
  - Check if **Request<sub>4</sub> ≤ Available**
    - $(3\ 3\ 0) \leq (2\ 2\ 0) \Rightarrow \text{false}$
    - The request cannot be granted.
- When system in this state, can request for (0 2 0) by  $P_0$  be granted?
  - Check if **Request<sub>0</sub> ≤ Available**
    - $(0\ 2\ 0) \leq (2\ 2\ 0) \Rightarrow \text{true}$
  - Check if **Request<sub>0</sub> ≤ Need<sub>0</sub>**
    - $(0\ 2\ 0) \leq (7\ 4\ 3) \Rightarrow \text{true}$
  - Pretend to grant the resources requested

## Resource-Request Algorithm Example – Cont...

### Updated Resource allocation state

	Max. Need			Allocation			Need			Available		
	A	B	C	A	B	C	A	B	C	A	B	C
P <sub>0</sub>	7	5	3	0	3	0	7	2	3	2	0	0
P <sub>1</sub>	3	2	2	3	0	2	0	2	0			
P <sub>2</sub>	9	0	2	3	0	2	6	0	0			
P <sub>3</sub>	2	2	2	2	1	1	0	1	1			
P <sub>4</sub>	4	3	3	0	0	2	4	3	1			

- However, since no sequence of processes exist satisfying the safe state requirement, this request cannot be granted as doing so will leave the system in an unsafe state.


# Deadlock Detection

- Allow system to enter deadlock state
- Use deadlock detection algorithm to check if a deadlock exists
- If deadlock exists, use a recovery scheme to recover from the deadlock




# Deadlock Detection - Single Instance of Each Resource Type

- A variant of the resource-allocation graph if all resources have only **a single instance** - used for deadlock detection
- Nodes are processes, and an edge  $P_i \rightarrow P_j$  in the wait for graph implies that  $P_i$  is waiting for  $P_j$  to release a resource that  $P_i$  needs.




Resource-Allocation Graph




Corresponding wait-for graph

# Deadlock Detection - Single Instance of Each Resource Type

***If a cycle exists in the wait-for graph, then the system is in deadlock.***



Resource-Allocation Graph



Corresponding wait-for graph

# Deadlock Detection Algorithm for Multiple Instances of a Resource Type Outline

- The algorithm needs to know
  - The number of *available resources* for each resource type.
  - The number of *allocated resources* for each resource type.
  - The number of *requested resources* by all processes in the system.
- Given the above,
  - The deadlock detection algorithm checks whether the system is *in a deadlocked* state or not.
  - If the system is in a deadlocked state, then the algorithm also identifies the processes involved in the deadlock.

# Deadlock Detection Algorithm

1. Let **Work** and **Finish** be vectors of length **m** and **n**, respectively

Initialize:

(a) **Work = Available**

(b) For  $i = 1, 2, \dots, n$ ,

i. if **Allocation<sub>i</sub> ≠ 0**, then **Finish[i] = false**;

ii. otherwise, **Finish[i] = true**

2. Find an index **i** such that both:

(a) **Finish[i] == false**

(b) **Request<sub>i</sub> ≤ Work**

If no such **i** exists, go to step 4

*Items in red highlight the differences in Deadlock detection algorithm and the safety algorithm described under Banker's algorithm.*

# Detection Algorithm (Cont.)

3.  **$Work = Work + Allocation_i$**   
 **$Finish[i] = true$**   
go to step 2
4. If  **$Finish[i] == false$** , for some  $i$ ,  $1 \leq i \leq n$ , then the system is in deadlock state. **Moreover, if  $Finish[i] == false$ , then  $P_i$  is deadlocked.**

# Example for Deadlock Detection Algorithm

- Five processes  $P_0$  through  $P_4$
- Three resource types  
A (7 instances), B (2 instances), and C (6 instances)
- Snapshot of the system at time  $T_0$ :

	<b>Allocation</b>	<b>Request</b>	<b>Available</b>
	A B C	A B C	A B C
$P_0$	0 1 0	0 0 0	0 0 0
$P_1$	2 0 0	2 0 2	
$P_2$	3 0 3	0 0 0	
$P_3$	2 1 1	1 0 0	
$P_4$	0 0 2	0 0 2	

# Example of Detection Algorithm Cont...

## Step 1:

1. Work = Available = (0 0 0)

2. Finish =

False	False	False	False	False
-------	-------	-------	-------	-------

## Step 2: Find an index $i$ such that both:

1. Finish[i] == false

2. Request <sub>$i$</sub>  ≤ Work

If no such  $i$  exists, go to step 4

**P<sub>0</sub> satisfies the above two conditions.**

## Example of Detection Algorithm Cont...

- Step 3:

- $\mathbf{Work} = \mathbf{Work} + \mathbf{Allocation}_0 = (0\ 0\ 0) + (0\ 1\ 0) = (0\ 1\ 0)$

- $\mathbf{Finish}[1] = \mathbf{true}$

- go to step 2

- We see that process  $P_2$ ,  $P_3$ ,  $P_1$ , and  $P_4$  all satisfy the conditions in step 2.

- Finally, in **Step 4: Finish =**

True	True	True	True	True
------	------	------	------	------

- Therefore, the system is ***not in a deadlocked state***, as the following sequence of processes results in all values of the **Finish** vector to be **True**:

**$\langle P_0, P_2, P_3, P_1, P_4 \rangle$**



# Example of Detection Algorithm Cont...

- Suppose  $P_2$  requests an additional instance of type C.
- Then below is the updated snapshot of the system including this request:

	Allocation	Request	Available
	A B C	A B C	A B C
$P_0$	0 1 0	0 0 0	0 0 0
$P_1$	2 0 0	2 0 2	
$P_2$	3 0 3	0 0 1	
$P_3$	2 1 1	1 0 0	
$P_4$	0 0 2	0 0 2	

- Can reclaim resources held by process  $P_0$ , but insufficient resources to fulfill other processes' requests
- **Deadlock exists, consisting of processes  $P_1$ ,  $P_2$ ,  $P_3$ , and  $P_4$**

# Recovery from Deadlock

## ■ Process Termination

- Abort all deadlocked processes
- Abort one process at a time until the deadlock cycle is eliminated

## ■ Resource Preemption

- **Selecting a victim** – minimize cost
- **Rollback** – return to some safe state, restart process for that state
- **Starvation** – same process may always be picked as victim.

Possible solution - include number of rollback in cost factor