

# Operating Systems: Mass Storage Structure – Part II

Neerja Mhaskar

Department of Computing and Software, McMaster University, Canada

**Acknowledgements:** Material based on the textbook Operating Systems Concepts (Chapter 11)

# Array of Disks

- **Storage array**: Array of disks that operate independently and in parallel.
- **Advantages:**
  - **Separate I/O requests can be handled in parallel** as long as the data required reside on separate disks.
  - **A single I/O request can be executed in parallel** if the block of data to be accessed is distributed across multiple disks
  - Reliability can be achieved via data mirroring.
    - **Data mirroring**: Duplicating identical data on multiple disks
      - Mirroring provides high reliability, but it is expensive.
- **Disadvantage**: Use of multiple disks increases the probability of failure

# Image of storage array



# Error Detection and Error Correction Code (ECC)

- **Error detection** determines if a problem has occurred
  - Parity used to detect (single bit) errors in memory systems.
  - **Parity** records whether the number of bits in the byte set to 1 is even (parity = 0) or odd (parity = 1).
    - Parity (10011000) = 1
    - Parity (11011000) = 0
- An **error-correction code (ECC)** not only detects the problem, but also corrects it – achieved using algorithms and extra amounts of storage.

# Parity

- In the memory system, *each byte has a parity bit associated with it.*
  - Every byte of memory needs an extra bit of memory to store the parity.
  - If one of the bits in the byte is damaged (either a 1 becomes a 0, or a 0 becomes a 1), the parity of the byte changes and thus does not match the stored parity, and vice a versa.
  - A double-bit-error *might* go undetected however.
  - Parity calculated by performing an XOR (“eXclusive OR”) operation of the bits in a byte.

# Error correction code

## How does it work?

- When the controller writes data on a sector, the ECC is calculated from all the bytes in the data and written on the sector.
- When the sector is read, ECC is recalculated and compared with the stored value.
- If the stored and calculated numbers are different => data corruption.
- If only a few bits of data have been corrupted, ECC can correct the errors. Otherwise, reports data error.

# RAID – Redundant Array of Independent Disks

- **RAID** - disk-organization techniques used to improve performance and reliability in a system *using an array of disks*.
- The different RAID levels share the below characteristics:
  1. Each RAID level consists of a set of physical disk drives viewed by the OS as a ***single logical drive***.
  2. Data are distributed across the physical drives of an array in a scheme known as **striping**.
  3. **Redundant disk capacity** is used to store duplicate data or parity information
    - Guarantees data recoverability in case of a disk failure.

# Data Striping

- **Data striping** is of two types:
  - **Bit-level striping** - splitting the bits of each byte across multiple disks
    - With 8 disks, the  $i$ -th bit of a byte goes to disk  $i$ .
  - **Block-level striping** – stripping the blocks of a file across multiple disks.
    - With  $n$  disks, block  $i$  of a file goes to disk  $(i \bmod n) + 1$ .
    - For example if  $n=4$  and  $i=5$ , then block 5 goes to disk  $(5 \bmod 4) + 1 = 1 + 1 = 2$ . The assumption here is that Disk numbering starts from 1 and block numbering starts from 0.
    - Most common



# RAID Levels

- **RAID 0**: has *block level striping* with no redundancy.
- **Raid 1**: *has mirroring* only, no striping.
- In figure on the left
  - C = copy of data
  - P = Parity in RAID level 4, 5
  - P, Q = ECC



(a) RAID 0: non-redundant striping.



(b) RAID 1: mirrored disks.



(c) RAID 4: block-interleaved parity.



(d) RAID 5: block-interleaved distributed parity.



(e) RAID 6: P + Q redundancy.

# RAID 4

- **RAID 4** (block-interleaved parity organization):
  - Uses **block-level striping**, and in addition keeps *a parity block on a separate additional disk* for corresponding blocks from  $N$  other disks
  - Therefore, RAID 4, has a dedicated block for parity blocks.
  - RAID 4 is also known as memory-style error-correcting-code (ECC) organization. ECC is also used in RAID 5 and 6.



(c) RAID 4: block-interleaved parity.

## RAID 5 (block-interleaved distributed parity) (most common):

- Spreads data and parity among all  $N+1$  disks, rather than storing data in  $N$  disks and parity in one disk.



- For each block, one of the disks stores the parity and the others store data.
- A parity block cannot store parity for blocks in the same disk
- For example, with an array of five drives, the parity for the  $n$ th block is stored in drive  $(n \bmod 5) + 1$ .
- By spreading the parity across all the drives in the set, RAID 5 avoids potential overuse of a single parity drive.

# RAID 6

- **RAID 6 (P + Q redundancy scheme)** - Like RAID level 5 but stores extra redundant information to guard against multiple disk failures.
- error-correcting codes used are used to calculate
- In the below RAID 6 example, 2 blocks of redundant data are stored for every 4 blocks of data, as opposed to one just parity block in level 5. This enables the system to recover from two drive failures.



(e) RAID 6: P + Q redundancy.

# RAID 4, 5 and 6 Analysis

## Reads: For a single block

A block read accesses only one disk, allowing other requests to be processed by the other disks.

- Thus, the data-transfer rate for each access is slower, but multiple read accesses can proceed in parallel, leading to a higher overall I/O rate.

## Reads: For many blocks

The transfer rates for large reads are high, since all the disks can be read in parallel.

# RAID 4, 5 and 6 Analysis Cont...

Writes: smaller than a block – Require significantly more time, as the OS needs to do the following:

- First read the block to which data is to be written, and its corresponding parity block – involves **2 reads (2 disk accesses)**
- Modify the block with new data, and written back. Modify parity and write it back – **involves 2 writes (2 disk accesses)**
- This is known as the **read-modify-write cycle**.

Writes: many blocks

- Large writes have high transfer rates, since the data and parity can be written in parallel.