

# File Input and Output

PHYS2G03

© James Wadsley,  
McMaster University

# File I/O

- Modern science produces huge amounts of data
- e.g. Experiments: Genomics, Supercolliders, Astronomy, Social Science, Health ...
- e.g. Simulations: Climate, Astrophysics, Fluid Dynamics, Physics, Molecular Dynamics ...
- Data is stored in files

# Data Formats

Data in files can be

- 1) Human readable: plain text  
e.g. what `std::cout` makes
- 2) Binary/Raw data  
e.g. Same format as computer memory

# Text Files

```
#include <iostream>
```

provides cout and cin for terminals

Programs can generate plain text using `std::cout` and read from the keyboard using `std::cin`

```
cp -r /home/2G03/fileio ~/
```

```
cd fileio
```

```
gedit cout.cpp
```

# cout.cpp: print 10 random numbers

```
#include <iostream>
#include <stdlib.h>

int main() {

    int i;
    const int n=10;

    std::cout << n << "\n"; // Write size of data

    for (i=0;i<n;i++) { // Write data (random numbers)
        std::cout << float(rand())/RAND_MAX << "\n";
    }
}
```

# cin.cpp: read in n numbers

```
#include <iostream>

int main() {

    int i;
    int n;

    // note -- no white space like "\n" or " " for reading!
    std::cin >> n; // read size

    float inputdata[n]; // Dynamically allocate array

    for (i=0;i<n;i++) { // Read data
        std::cin >> inputdata[i];
    }

    std::cout << "Input complete: Last value read in was " << inputdata[n-1] << "\n";
}
```

# cin and cout

Try:

make cout

cout

cout > file.data

Put output into a file

more file.data

make cin

cin < file.data

Read input from a file

cout | cin

Pipe output from cout to cin

# cout, cin > < and pipe |

```
[wadsley@phys-ugrad ~/fileio]$ cout  
10  
0.840188  
0.394383  
0.783099  
0.79844  
0.911647  
0.197551  
0.335223  
0.76823  
0.277775  
0.55397  
[wadsley@phys-ugrad ~/fileio]$ cout > myfile.data  
[wadsley@phys-ugrad ~/fileio]$ cin < myfile.data  
Input complete: Last value read in was 0.55397  
[wadsley@phys-ugrad ~/fileio]$ cout | cin  
Input complete: Last value read in was 0.55397  
[wadsley@phys-ugrad ~/fileio]$
```



# Working with files

cin and cout are not very flexible. In general you want to be able to open and close files with a program to read or write from them – without needing a terminal

```
#include <fstream>
```

Extends io to use files instead of terminal as source of data

# fstream

Instead of using cout and cin, you make a new object that represents the file

```
#include <iostream>    Must include iostream too!
```

```
#include <fstream>
```

```
ofstream outputfile;    a stream object for output
```

```
ifstream inputfile;     a stream object for input
```

# fout.cpp: write to “output.data”

```
#include <iostream>
#include <fstream>
```

```
#include <stdlib.h>
```

```
int main() {
```

```
    int i;
```

```
    const int n=10;
```

I want to write to the file – so I use ofstream



```
    std::ofstream outputfile;
```

```
    outputfile.open( "output.data" );
```

```
    outputfile << n << "\n"; // Write size of data
```

```
    for (i=0;i<n;i++) { // Write data (random numbers)
```

```
        outputfile << float(rand())/RAND_MAX << "\n";
```

```
    }
```

```
    outputfile.close();
```

```
}
```

# ofstream

Direct replacement for `std::cout`

Extra step: open a file first!

Try:

make fout

fout

more output.data

cin < output.data

**fout writes directly to a file**

**The program chooses the filename**

**Unlike cout, I can open as many  
files at once as I like**

# fout

```
[wadsley@phys-ugrad ~/fileio]$ [wadsley@phys-ugrad ~/fileio]$ make fout
c++ -c fout.cpp
c++ -o fout fout.o -ltrapfpe
[wadsley@phys-ugrad ~/fileio]$ fout
[wadsley@phys-ugrad ~/fileio]$ more output.data
10
0.840188
0.394383
0.783099
0.79844
0.911647
0.197551
0.335223
0.76823
0.277775
0.55397
[wadsley@phys-ugrad ~/fileio]$
```

# fin.cpp: read from “output.data”

```
#include <iostream>
#include <fstream>

int main() {
    int i;
    int n;
    std::ifstream inputfile;
    inputfile.open("output.data" );
    // note -- no white space like "\n" or " " for reading!
    inputfile >> n; // read size
    float inputdata[n]; // Dynamically allocate array

    for (i=0;i<n;i++) { // Read data
        inputfile >> inputdata[i];
    }
    inputfile.close();
    std::cout << "Input complete: Last value read in was " << inputdata[n-1] << "\n";
}
```

# ifstream

Direct replacement for `std::cin`

Extra step: open a file first (file must exist!)

Try:

make fin

fout

fin      Note: fin is hardwired to read output.data

# fin\_ask.cpp    Ask the user

```
#include <iostream>
#include <fstream>
#include <string>
int main() {
    int i, n;
    std::string filename;
    std::cout << "What file do you want to open?\n";
    std::cin >> filename;
    std::ifstream inputfile;

    inputfile.open(filename.c_str());    Note: open() needs a char string: convert
                                        filename using .c_str() function

    inputfile >> n; // read size
    float inputdata[n]; // Dynamically allocate array
    for (i=0;i<n;i++) { // Read data
        inputfile >> inputdata[i];
    }
    inputfile.close();
    std::cout << "Input complete: Last value read in was " << inputdata[n-1] << "\n";
```



# fstream: Did it work?

fstream objects have extra functions to test if things worked.

e.g.

```
#include <iostream>
```

```
#include <fstream>
```

```
int main () {
```

```
    std::ofstream myfile ("example.txt");
```

```
    if (myfile.is_open()) {
```

```
        // Test that it opened
```

```
    }
```

```
    else std::cout << "Unable to open file";
```

# fin\_ask.cpp Test if open worked

```
#include <iostream>
#include <fstream>
#include <string>
int main() {
    int i,n;;
    std::string filename;
    std::cout << "What file do you want to open?\n";
    std::cin >> filename;

    std::ifstream inputfile;
    inputfile.open(filename.c_str());
    if (!inputfile.is_open()) {
        std::cout << "Could not open file: " << filename << "\n";
        return 1;
    }
}
```

# ifstream

fin\_ask.cpp uses cin to ask the user for a filename

Try:

make fin\_ask

fin\_ask      ask it to open output.data

# fstream basic functions

<code>open("name")</code>	open a file
<code>close()</code>	close the file
<code>is_open()</code>	Is it open?
<code>bad()</code>	Something went wrong?
<code>fail()</code>	Failed to read/write or bad?
<code>eof()</code>	Are we at the end of the file?
<code>good()</code>	OK to keep reading/writing?

# Example of a problem

Try:

fin\_ask            Tell it: short.data

short.data sets size 10 but only has 4 numbers  
to read

make fin\_test

fin\_test           Tell it: short.data

This program checks after each read

# fin\_test.cpp Test if read worked

```
inputfile >> n; // read size
float inputdata[n]; // Dynamically allocate array

for (i=0;i<n;i++) { // Read data
    inputfile >> inputdata[i];
    if (inputfile.fail()) {
        std::cout << "Read incomplete: at line: " << i << "\n";
        if (i>0) std::cout << "Last good data value " <<
            inputdata[i-1] << "\n";
        return 1;
    }
}
inputfile.close();
```

# File I/O

You can open as many files as you want at once.

You can control the way the files are used and how the system responds if files do or don't exist, can't be written etc...

# Options for fstream

Open can be used with 2 arguments

e.g.

```
ifstream myfile;
```

```
myfile.open("somefile.dat", mode );
```

With mode you can specify what to do about the file if it exists, expectations about data format and so on



# Mode options

- `std::ios::in` Open for input operations.
- `std::ios::out` Open for output operations.
- `std::ios::binary` Open in binary mode.
- `std::ios::ate` Set the initial position at the end of the file.
- `std::ios::app` All output operations are appended to end of file (only for output)
- `std::ios::trunc` If the file opened for output operations, delete previous

# Using mode options

```
std::ofstream myfile;  
myfile.open("somefile.dat", std::ios::app );
```

Open file to append to end of file (don't overwrite what's already there)

Default is to overwrite from the start

# fstream

`fstream` is the general version of a file stream

- You can create a stream for input or output as you wish

```
fstream myfile;  
myfile.open("somefile.dat", std::ios::in );
```

- `ifstream` assumes input (reading a file)
- `ofstream` assumes output (writing a file)

# Binary output

Writing text is only useful if a human wants to read it. For the computer, it requires translating what is in its memory into characters

You can write memory contents directly. It is faster and more efficient (smaller files).

It is also exact – you save precisely what the computer stored in memory

# binaryout.cpp

```
#include <iostream>
#include <fstream>
#include <stdlib.h>

int main() {
    int i;
    const int n=10;

    std::ofstream outputfile;
    outputfile.open( "output.bin", std::ios::binary );
    outputfile.write( (char *) &n, sizeof(n) );

    for (i=0;i<n;i++) { // Write data (random numbers)
        float data = float(rand())/RAND_MAX;
        outputfile.write( (char *) &data, sizeof(data) );
    }
    outputfile.close();
}
```

# Working with binary

Don't use << >> operators. They are for translating variables into text

Directly read and write memory from/to a file:

```
outputfile.open( "output.bin", std::ios::binary );  
outputfile.write( (char *) stuff, nBytes );
```

C/C++ char variables are equivalent to bytes

The fstream read/write just want a memory location and how many bytes to read/write!

# Working with binary

Direct write memory to a file:

Write one integer n:

```
int n = 42;
```

```
outputfile.write( (char *) &n, sizeof(n) );
```

(char \*) says “treat the address of n as a pointer to raw bytes”. It doesn’t change the address.

sizeof(n) is how many bytes in an int ( 4 )

This writes 4 bytes directly from memory to the file

# Working with binary

Direct write memory to a file:

Write an array of 100 floats:

```
float a[100];
```

```
outputfile.write( (char *) a, 100*sizeof(a[0]) );
```

(char \*) => treat array as a pointer to raw bytes

sizeof(a[0]) is how many bytes in an float ( 4 )

Writes 400 bytes directly from memory to the file

```
outputfile.write( (char *) a, 100*sizeof(float) );
```

also ok



# Reading Binary

Direct read memory to a file:

Write an array of 100 floats:

```
float a[100];
```

```
outputfile.read( (char *) a, 100*sizeof(a[0]) );
```

(char \*) converts array a into a pointer to raw bytes

sizeof(a[0]) is how many bytes in an float ( 4 )

Writes 400 bytes directly from the file to memory

# Reading Binary

Direct read memory to a file:

Write an array of n floats:

```
float a[n];
```

```
outputfile.read( (char *) a, n*sizeof(a[0]) );
```

C/C++ lets you decide the array size a the last minute and then read data into that memory

(char \*) => treat array a into a pointer to raw bytes

sizeof(a[0]) is how many bytes in an float ( 4 )

Writes 4 n bytes directly from the file to memory

# binaryin.cpp

```
#include <iostream>
#include <fstream>

int main() {
    int i;
    int n;

    std::ifstream inputfile;

    inputfile.open("output.bin" );
    inputfile.read( (char *) &n, sizeof(n)); // read size

    float inputdata[n]; // Dynamically allocate array
    inputfile.read( (char *) inputdata, n*sizeof(inputdata[0]));
    inputfile.close();

    std::cout << "Input complete: Last value read in was " << inputdata[n-1] << "\n";
}
```

# binaryin.cpp

```
#include <iostream>
#include <fstream>

int main() {
    int i;
    int n;


    std::ifstream inputfile;

    inputfile.open("output.bin" );
    inputfile.read( (char *) &n, sizeof(n)); // read size

    float inputdata[n]; // Dynamically allocate array
    inputfile.read( (char *) inputdata, n*sizeof(inputdata[0]));
    inputfile.close();

    std::cout << "Input complete: Last value read in was " << inputdata[n-1] << "\n";
}
```

Note how compact and efficient this binary read is -- just read all 4n bytes into memory in one go!



# Look at the binary files!

Try:

make binaryout

binaryout

fout

ls -l output\*      Which is bigger?

make binaryin

binaryin

make binaryin\_test      testing version

# Raw text is bigger

```
[wadsley@phys-ugrad fileio]$ ls -l out* short*  
-rw-rw-r-- 1 wadsley wadsley 44 Oct  5 02:13 output.bin  
-rw-rw-r-- 1 wadsley wadsley 90 Oct  5 02:13 output.data  
-rw-rw-r-- 1 wadsley wadsley 20 Oct  5 02:13 short.bin  
-rw-rw-r-- 1 wadsley wadsley 38 Oct  5 02:13 short.data
```

If you write more decimal places text (output.data) would be even bigger

Binary data contains exactly what is in memory and no more (output.bin)

# Looking at binary: od command

```
[wadsley@phys-ugrad ~/fileio]$ od -i output.bin
```

Integers

```
0000000  10 1062672011 1053420687 1061714225
0000020 1061971601 1063870905 1045056232 1051435601
0000040 1061464754 1049507965 1057870074
0000054
```

floats

```
[wadsley@phys-ugrad ~/fileio]$ od -f output.bin
```

```
0000000  1.401298e-44  8.401877e-01  3.943829e-01  7.830992e-01
0000020  7.984400e-01  9.116474e-01  1.975514e-01  3.352228e-01
0000040  7.682296e-01  2.777747e-01  5.539700e-01
0000054
```

```
[wadsley@phys-ugrad ~/fileio]$ more output.data
```

```
10
0.840188
0.394383
0.783099
0.79844
0.911647
```

Compare to  
text file

# binaryin\_test.cpp

```
inputfile.read( (char *) &n, sizeof(n)); // read size
std::cout << n << " data values to read\n";

float inputdata[n]; // Dynamically allocate array

if (!inputfile.read( (char *) inputdata, n*sizeof(inputdata[0]))) {
    std::cout << "Read incomplete: only read " << inputfile.gcount() <<
        " bytes \n";
    int nread = inputfile.gcount()/sizeof(inputdata[0]);
    if (nread > 0)
        std::cout << "Last good data value " << inputdata[nread-1] << "\n";
    return 1;
}
inputfile.close();
```



# Testing for success

`read()` and `write()` return 0 if there was a problem

```
if (myfile.read( stuff, nBytes )) { // success! }
```

You can also ask how much was read

```
std::cout << myfile.gcount()) << " bytes read";
```

This can be converted into a number of values with `sizeof`

Try `binaryin_test` with file: `short.bin`

# binaryin\_test.cpp

```
inputfile.read( (char *) &n, sizeof(n)); // read size
std::cout << n << " data values to read\n";

float inputdata[n]; // Dynamically allocate array

if (!inputfile.read( (char *) inputdata, n*sizeof(inputdata[0]))) {
    std::cout << "Read incomplete: only read " << inputfile.gcount() <<
        " bytes \n";
    int nread = inputfile.gcount()/sizeof(inputdata[0]);
    if (nread > 0)
        std::cout << "Last good data value " << inputdata[nread-1] << "\n";
    return 1;
}
inputfile.close();
```

# Binary:

## Big Endian and Little Endian

With text there are conventions on how to order numbers: 100 means one hundred (rather than writing 001)

Files consist of bytes

With raw binary of a 4 byte number do you put the big bytes first or the little ones?

Unfortunately chip makers didn't agree

# endian.cpp

```
#include <iostream>
#include <fstream>
int main() {
```

```
    int i=54321;
```

```
    std::ofstream outputfile;
    outputfile.open( "endian.bin", std::ios::binary );
    outputfile.write( (char *) &i, sizeof(i) );
    outputfile.close();
}
```

Writes a 4-byte integer to  
the file: endian.bin  
value: 54321

# Unformatted:

## Big Endian and Little Endian

$$54321 = 256 * 212 + 49$$

Big Endian

0	0	212	49
---	---	-----	----

Little Endian

49	212	0	0
----	-----	---	---

Many compilers have flags to choose how files are written

Try `od -t u1 endian.bin`

What endianness is Intel (phys-ugrad)?

# Endianness

- Big Endian

Apple\*, SUN (Sparc), Internet Standard

- Little Endian

HP (alpha), **Intel, AMD**      \*macbooks use intel

Internet packets are converted to big-endian  
regardless of your computer for compatibility  
(using the xdr library)

For most purposes you don't have to care 😊